

File Name: T-ICML-O_3_I1_the_data_scarcity_problem

Format: Presenter in Studio

Presenter: Max Lotstein



Google Cloud

Dealing with Data Scarcity

Max Lotstein

Agenda

The Data Scarcity Problem

Data Augmentation

Transfer Learning

No Data, No Problem

Learn how to...

Calculate the number of parameters for ML models
and understand why the data needs grow with the
number of parameters

Understand how data augmentation can improve
performance if done thoughtfully

Understand how to add data augmentation to a
model

Apply understanding of CNNs to the concept of task
dependence

Understand how transfer learning decreases the
need for data

Data Need Grows with
Model Complexity

A simple linear model

```
def linear_model(img):
    X = tf.reshape(img, [-1,HEIGHT*WIDTH])
    W = tf.get_variable("W",
        [HEIGHT*WIDTH, NCLASSES])
    b = tf.get_variable("b",NCLASSES)
    ylogits = tf.matmul(X,W)+b
    return ylogits, NCLASSES
```

How many parameters in our linear model?

1. HEIGHT * WIDTH * batch_size
2. HEIGHT * WIDTH * NCLASSES +
NCLASSES
3. HEIGHT * WIDTH * NCLASSES
4. HEIGHT * WIDTH + NCLASSES

How many parameters in our linear model?

1. HEIGHT * WIDTH * batch_size
2. **HEIGHT * WIDTH * NCLASSES +
NCLASSES**
3. HEIGHT * WIDTH * NCLASSES
4. HEIGHT * WIDTH + NCLASSES

A simple DNN

```
def dnn_model(img, mode, hparams):
    X = tf.reshape(img, [-1, HEIGHT*WIDTH])
    h1 = tf.layers.dense(X, 300)
    h2 = tf.layers.dense(h1, 100)
    h3 = tf.layers.dense(h2, 30)
    ylogits = tf.layers.dense(h3, NCLASSES,
                             activation= None)
    return ylogits, NCLASSES
```

How many parameters are in the h1 layer?

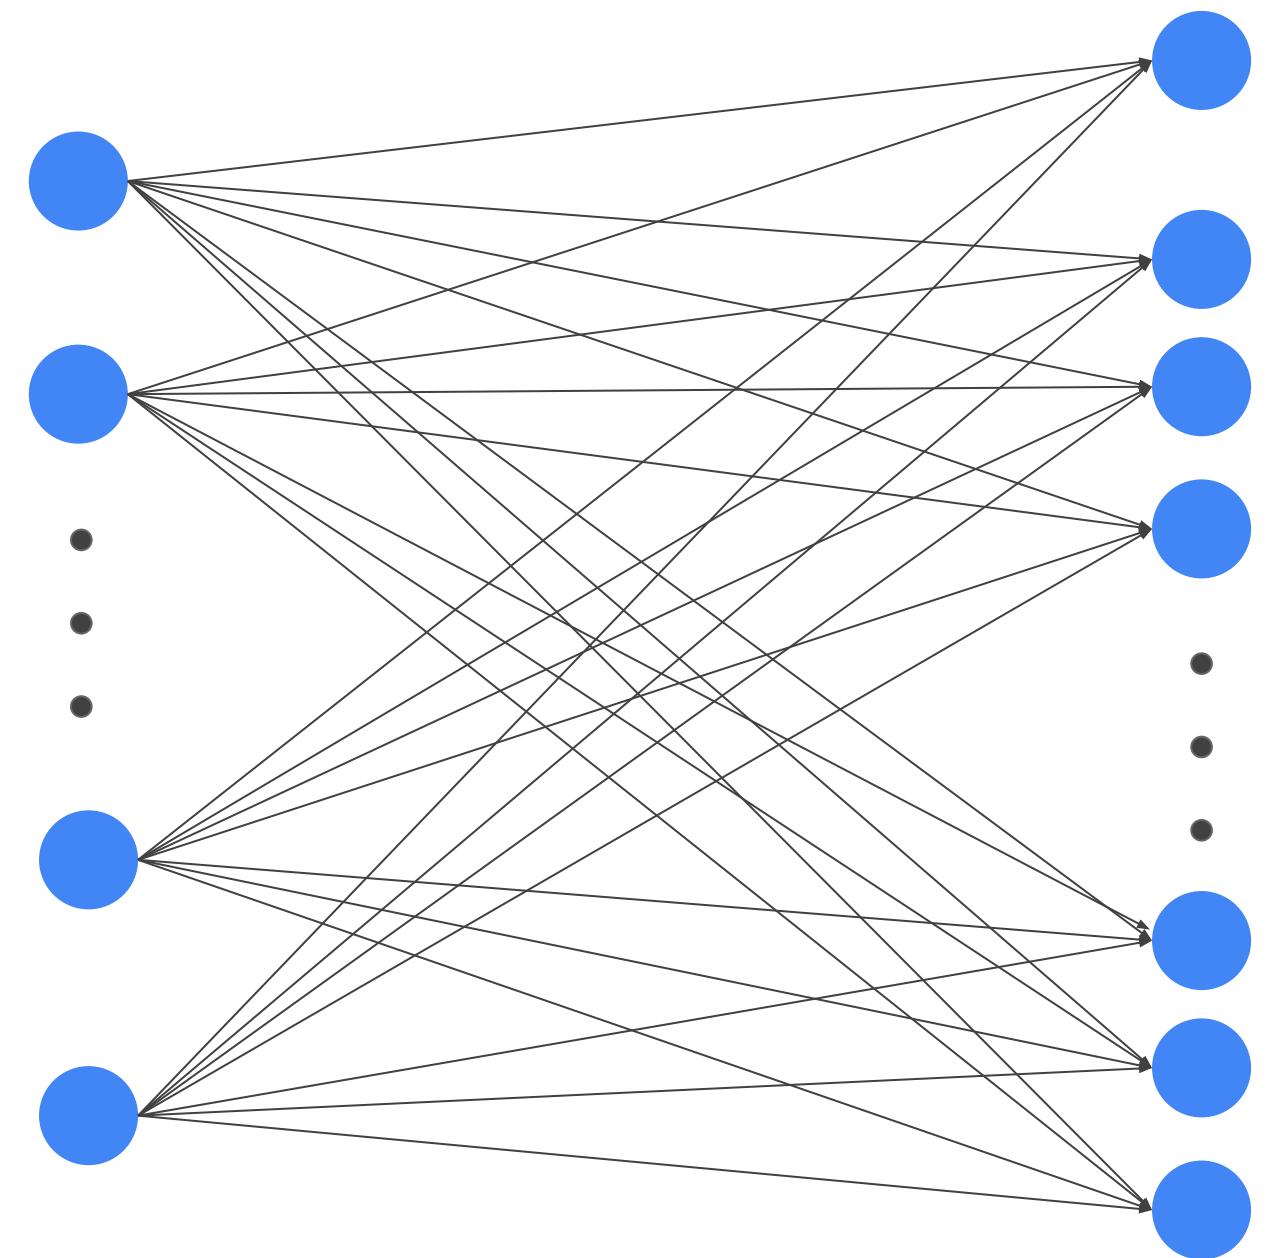
```
x = tf.reshape(img, [-1, HEIGHT*WIDTH])  
h1 = tf.layers.dense(x, 300)
```

1. HEIGHT * WIDTH * 300
2. HEIGHT * WIDTH * 100 + 100
3. HEIGHT * WIDTH * 300 * NCLASSES +
300
4. HEIGHT * WIDTH * 300 + 300

How many parameters are in the h1 layer?

```
# X.shape = [-1, HEIGHT*WIDTH])  
h1 = tf.layers.dense(X, 300)
```

1. HEIGHT * WIDTH * 300
2. HEIGHT * WIDTH * 100 + 100
3. HEIGHT * WIDTH * 300 * NCLASSES +
300
4. HEIGHT * WIDTH * 300 + 300



HEIGHT *
WIDTH

300

Parameters by model type for MNIST

Model Type	Number of Parameters
Toy linear model	7850
Toy DNN	270,000

A simple CNN model

```
import tensorflow.layers
def cnn_model(img, mode, params):
    c1 = conv2d(img, filters=10,
               kernel_size=3, strides=1,
               padding='same')
    p1 = max_pooling2d(c1,pool_size=2, strides=2)
    c2 = conv2d(p1, filters=20,
               kernel_size=3, strides=1,
               padding='same')
    p2 = max_pooling2d(c2,pool_size=2, strides=2)

    outlen = p2.shape[1]*p2.shape[2]*p2.shape[3]
    p2flat = tf.reshape(p2, [-1, outlen]) # flattened

    h3 = dense(p2flat, 300)
    ylogits = dense(h3, NCLASSES, activation=None)
    return ylogits, NCLASSES
```

How many parameters are in the c1 layer?

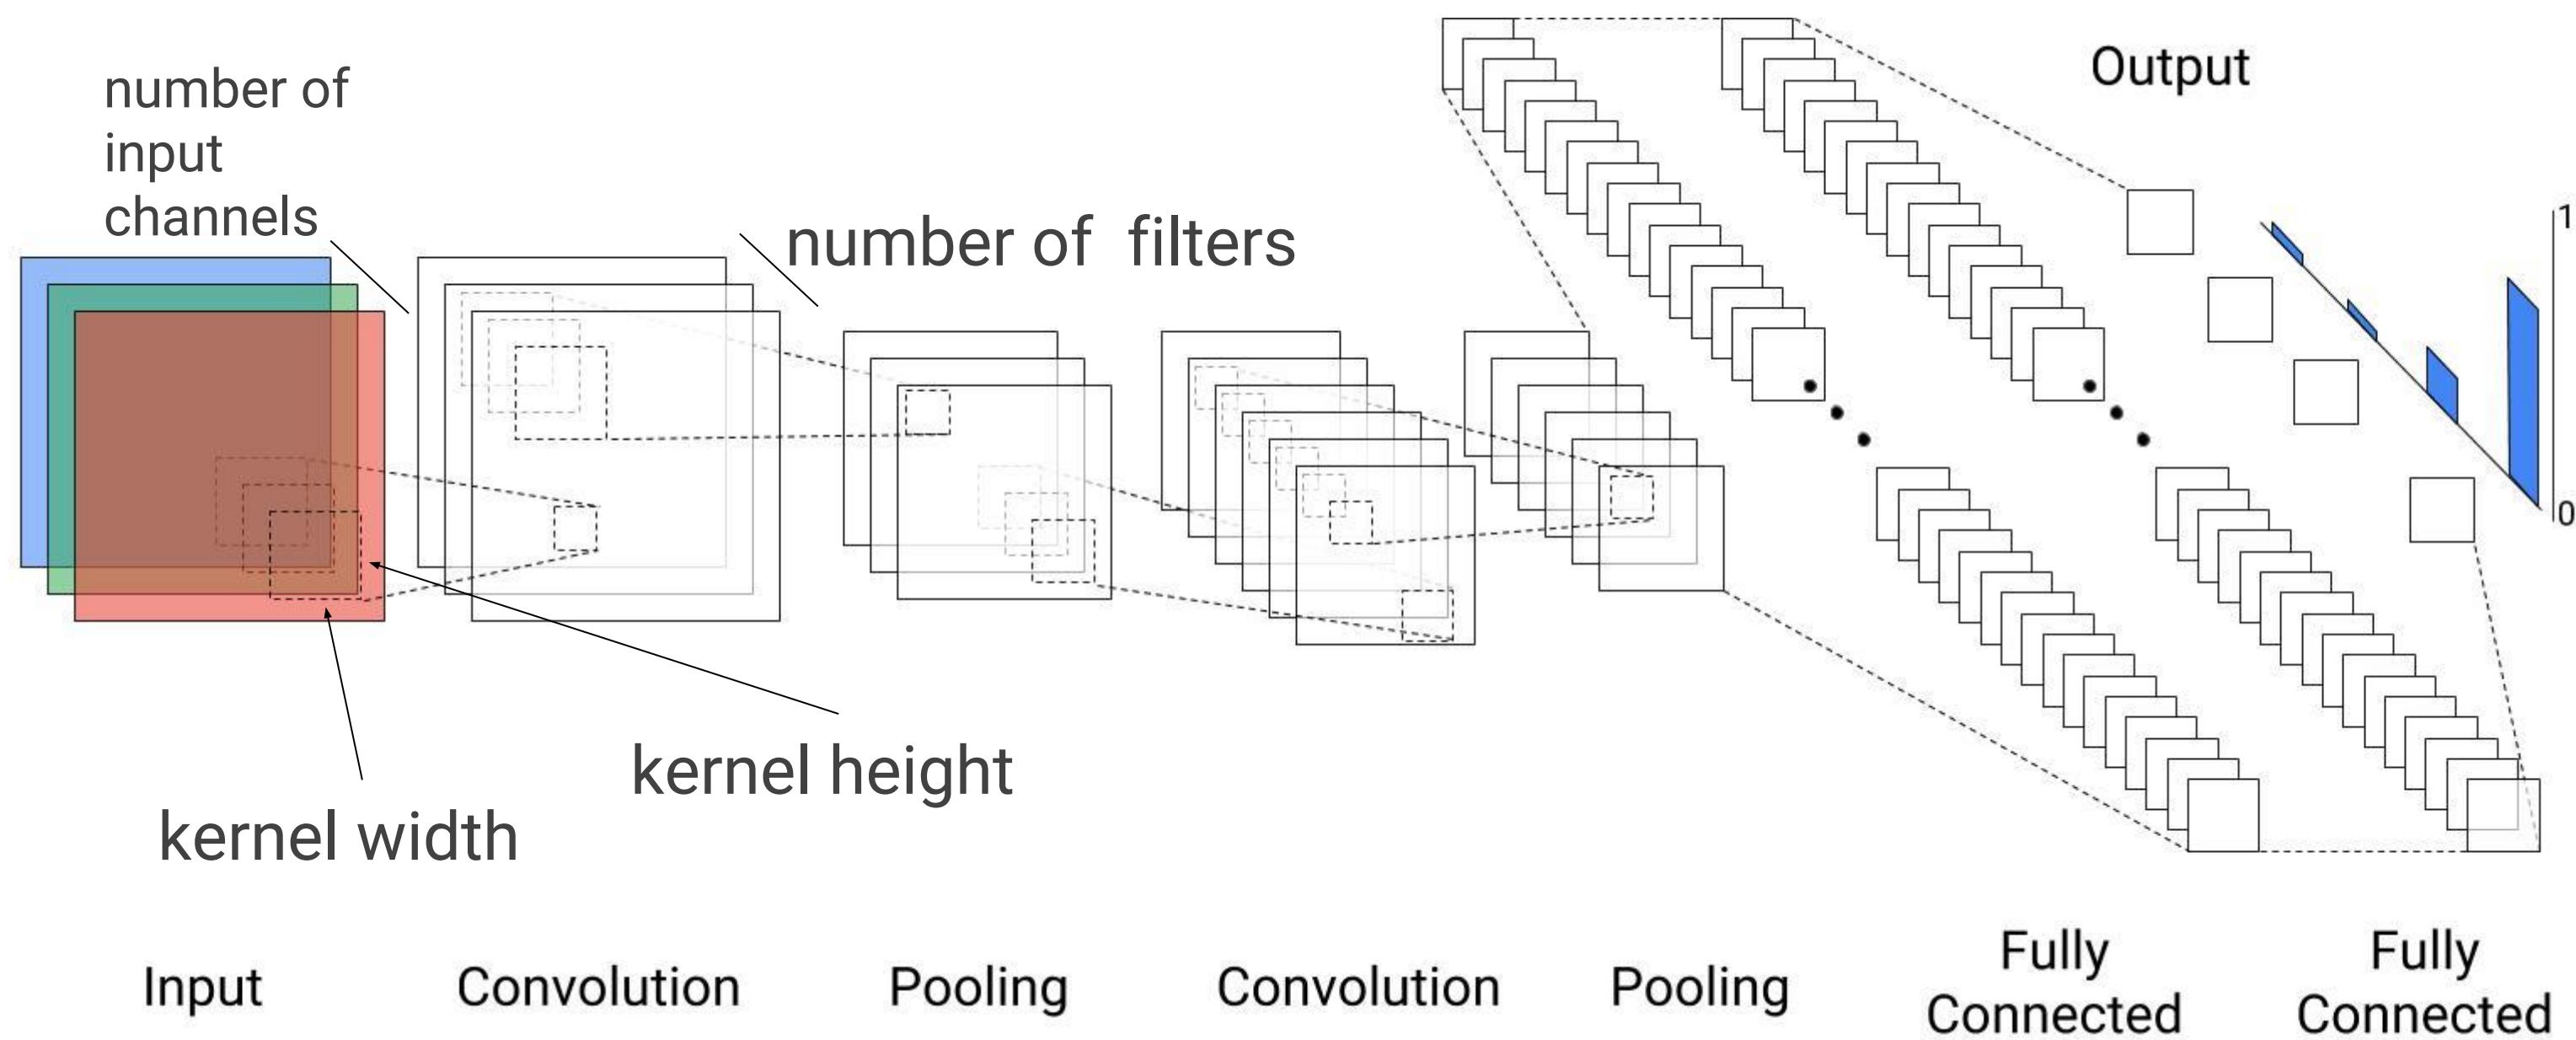
```
# img.shape = [-1, HEIGHT, WIDTH]
c1 = conv2d(img, filters=10,
            kernel_size=3, strides=1,
            padding='same')
```

1. $\text{HEIGHT} * \text{WIDTH} * 10 * 3 + 10$
2. $\text{HEIGHT} * \text{WIDTH} * 10 * 9 + 90$
3. $20 * 9 + 20$
4. $10 * 9 + 10$

How many parameters are
in the c1 layer?

1. HEIGHT * WIDTH * 10 * 3 + 10
2. HEIGHT * WIDTH * 10 * 9 + 90
3. 20 * 9 + 20
4. **10 * 9 + 10**

Computing the number of parameters in a CNN



Parameters by model type for MNIST

Model Type	Number of Parameters
Toy linear model	7,850
Toy DNN	270,000
Toy CNN	300,000

Real-world models have even more parameters

Model	Year	Number of Parameters
AlexNet	2012	60 million
VGGNet	2014	138 million
GoogLeNet	2014	23 million
ResNet	2015	25 million

What to do when you don't have enough labeled data?

When you have some data:

- Data augmentation
- Transfer learning

File Name: T-ICML-O_3_I2_data_augmentation

Format: Presenter in Studio

Presenter: Max Lotstein

Agenda

The Data Scarcity Problem

Data Augmentation

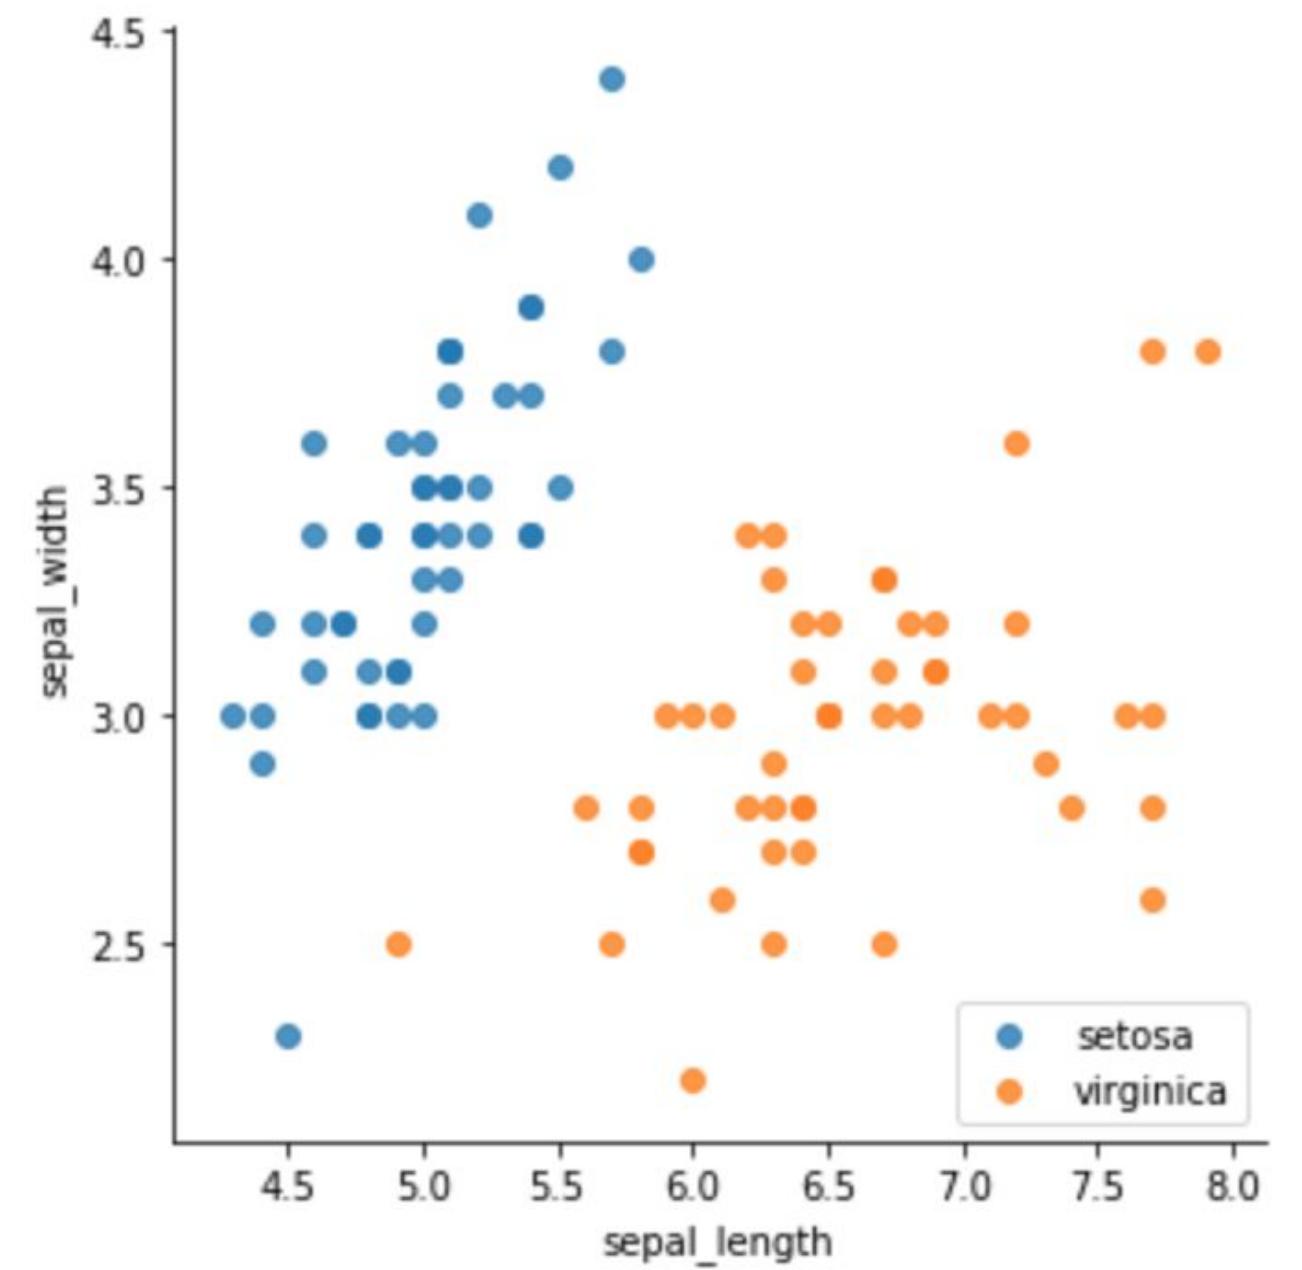
Transfer Learning

No Data, No Problem

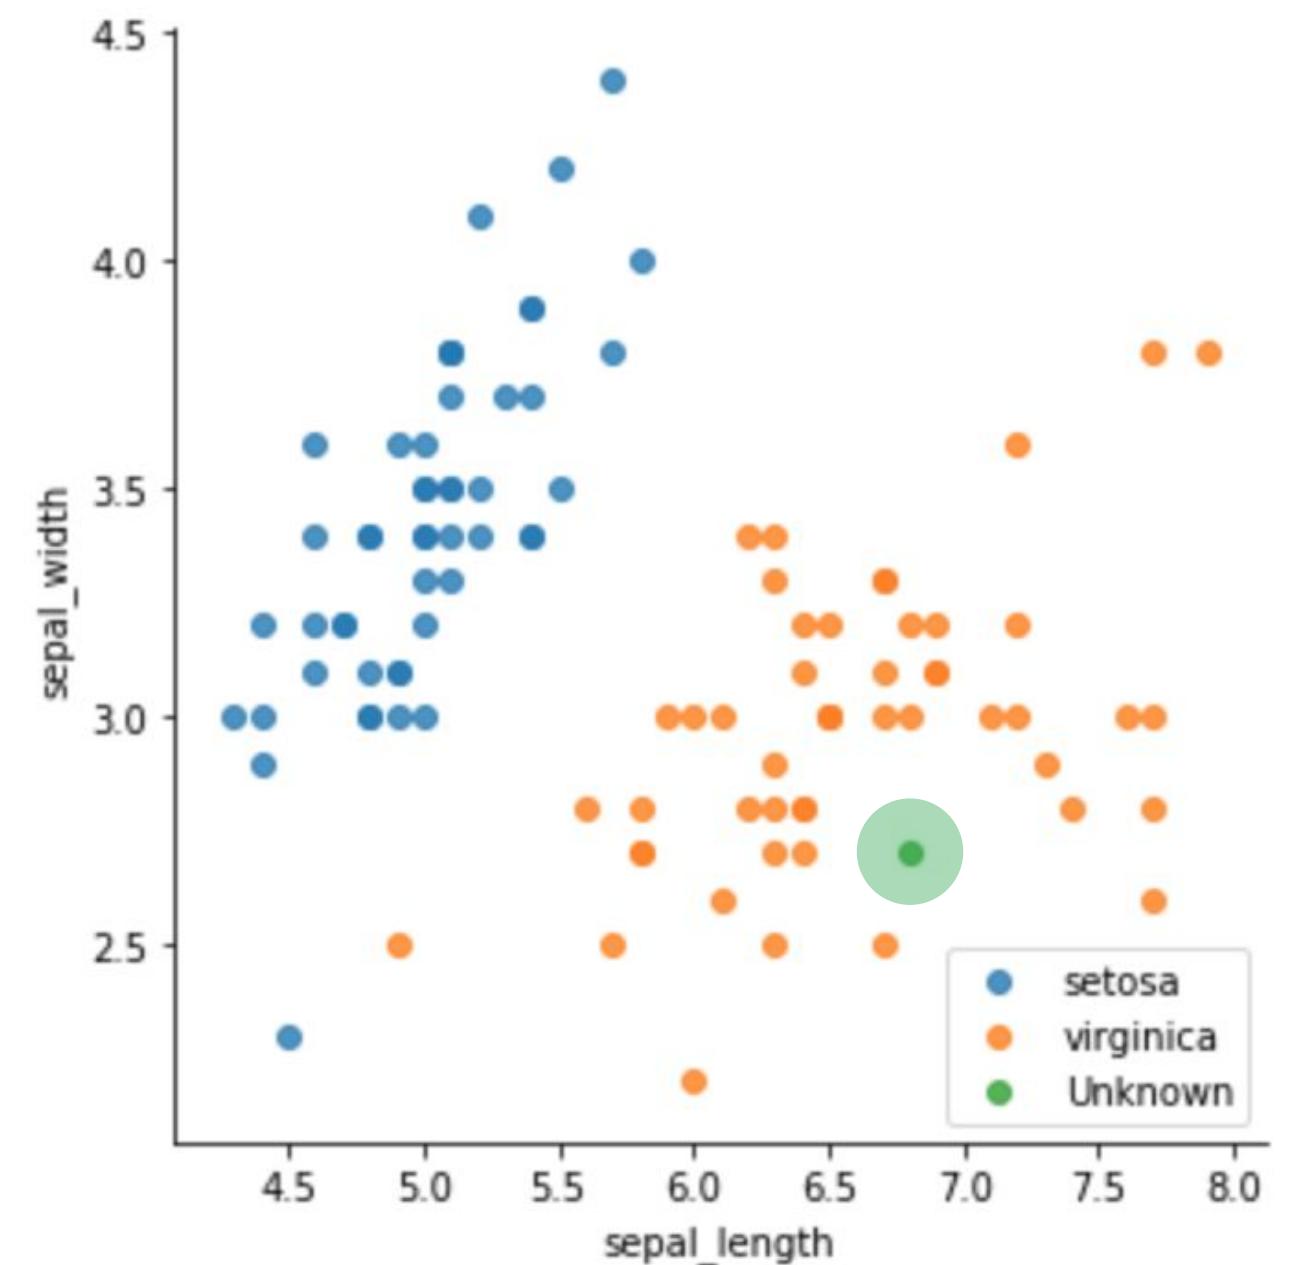
The iris dataset is structured

sepal_length	sepal_width	species
5.1	3.5	setosa
4.9	3.0	setosa
4.7	3.2	setosa
4.6	3.1	setosa
5.0	3.6	setosa

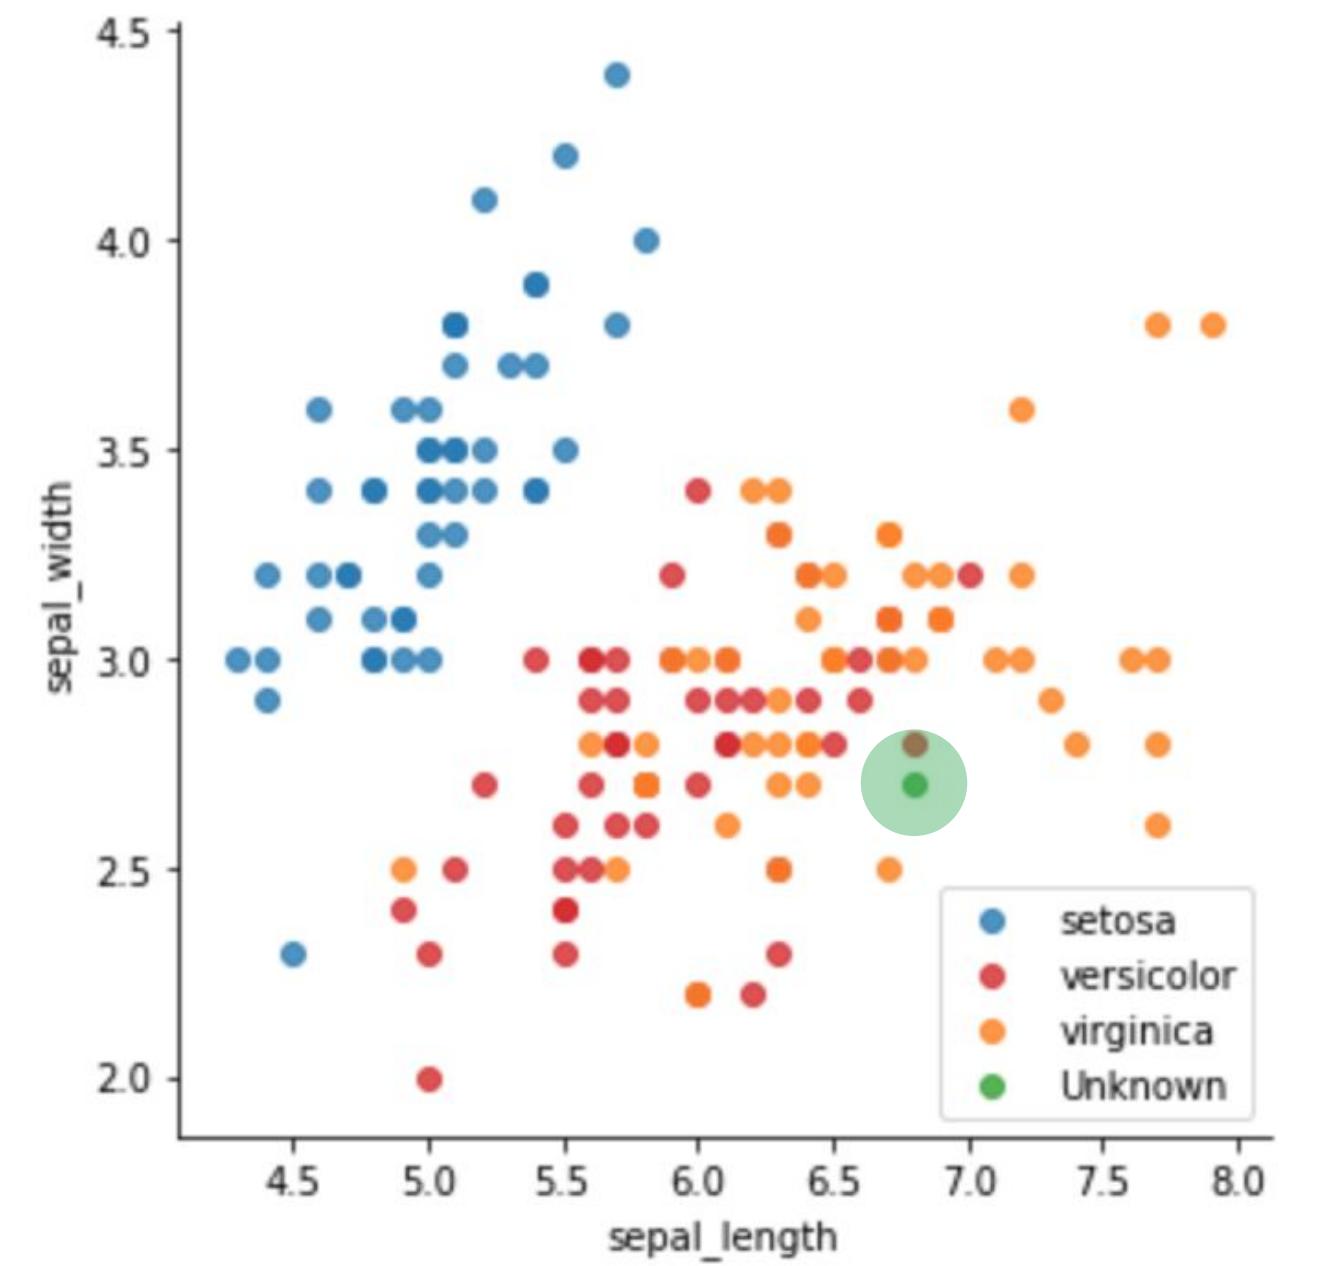
A typical labeled, structured dataset



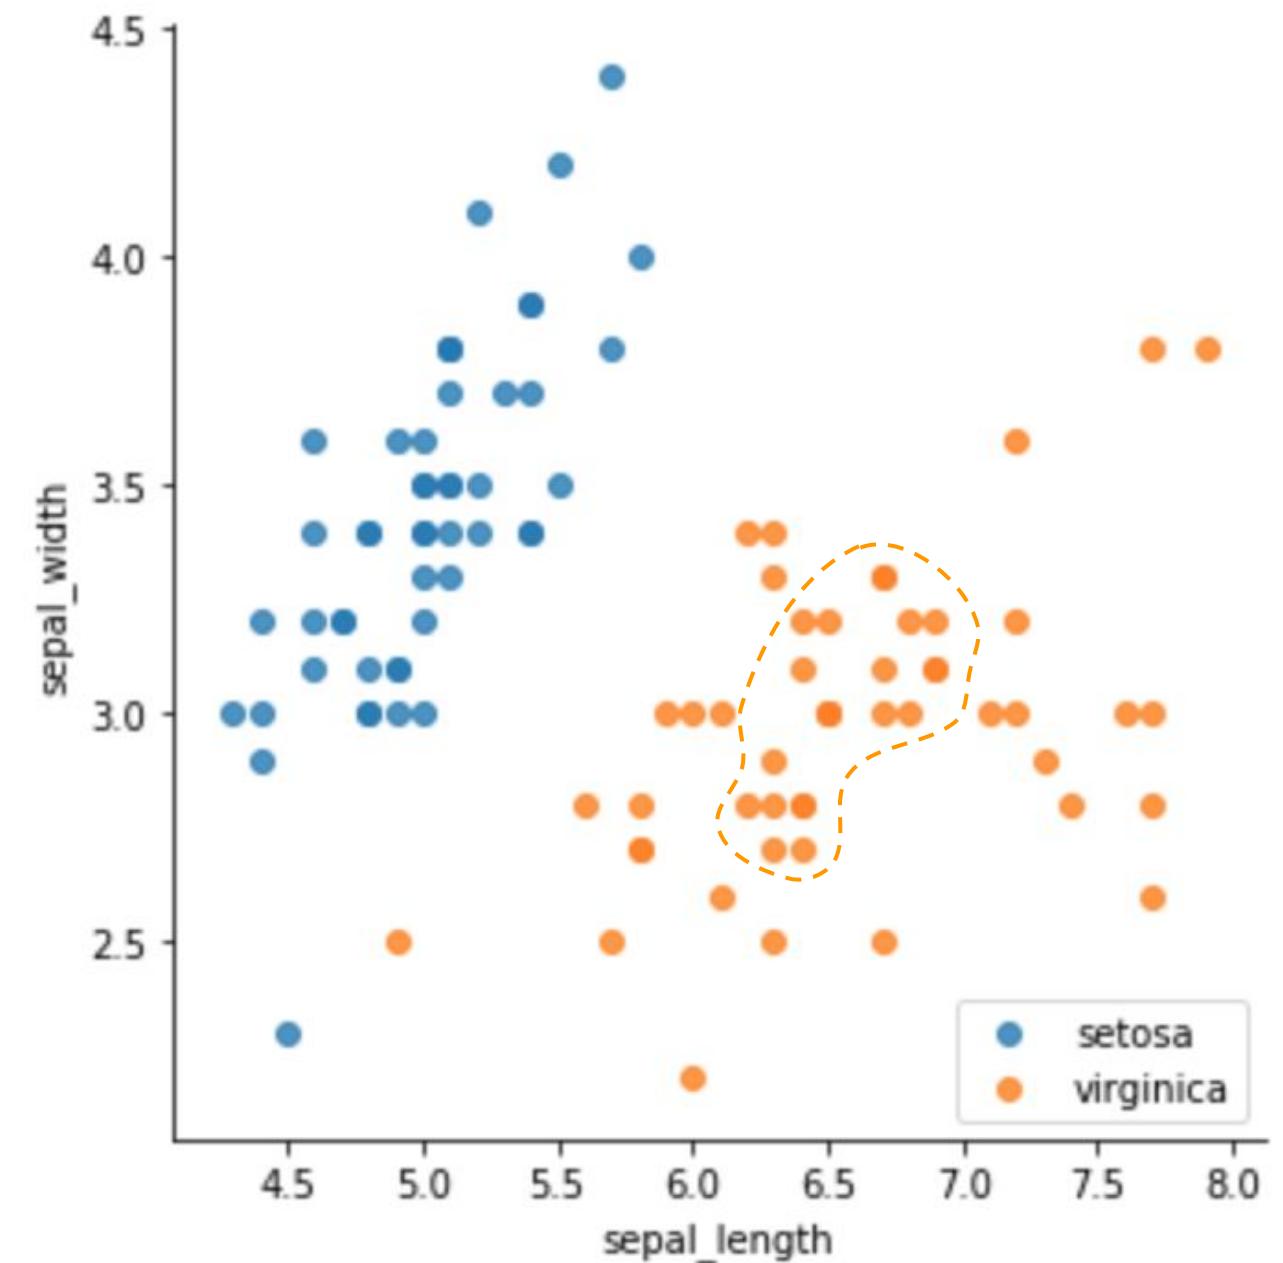
What about this new point?



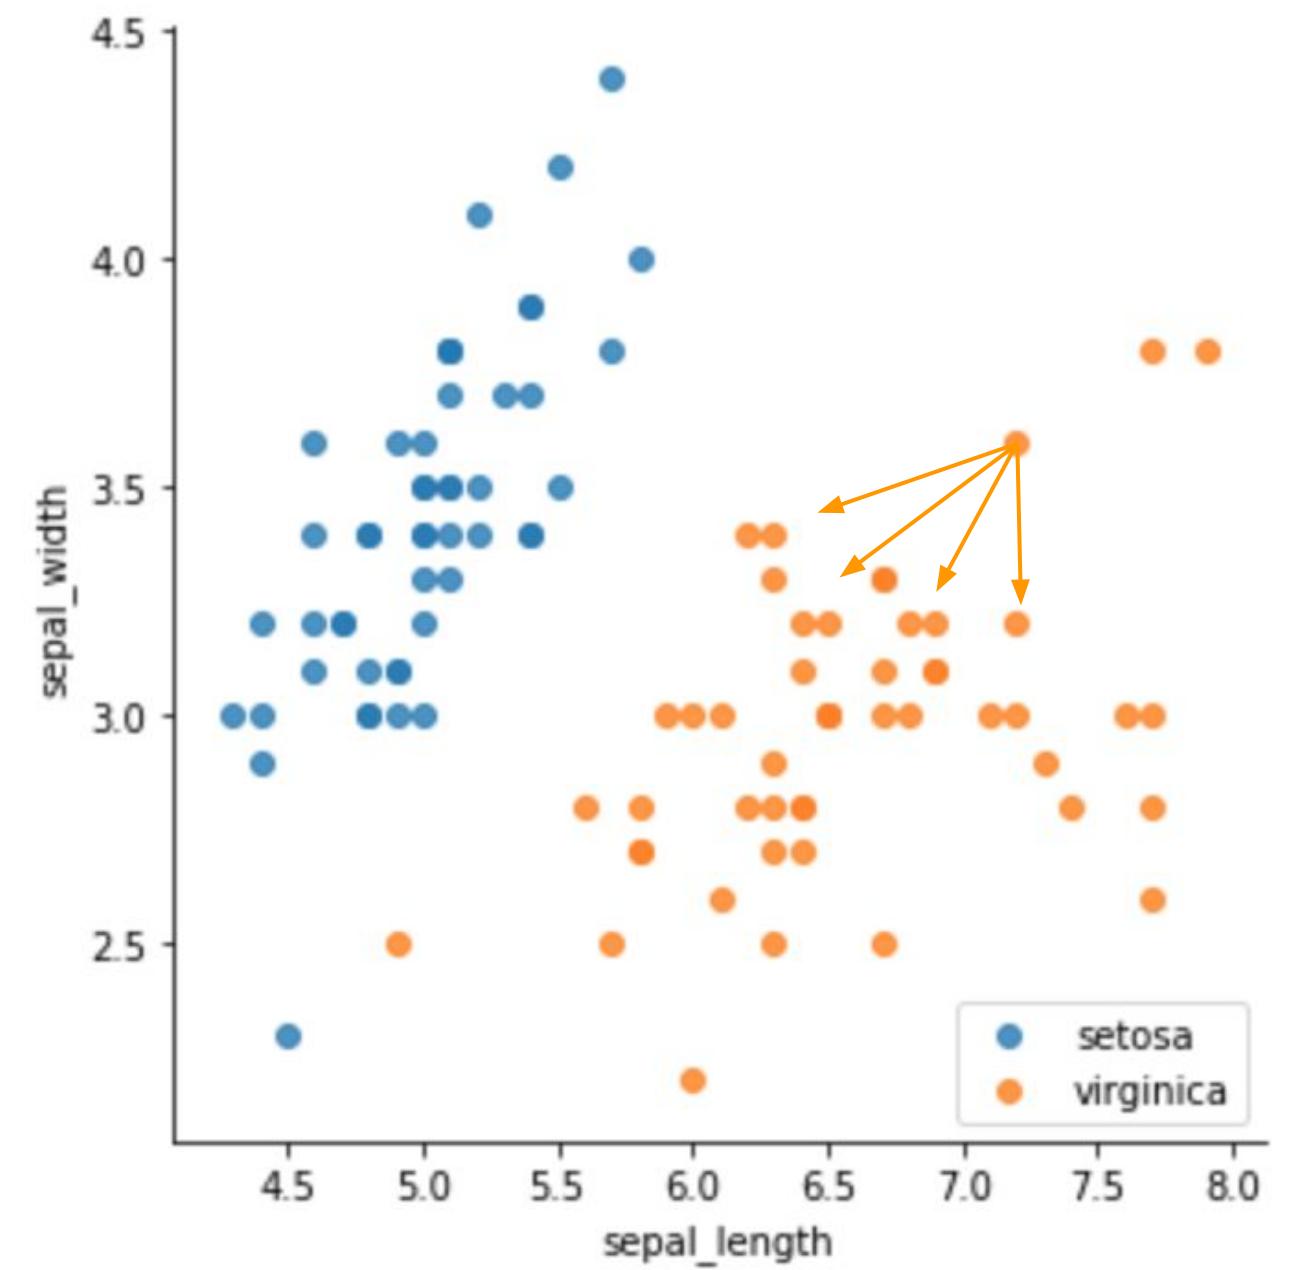
What about now?



One strategy is took look
for clusters



Another is to perturb
individual data points



Common image augmentation techniques



Common image augmentation techniques



Common image augmentation techniques



Common image augmentation techniques



Common image augmentation techniques



Common image augmentation techniques



Common image augmentation techniques



Common image augmentation techniques



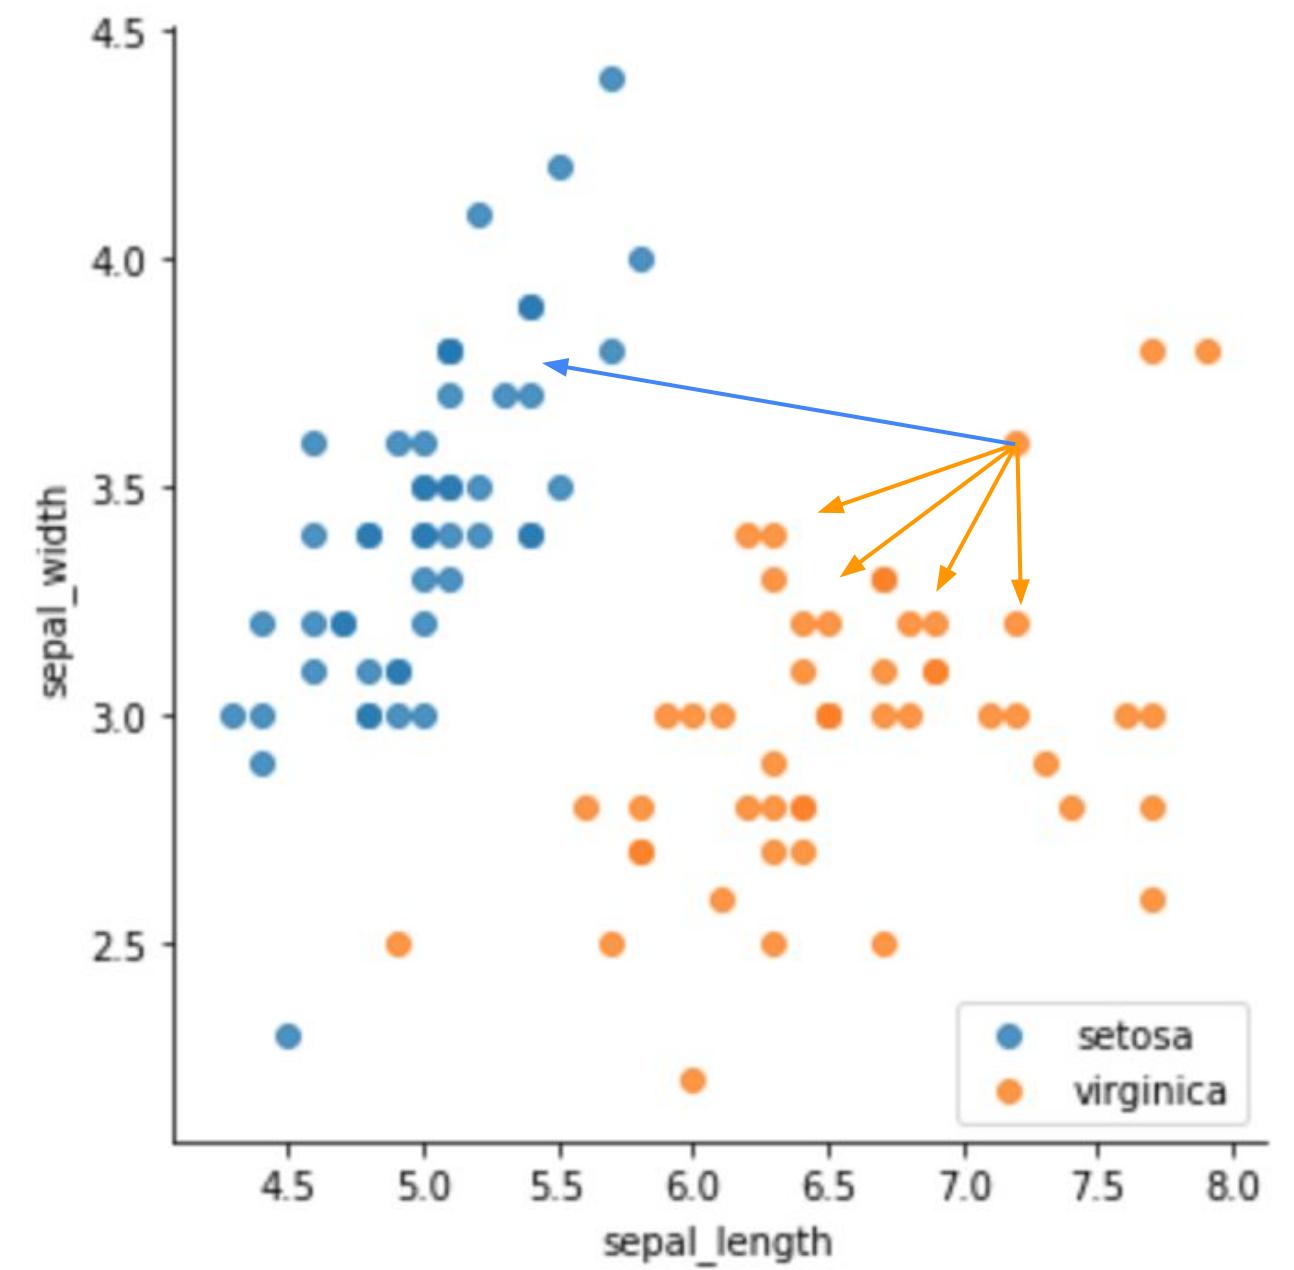
Common image augmentation techniques



Common image augmentation techniques



Data augmentation requires care



Sometimes color is
informative



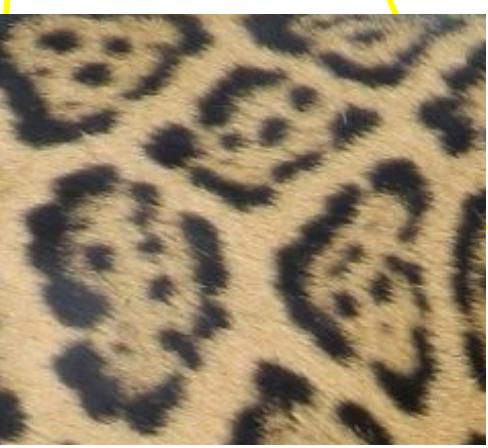
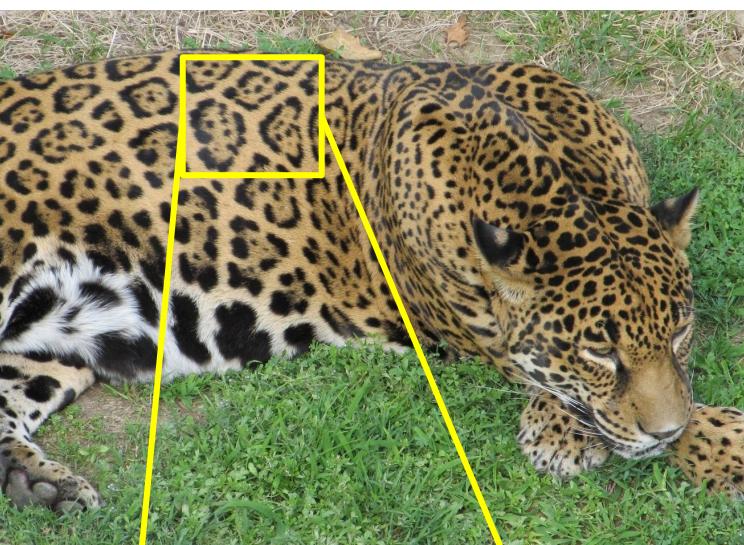
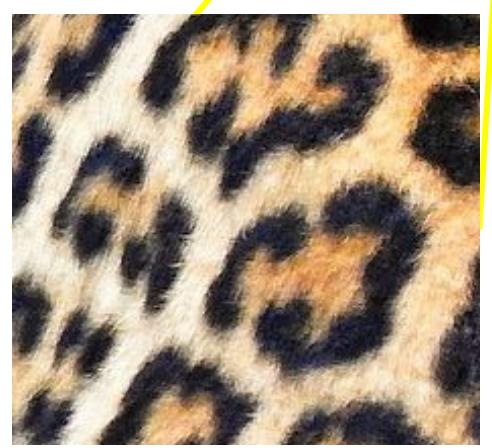
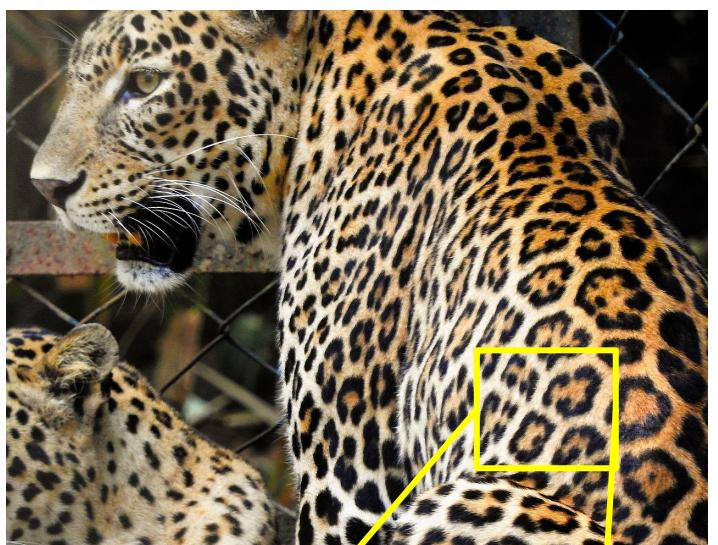
Sometimes color is
informative



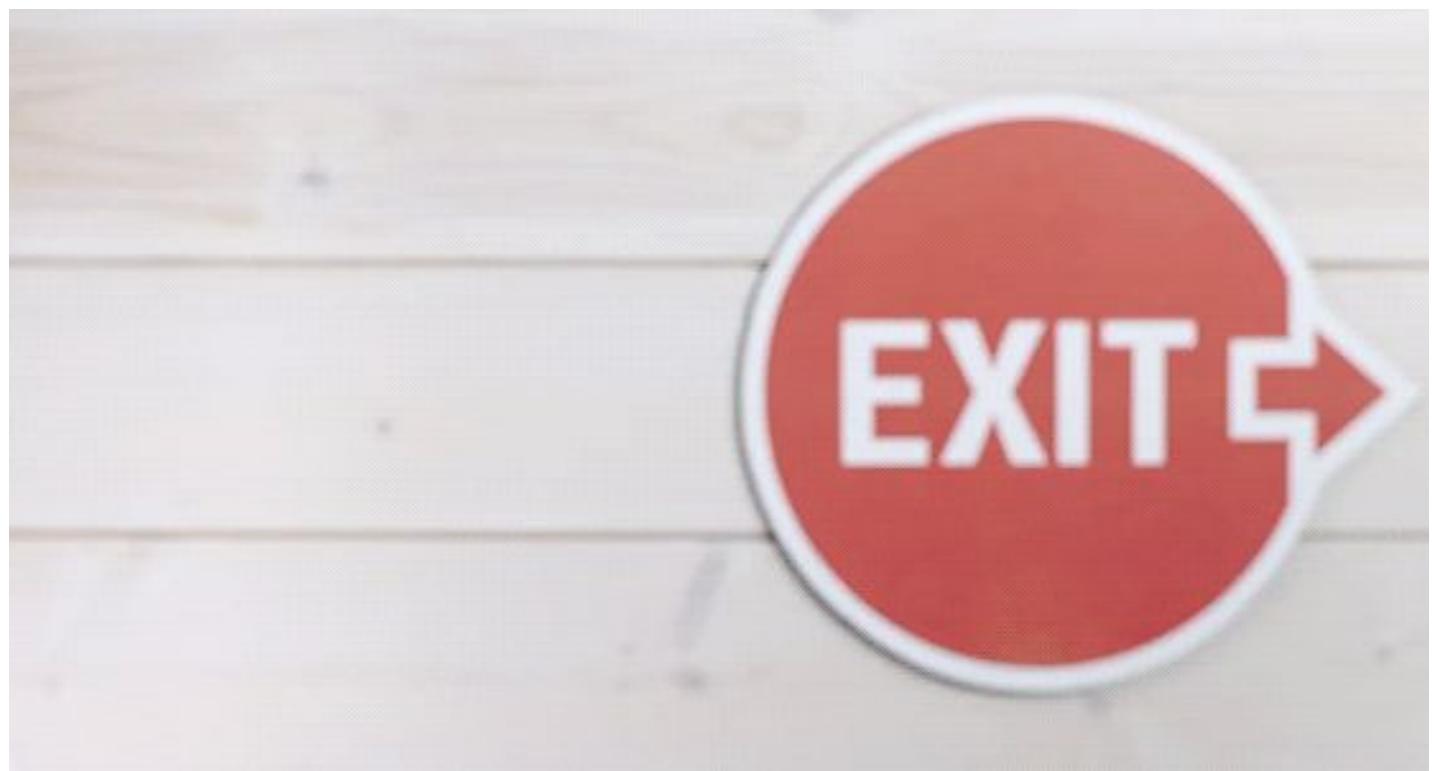
Sometimes orientation is informative



Sometimes small detail is informative



Don't make the problem
harder than it needs to be



What would happen if you
augmented data at
decision time?

Would this be a good idea?

Let's say that our task is to accept a professionally photographed flower and classify it by species.

Would randomly changing the brightness and contrast during data augmentation likely improve performance?

1. Yes
2. No

Would this be a good idea?

Let's say that our task is to accept a professionally photographed flower and classify it by species.

Would randomly changing the brightness and contrast during data augmentation likely improve performance?

1. Yes
2. No

Implementing image augmentation

```
def make_input_fn(csv_of_filenames, batch_size, mode,augment):
    def _input_fn():
        def decode_csv(csv_row):
            filename, label = tf.decode_csv(
                csv_row, record_defaults = [[],[]])
            image_bytes = tf.read_file(filename)
            return image_bytes, label

        dataset = tf.data.TextLineDataset(csv_of_filenames)
        .map(decode_csv).map(decode_jpeg).map(resize_image)
        if augment:
            dataset = dataset.map(augment_image)
        dataset = dataset.map(postprocess_image)
        ...
        return dataset.make_one_shot_iterator().get_next()
    return _input_fn
```

TensorFlow has a robust set of image functions

- ▶ tf.gfile
- ▶ tf.graph_util
- ▼ tf.image
 - Overview
 - adjust_brightness
 - adjust_contrast
 - adjust_gamma
 - adjust_hue
 - adjust_jpeg_quality
 - adjust_saturation
 - central_crop
 - convert_image_dtype
 - crop_and_resize
 - crop_to_bounding_box
 - decode_and_crop_jpeg
 - decode_bmp
 - decode_gif
 - decode_image
 - decode_jpeg
 - decode_png

Images

☆☆☆☆☆

Contents ▾

- Encoding and Decoding
- Resizing
- Cropping
- Flipping, Rotating and Transposing
- ...

★ **Note:** Functions taking `Tensor` arguments can also take anything accepted by [tf.convert_to_tensor](#).

Encoding and Decoding

↑

TensorFlow provides Ops to decode and encode JPEG and PNG formats. Encoded images are represented by scalar string Tensors, decoded images by 3-D uint8 tensors of

Implementing image augmentation

```
def make_input_fn(csv_of_filenames, batch_size, mode,augment):
    def _input_fn():
        def decode_csv(csv_row):
            filename, label = tf.decode_csv(
                csv_row, record_defaults = [[],[]])
            image_bytes = tf.read_file(filename)
            return image_bytes, label

        dataset = tf.data.TextLineDataset(csv_of_filenames)
        .map(decode_csv).map(decode_jpeg).map(resize_image)
        if augment:
            dataset = dataset.map(augment_image)
        dataset = dataset.map(postprocess_image)
        ...
        return dataset.make_one_shot_iterator().get_next()
    return _input_fn
```

Implementing image augmentation

JPEG URL	Label
gs://path/to/image_1.jpg	daisy
gs://path/to/image_2.jpg	dandelion
gs://path/to/image_3.jpg	dandelion

Implementing decode_csv()

```
def make_input_fn(csv_of_filenames, batch_size, mode,augment):
    def _input_fn():
        def decode_csv(csv_row):
            filename, label = tf.decode_csv(
                csv_row, record_defaults = [[],[]])
            image_bytes = tf.read_file(filename)
            return image_bytes, label

        dataset = tf.data.TextLineDataset(csv_of_filenames)
        .map(decode_csv).map(decode_jpeg).map(resize_image)
        if augment:
            dataset = dataset.map(augment_image)
        dataset = dataset.map(postprocess_image)
        ...
        return dataset.make_one_shot_iterator().get_next()
    return _input_fn
```

Implementing decode_jpeg()

```
def make_input_fn(csv_of_filenames, batch_size, mode,augment):
    def _input_fn():
        def decode_csv(csv_row):
            filename, label = tf.decode_csv(
                csv_row, record_defaults = [[],[]])
            image_bytes = tf.read_file(filename)
            return image_bytes, label

        dataset = tf.data.TextLineDataset(csv_of_filenames)
        .map(decode_csv).map(decode_jpeg).map(resize_image)
        if augment:
            dataset = dataset.map(augment_image)
        dataset = dataset.map(postprocess_image)
        ...
        return dataset.make_one_shot_iterator().get_next()
    return _input_fn
```

Implementing decode_jpeg()

```
import tensorflow.image as tfl
def decode_jpeg(image_bytes, label):
    image = tfl.decode_jpeg(image_bytes, channels=NUM_CHANNELS)
    image = tfl.convert_image_dtype(image, dtype=tf.float32) # 0-1
    return image, label
```

Implementing resize_image()

```
def make_input_fn(csv_of_filenames, batch_size, mode,augment):
    def _input_fn():
        def decode_csv(csv_row):
            filename, label = tf.decode_csv(
                csv_row, record_defaults = [[],[]])
            image_bytes = tf.read_file(filename)
            return image_bytes, label

        dataset = tf.data.TextLineDataset(csv_of_filenames)
        .map(decode_csv).map(decode_jpeg).map(resize_image)
        if augment:
            dataset = dataset.map(augment_image)
        dataset = dataset.map(postprocess_image)
        ...
        return dataset.make_one_shot_iterator().get_next()
    return _input_fn
```

Implementing resize_image()

```
import tensorflow.image as tfi
def resize_image(image, label):
    image = tf.expand_dims(image, 0) # add batch dimension
    image = tfi.resize_bilinear(image, [HEIGHT, WIDTH],
                                align_corners=False)
    image = tf.squeeze(image) # throw away batch dimension
    return image, label
```

Implementing augment_image()

```
def make_input_fn(csv_of_filenames, batch_size, mode,augment):
    def _input_fn():
        def decode_csv(csv_row):
            filename, label = tf.decode_csv(
                csv_row, record_defaults = [[],[]])
            image_bytes = tf.read_file(filename)
            return image_bytes, label

        dataset = tf.data.TextLineDataset(csv_of_filenames)
        .map(decode_csv).map(decode_jpeg).map(resize_image)
        if augment:
            dataset = dataset.map(augment_image)
        dataset = dataset.map(postprocess_image)
        ...
        return dataset.make_one_shot_iterator().get_next()
    return _input_fn
```

Implementing augment_image()

```
def augment_image(image_dict, label=None):
    image = image_dict['image']
    image = tf.expand_dims(image, 0) # resize_bilinear needs batches
    image = tfl.resize_bilinear(
        image, [HEIGHT+10, WIDTH+10], align_corners=False)
    image = tf.squeeze(image) #remove batch dimension
    image = tf.random_crop(image, [HEIGHT, WIDTH, NUM_CHANNELS])
    image = tfl.random_flip_left_right(image)
    image = tfl.random_brightness(image, max_delta=63.0/255.0)
    image = tfl.random_contrast(image, lower=0.2, upper=1.8)

    return image, label
```

Implementing augment_image()

```
def augment_image(image_dict, label=None):
    image = image_dict['image']
    image = tf.expand_dims(image, 0) # resize_bilinear needs batches
    image = tfl.resize_bilinear(
        image, [HEIGHT+10, WIDTH+10], align_corners=False)
    image = tf.squeeze(image) #remove batch dimension
    image = tf.random_crop(image, [HEIGHT, WIDTH,
NUM_CHANNELS])
    image = tfl.random_flip_left_right(image)
    image = tfl.random_brightness(image, max_delta=63.0/255.0)
    image = tfl.random_contrast(image, lower=0.2, upper=1.8)

    return image, label
```

Implementing postprocess_image()

```
def make_input_fn(csv_of_filenames, batch_size, mode,augment):
    def _input_fn():
        def decode_csv(csv_row):
            filename, label = tf.decode_csv(
                csv_row, record_defaults = [[],[]])
            image_bytes = tf.read_file(filename)
            return image_bytes, label

        dataset = tf.data.TextLineDataset(csv_of_filenames)
        .map(decode_csv).map(decode_jpeg).map(resize_image)
        if augment:
            dataset = dataset.map(augment_image)
        dataset = dataset.map(postprocess_image)
        ...
        return dataset.make_one_shot_iterator().get_next()
    return _input_fn
```

Implementing postprocess_image()

```
def postprocess_image(image, label):
    #pixel values are in range [0,1], convert to [-1,1]
    image = tf.subtract(image, 0.5)
    image = tf.multiply(image, 2.0)
    return {'image' : image}, label
```

File Name:
T-ICML-O_3_I3_lab_intro:_implementing_image_augmentation

Format: Presenter in Studio

Presenter: Max Lotstein

Lab

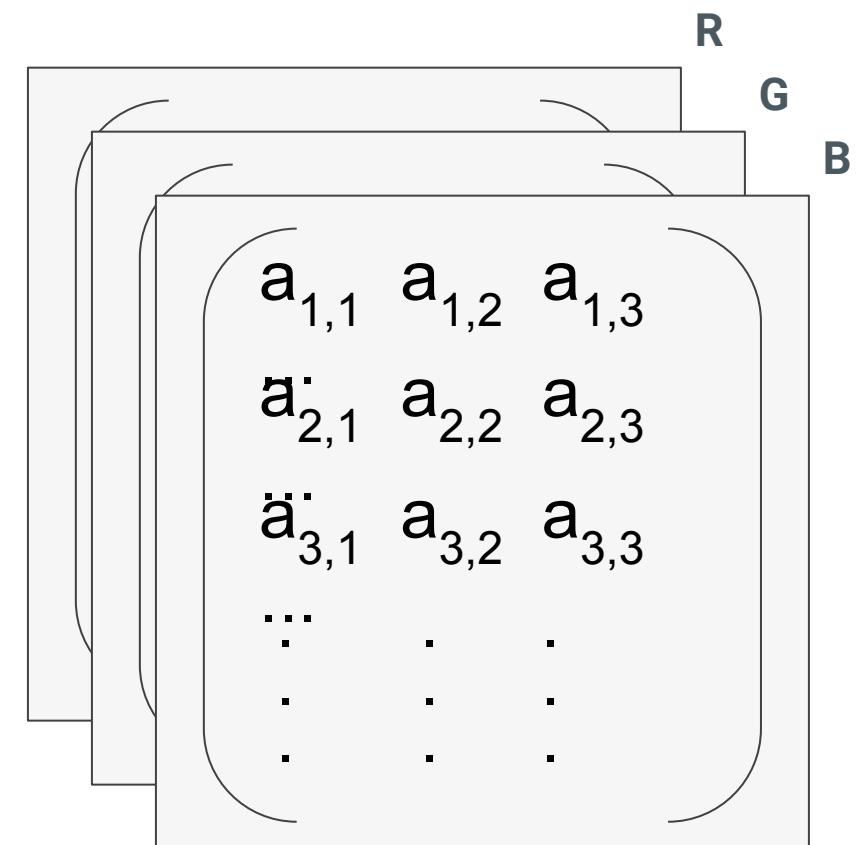
Implementing image
augmentation

Max Lotstein

Working with color images



Color images are still just tensors



Color images are still just tensors

JPEG URL	Label
gs://path/to/image_1.jpg	daisy
gs://path/to/image_2.jpg	dandelion
gs://path/to/image_3.jpg	dandelion

Extra Credit: implement blur using convolution



?	?	?
?	?	?
?	?	?



File Name: T-ICML-O_3_l9_image_augmentation_lab_intro

Format: Presenter in Studio

Presenter: Max Lotstein

File Name:
T-ICML-O_3_I4_lab_solution:_implementing_image_augmentation

Format: Presenter in Studio

Presenter: Max Lotstein



Title Safe >

< Action Safe

File Name: T-ICML-O_3_I5_transfer_learning

Format: Presenter in Studio

Presenter: Max Lotstein

Agenda

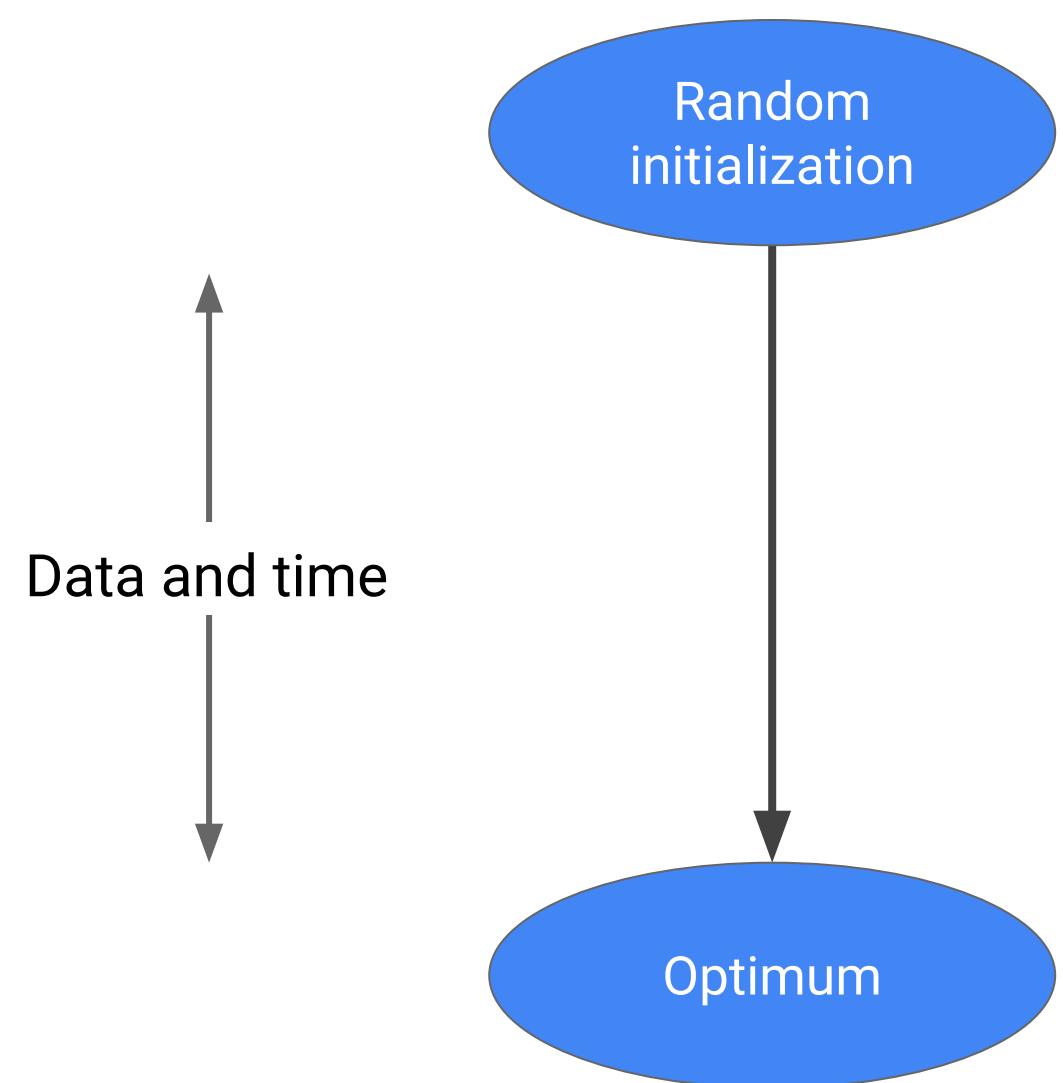
The Data Scarcity Problem

Data Augmentation

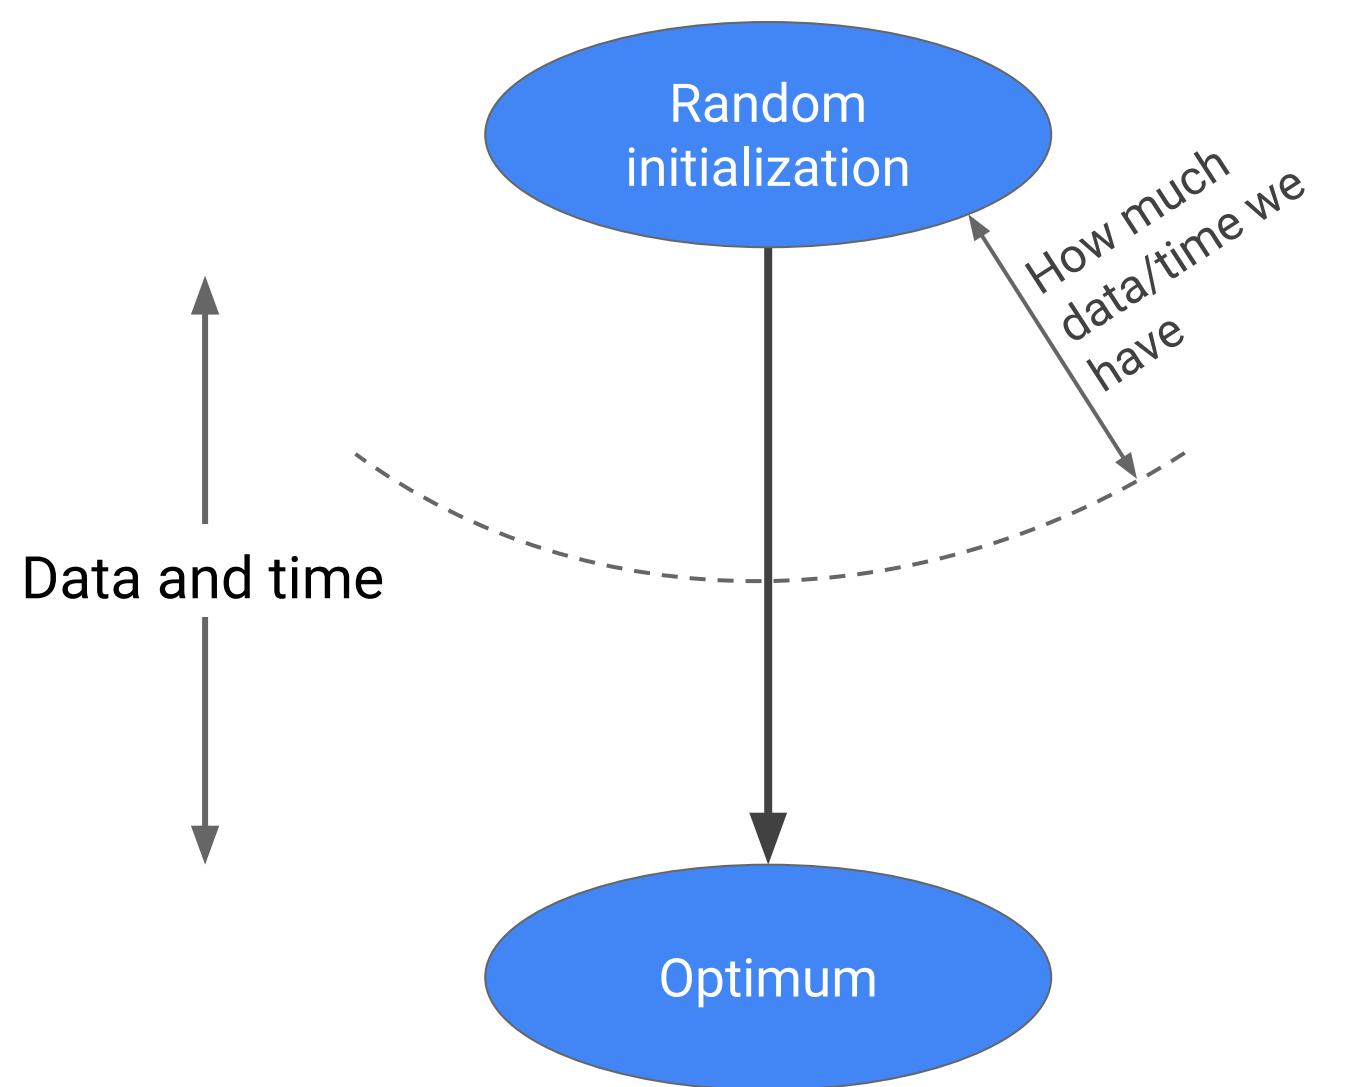
Transfer Learning

No Data, No Problem

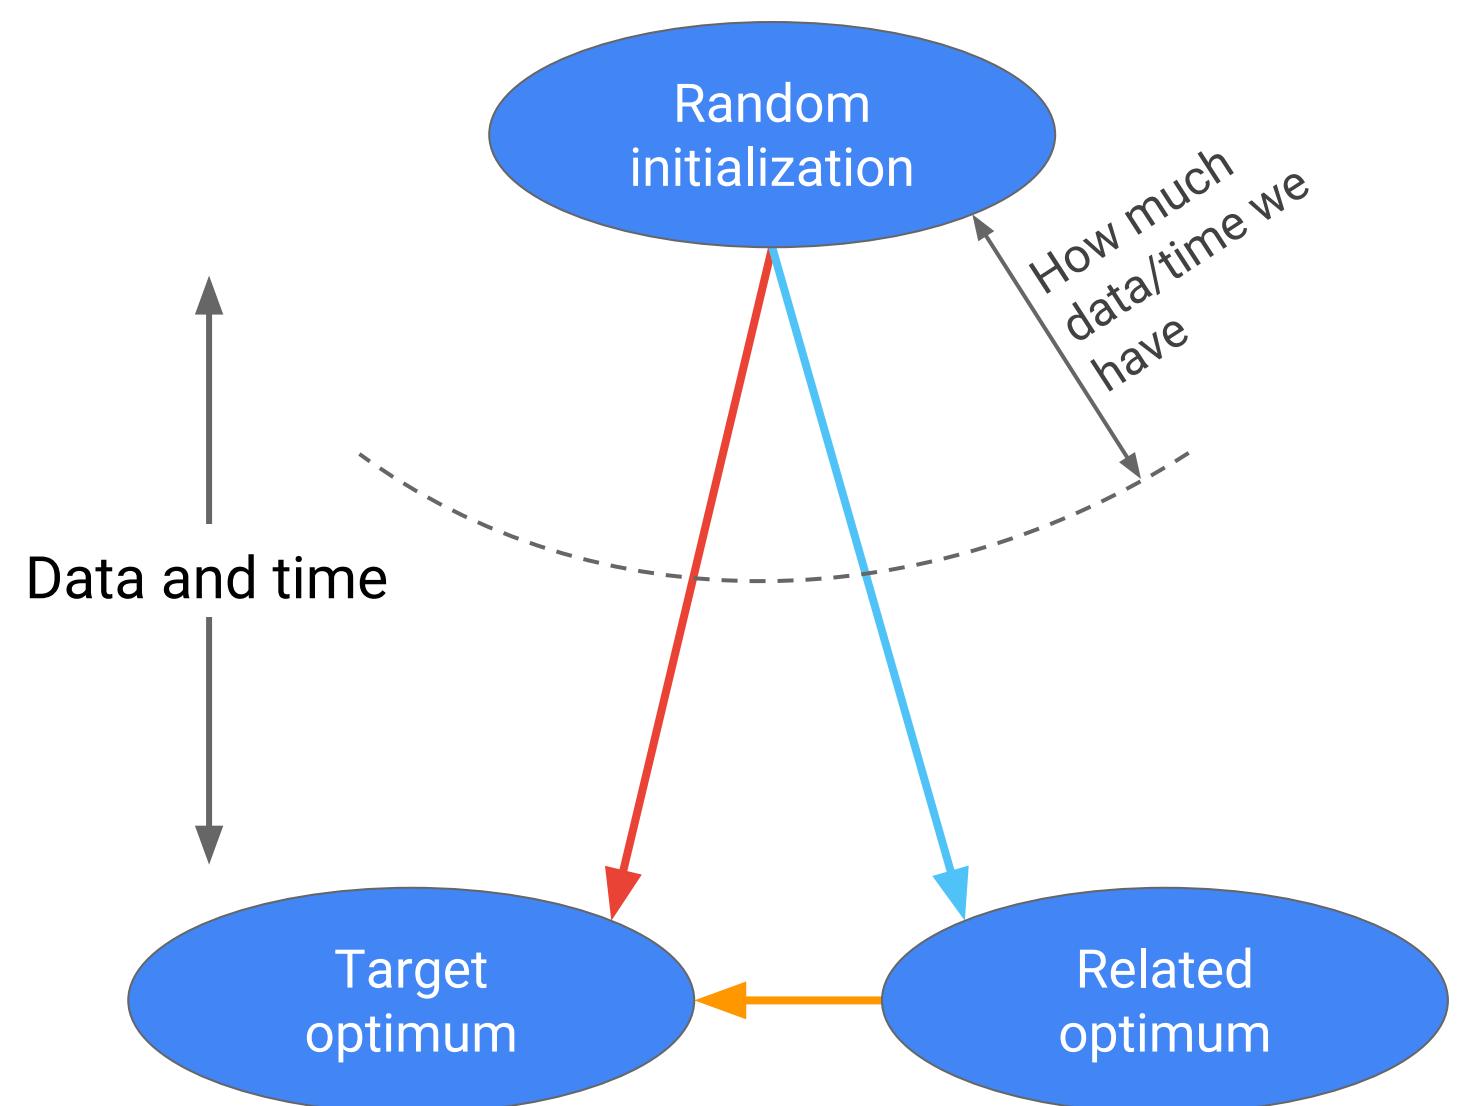
Optimization is a journey through weight space



Unstructured domains limit our optimization range



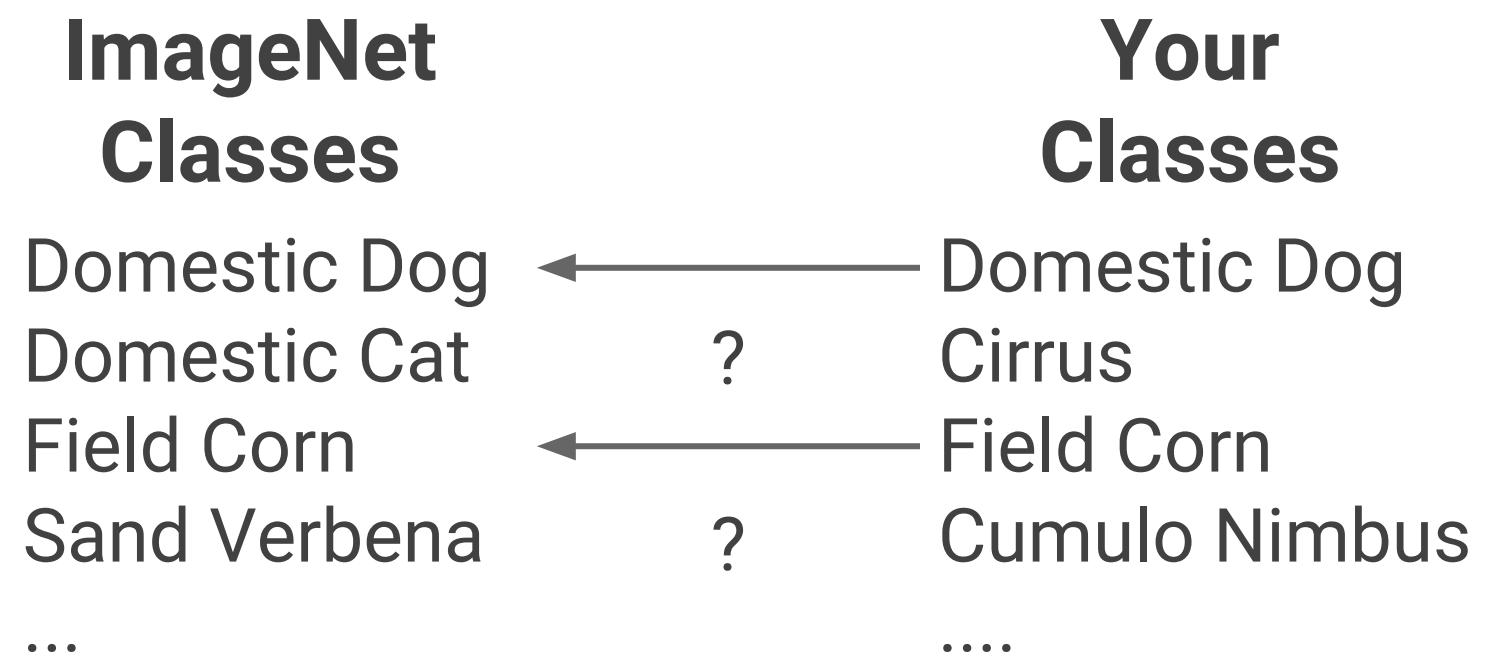
Transfer learning is like a shortcut

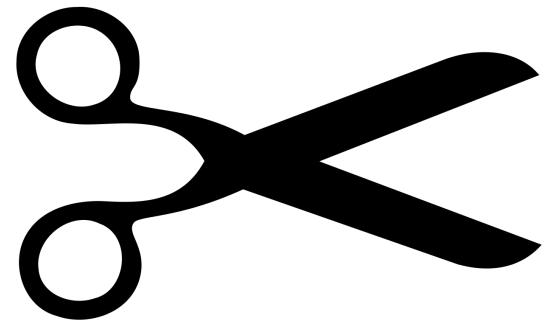


ImageNet

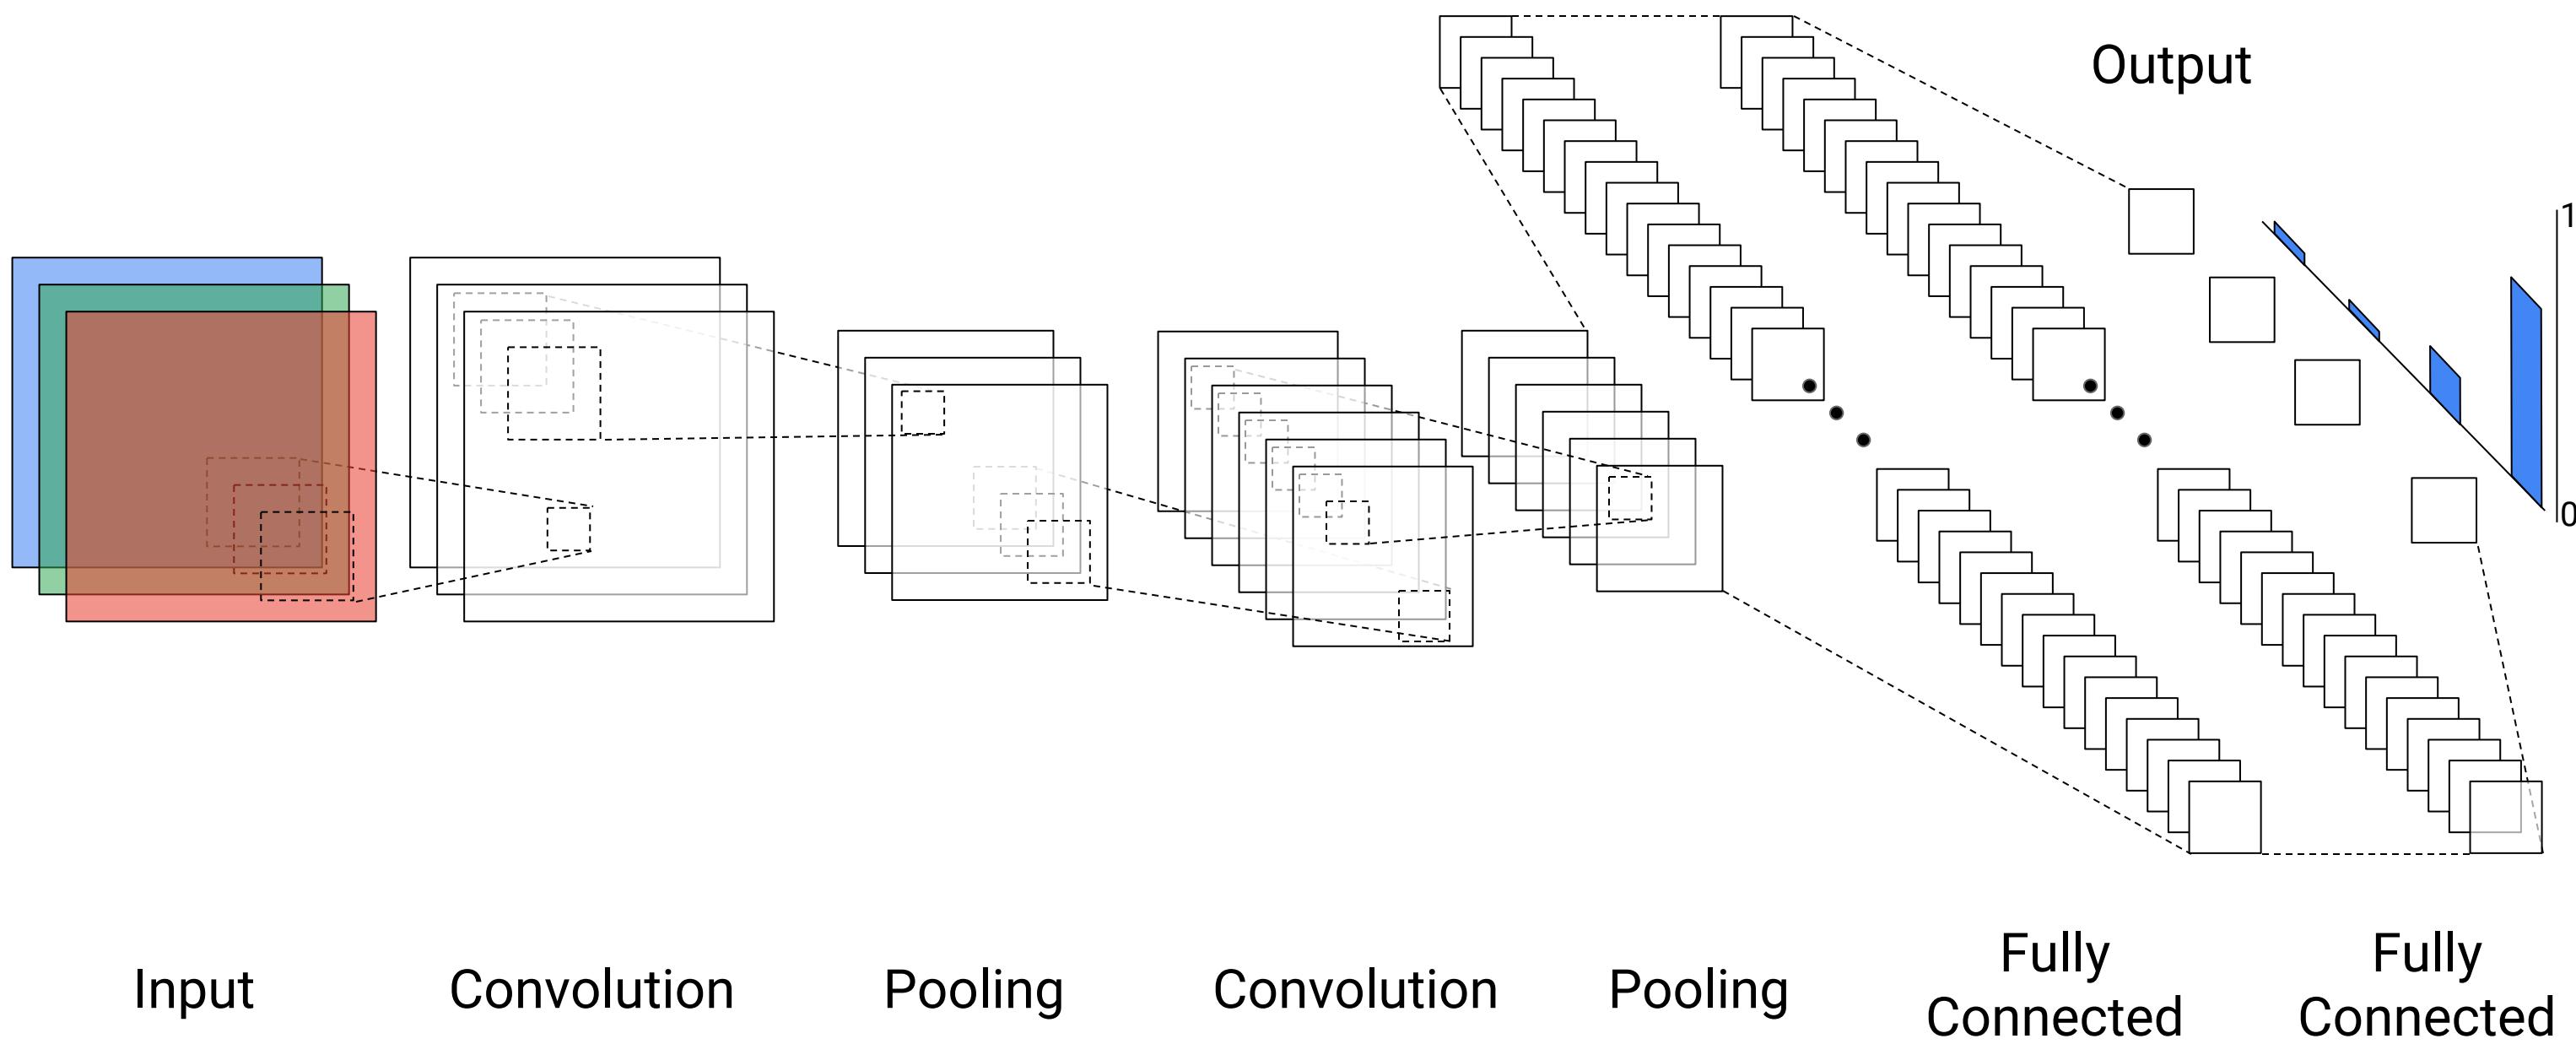


Naive Transfer Learning





Where to cut?



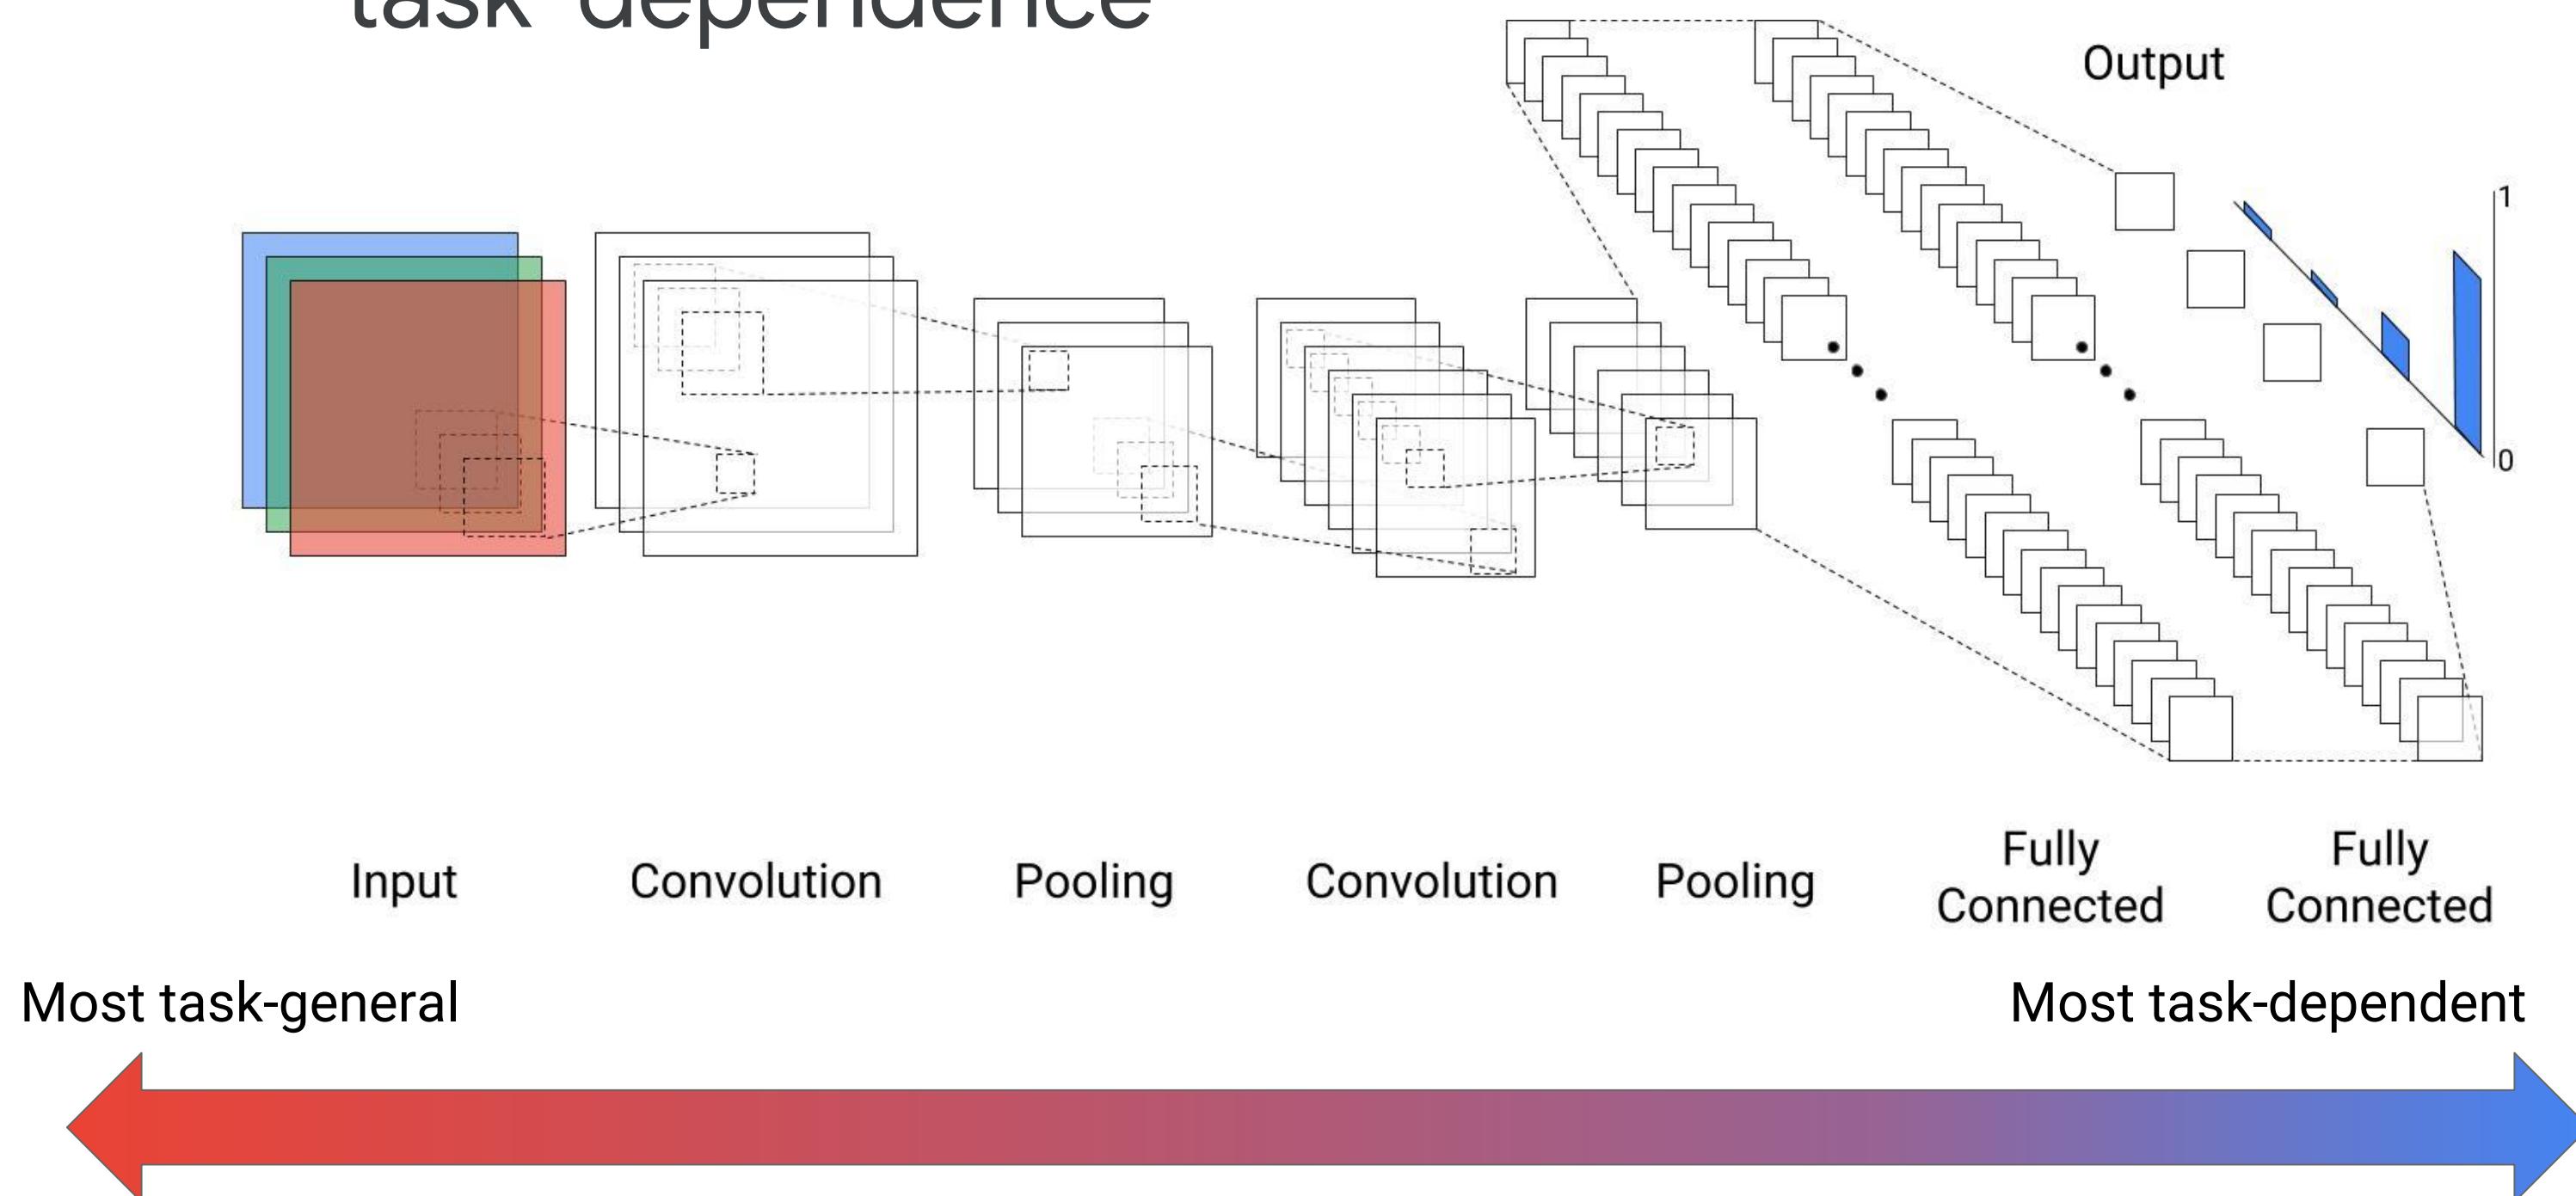
Which parts of the source image model are most closely aligned to the source task?

1. The input to the model
2. The convolutional layers
3. The layers near the output

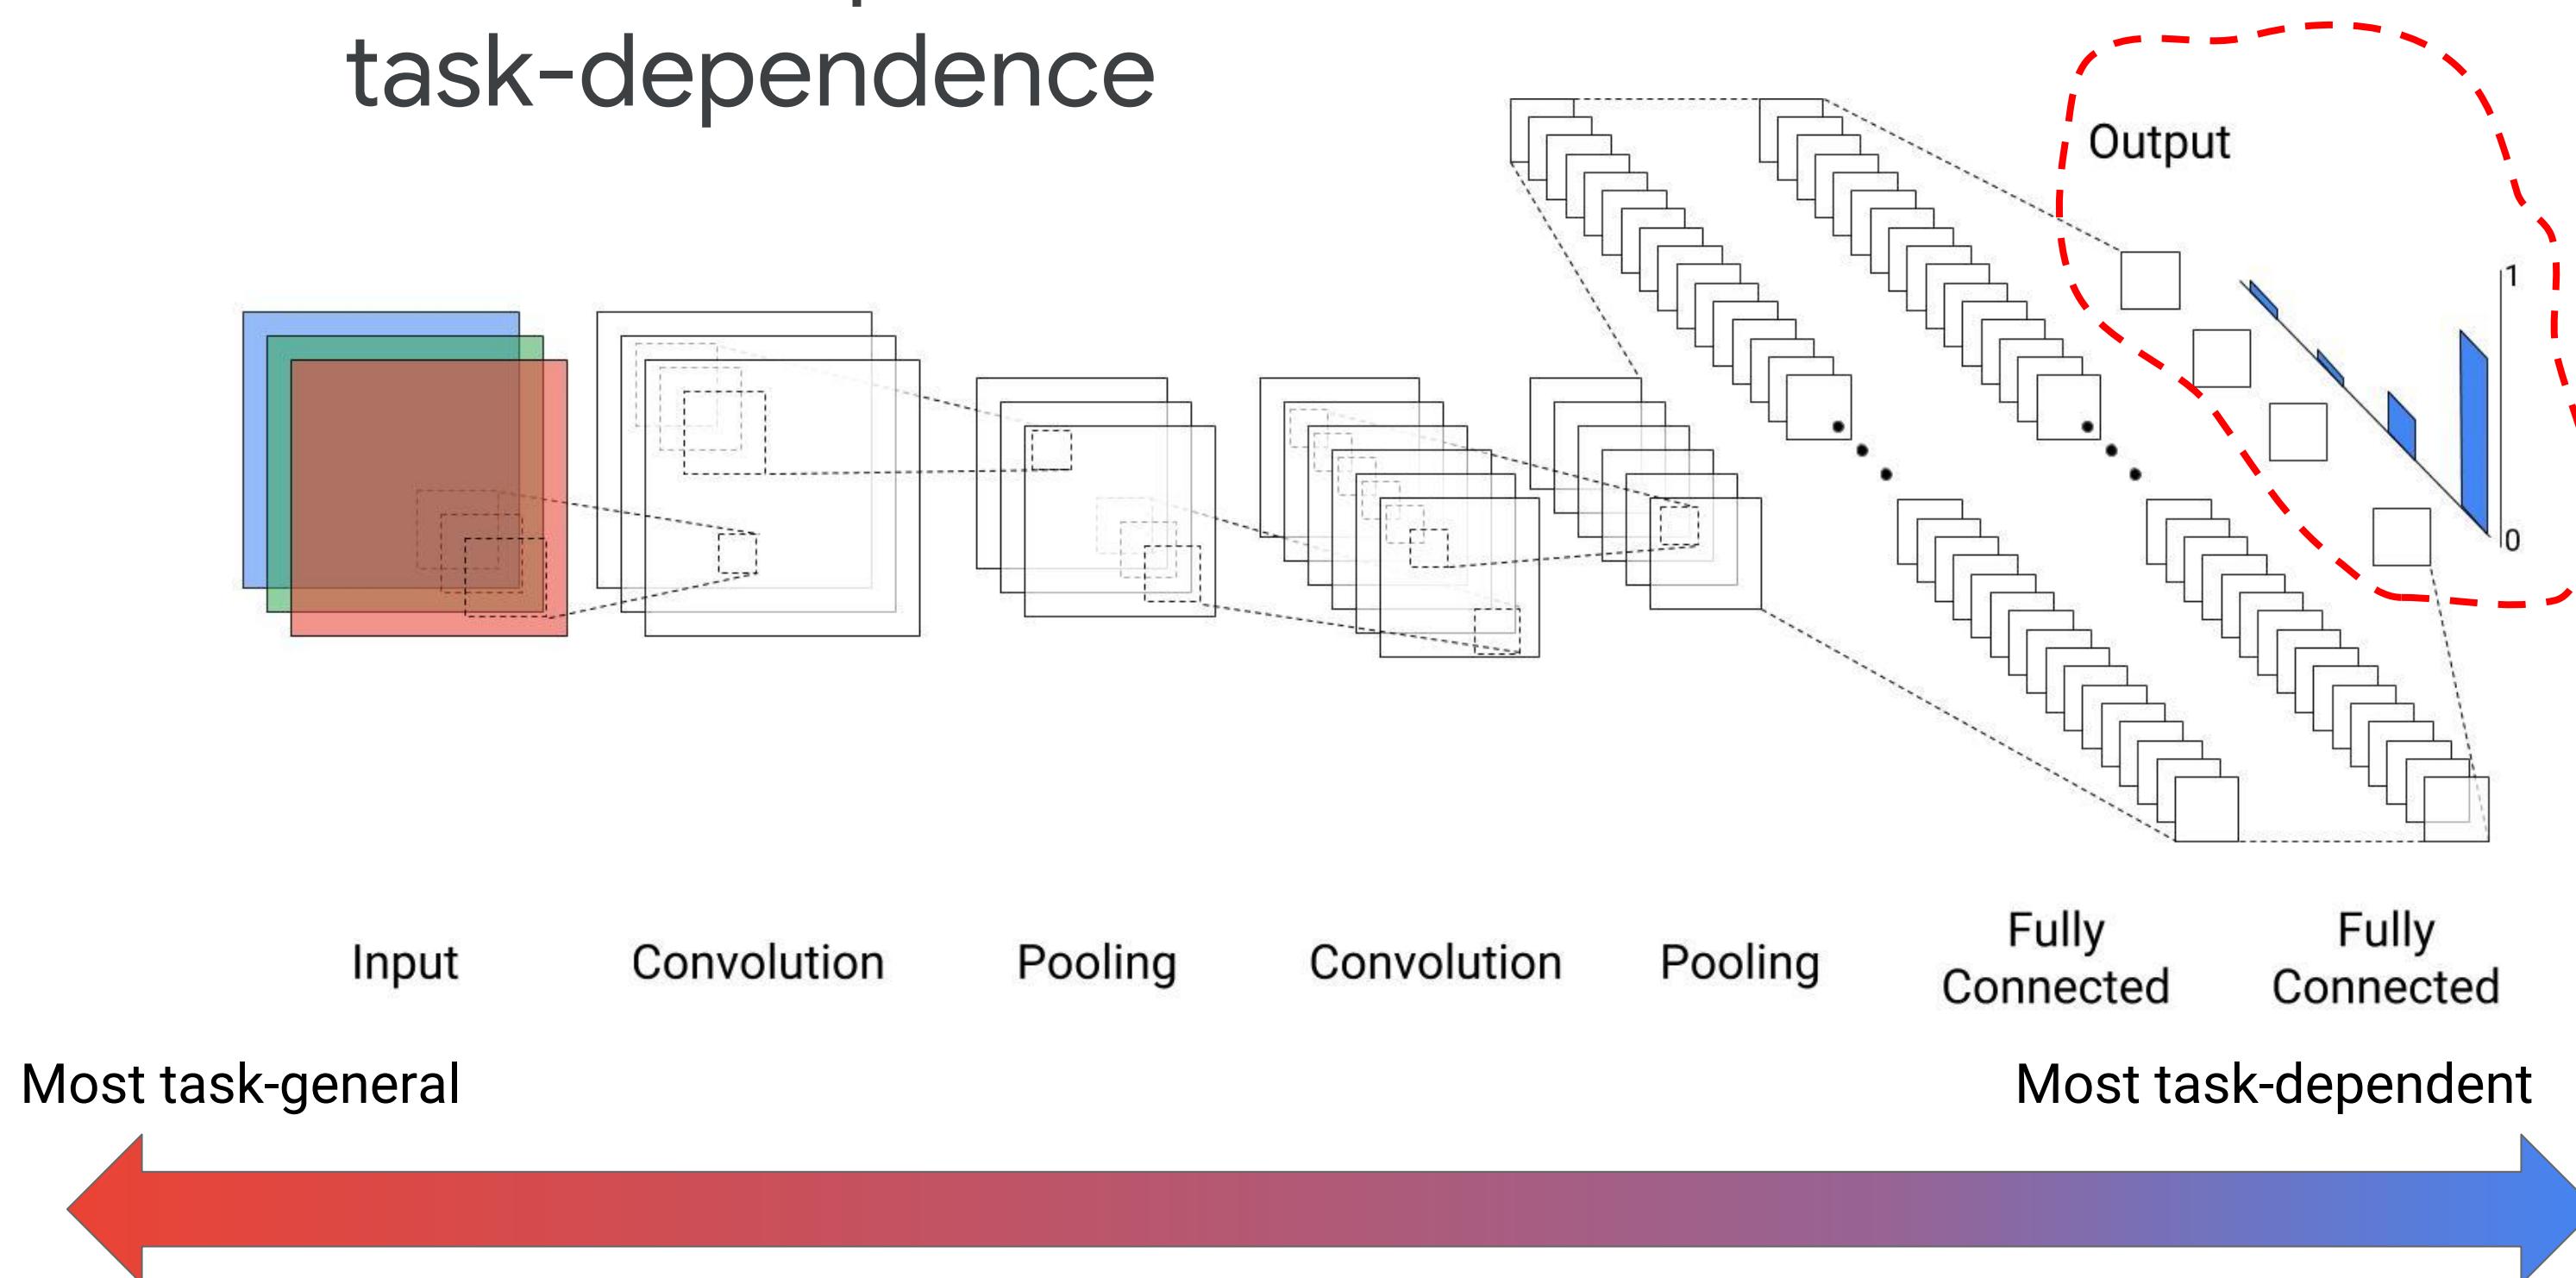
Which parts of the source image model are most closely aligned to the source task?

1. The input to the model
2. The convolutional layers
- 3. The layers near the output**

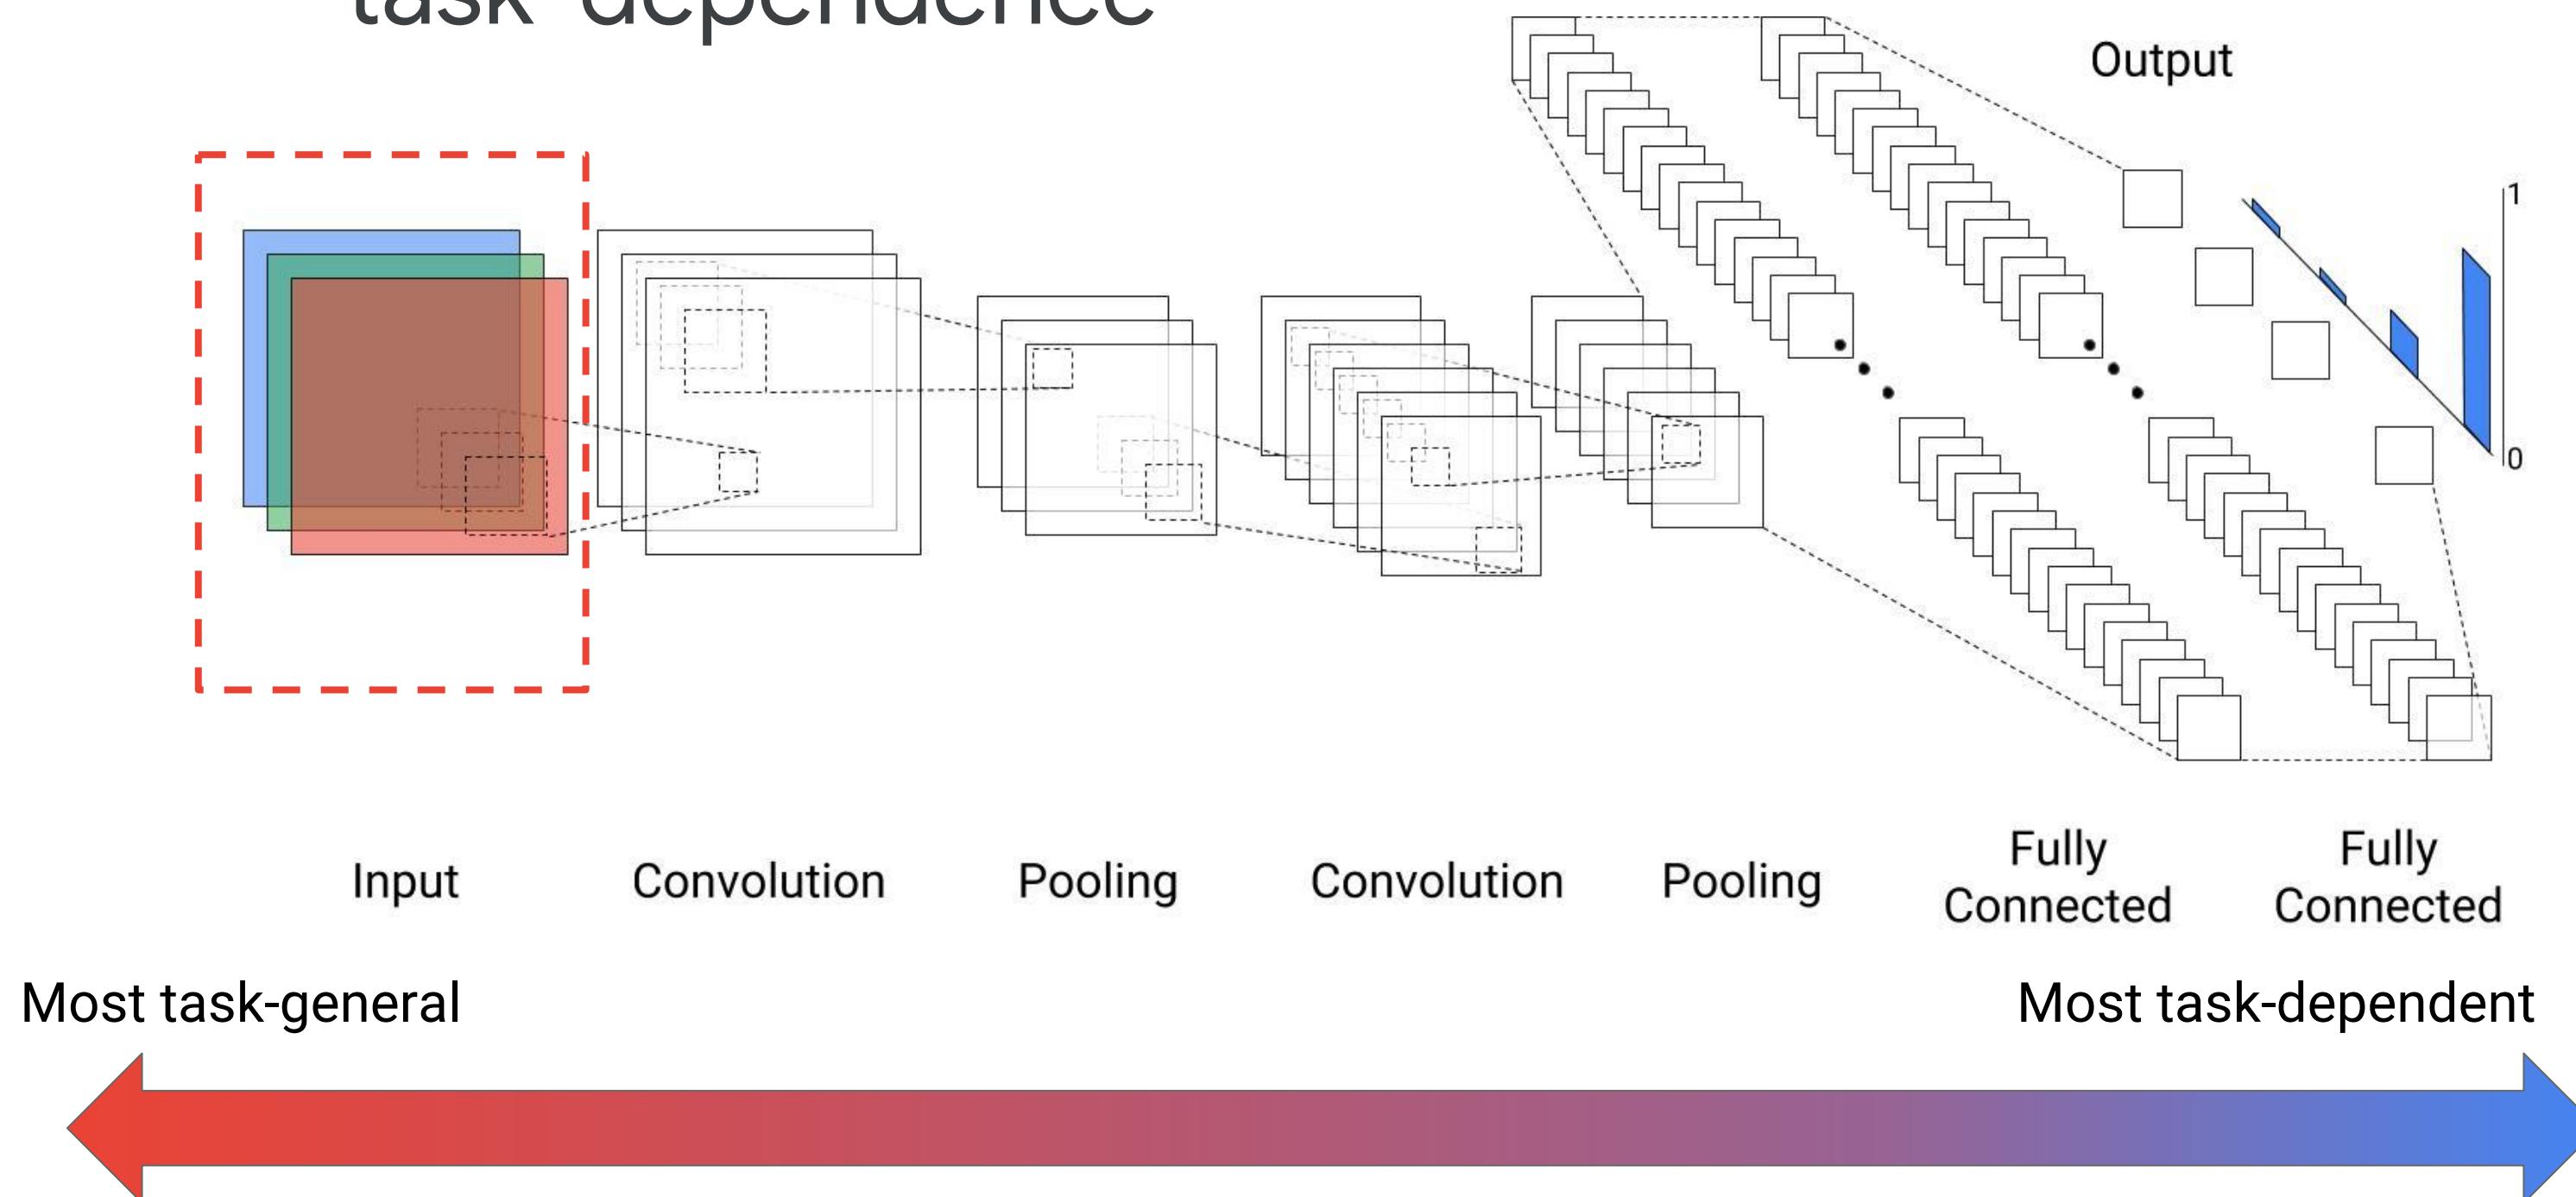
CNNs and the spectrum of task-dependence



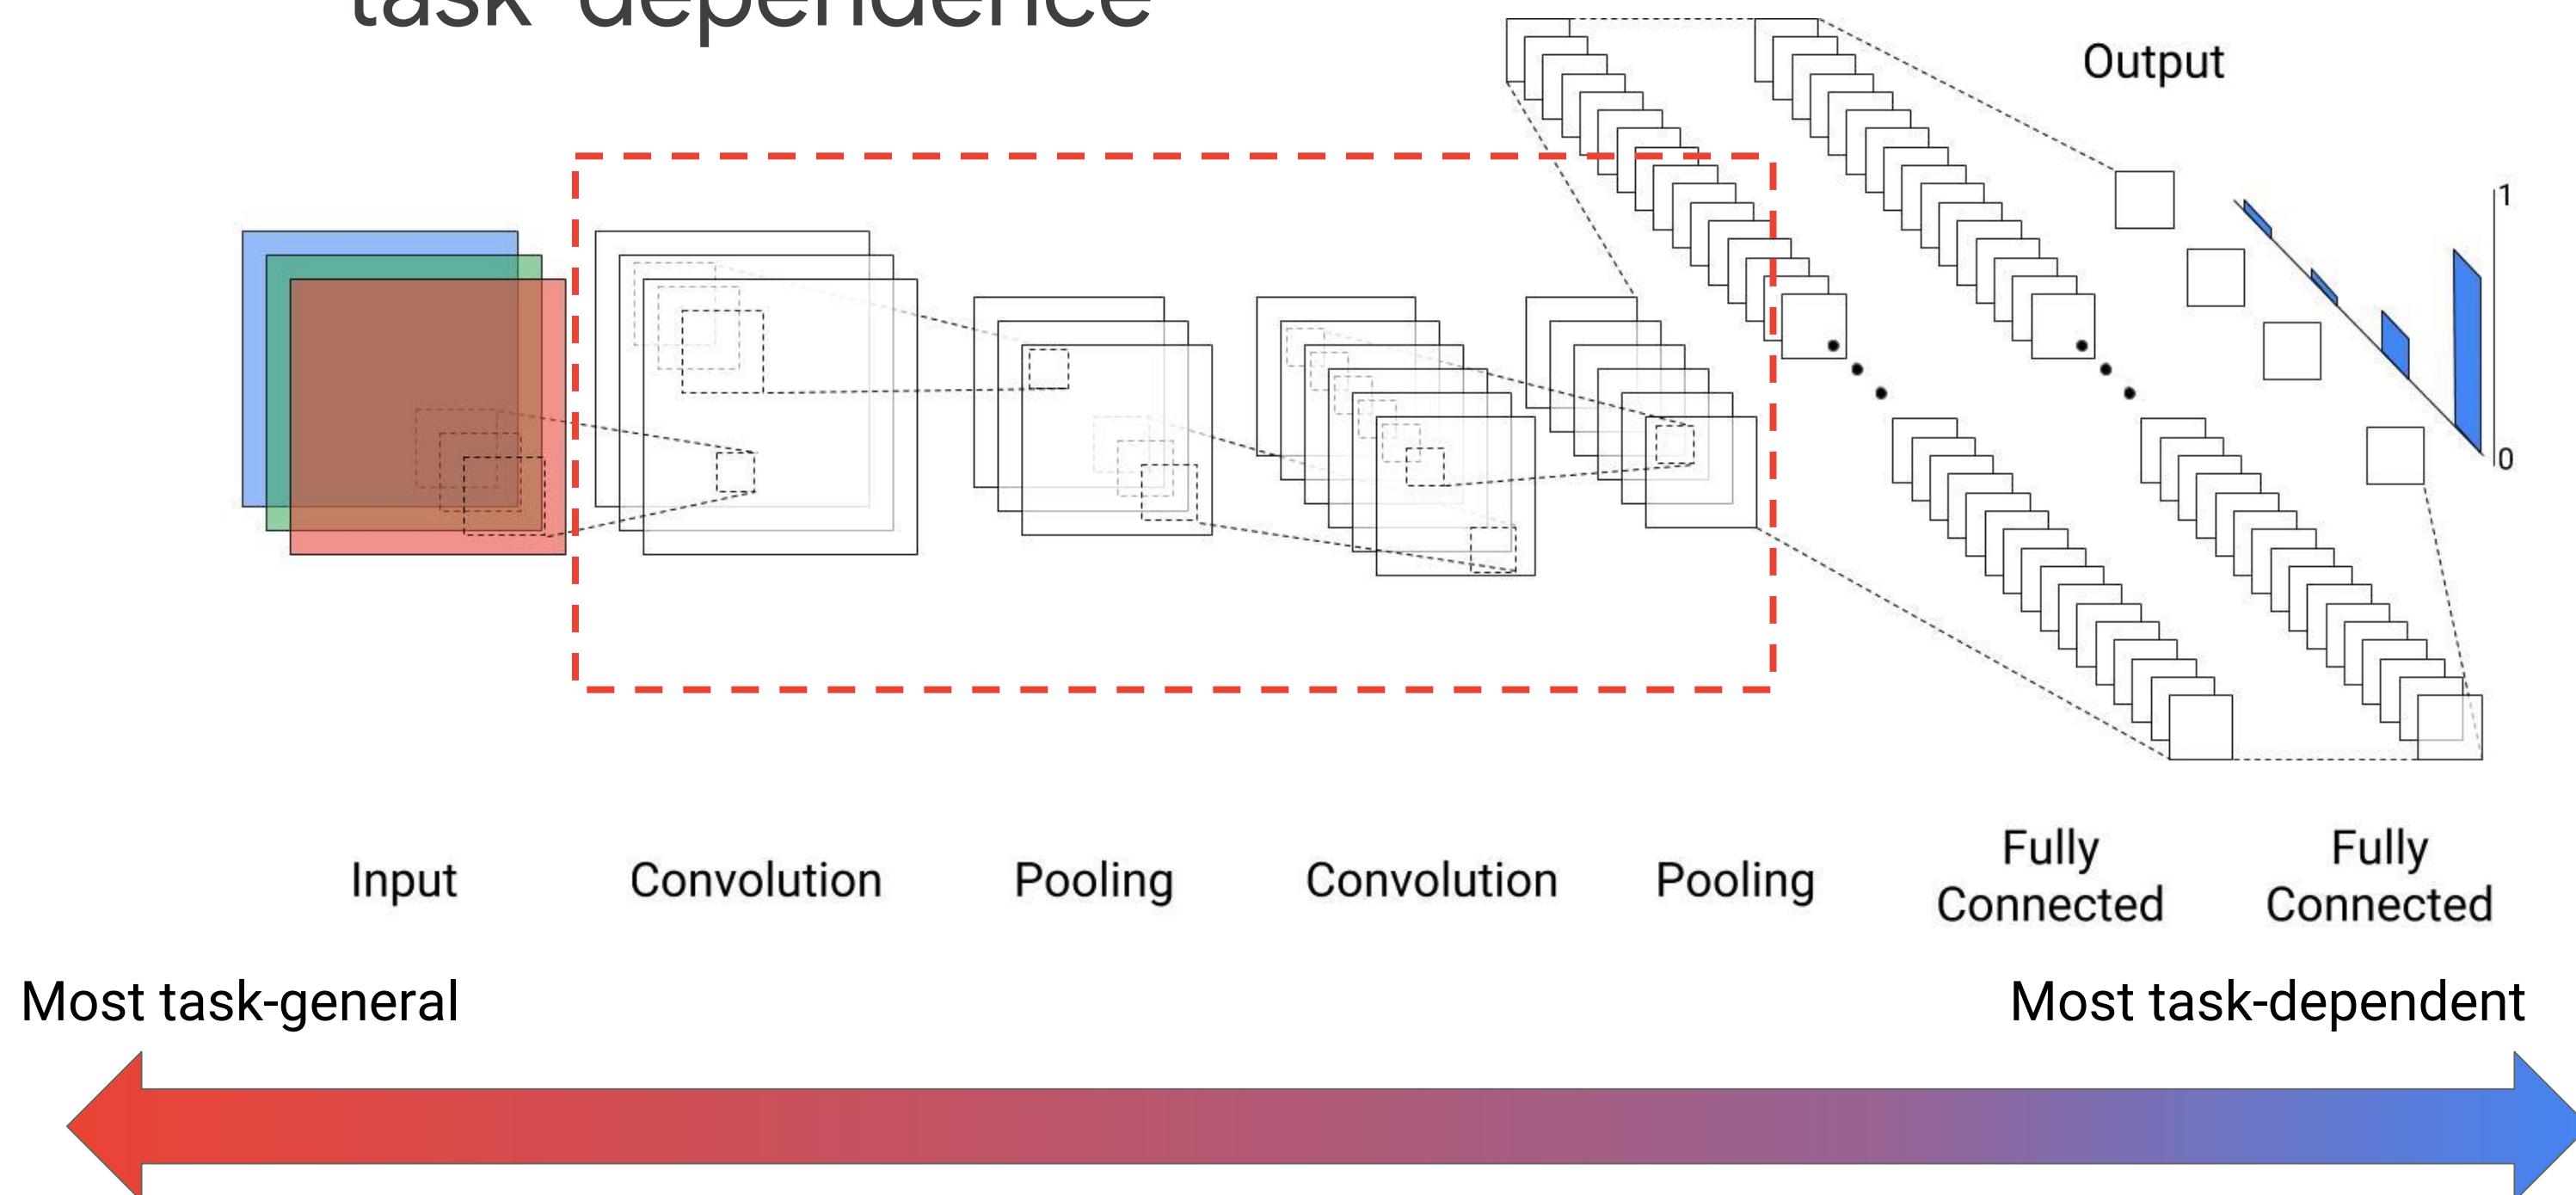
CNNs and the spectrum of task-dependence



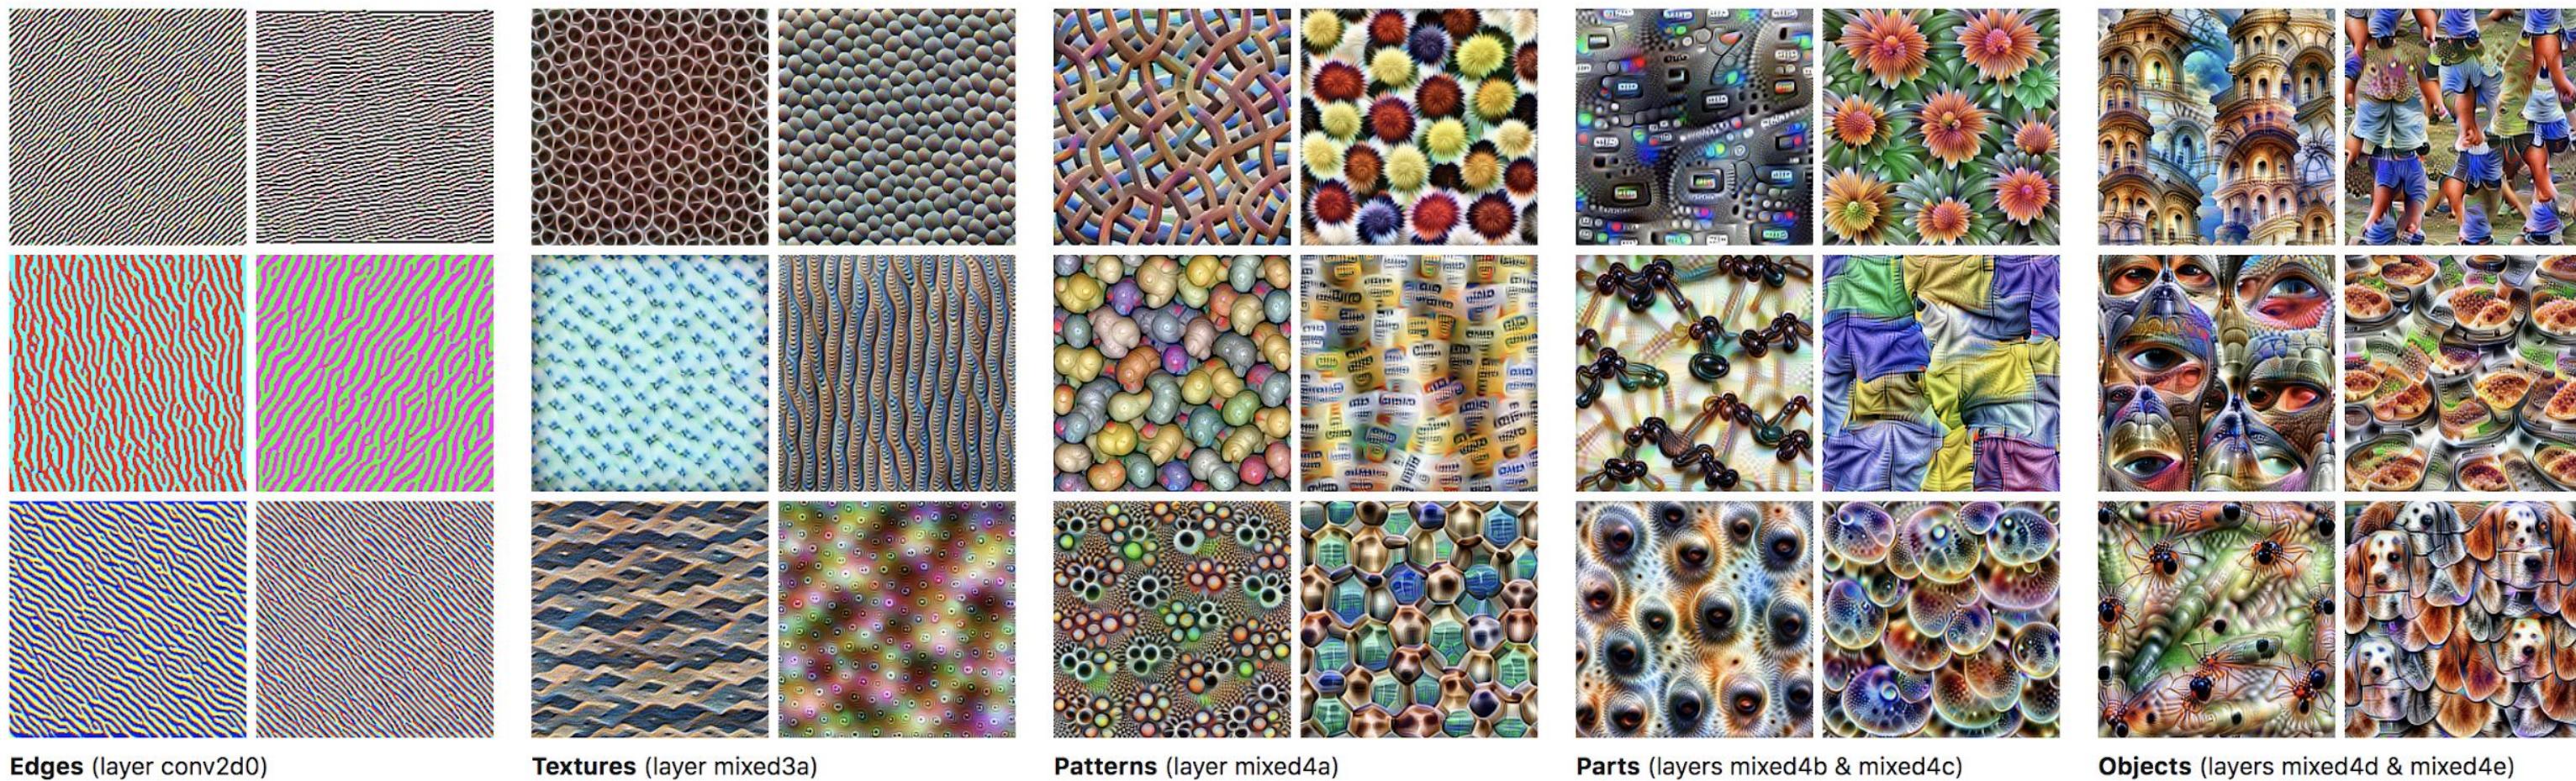
CNNs and the spectrum of task-dependence



CNNs and the spectrum of task-dependence

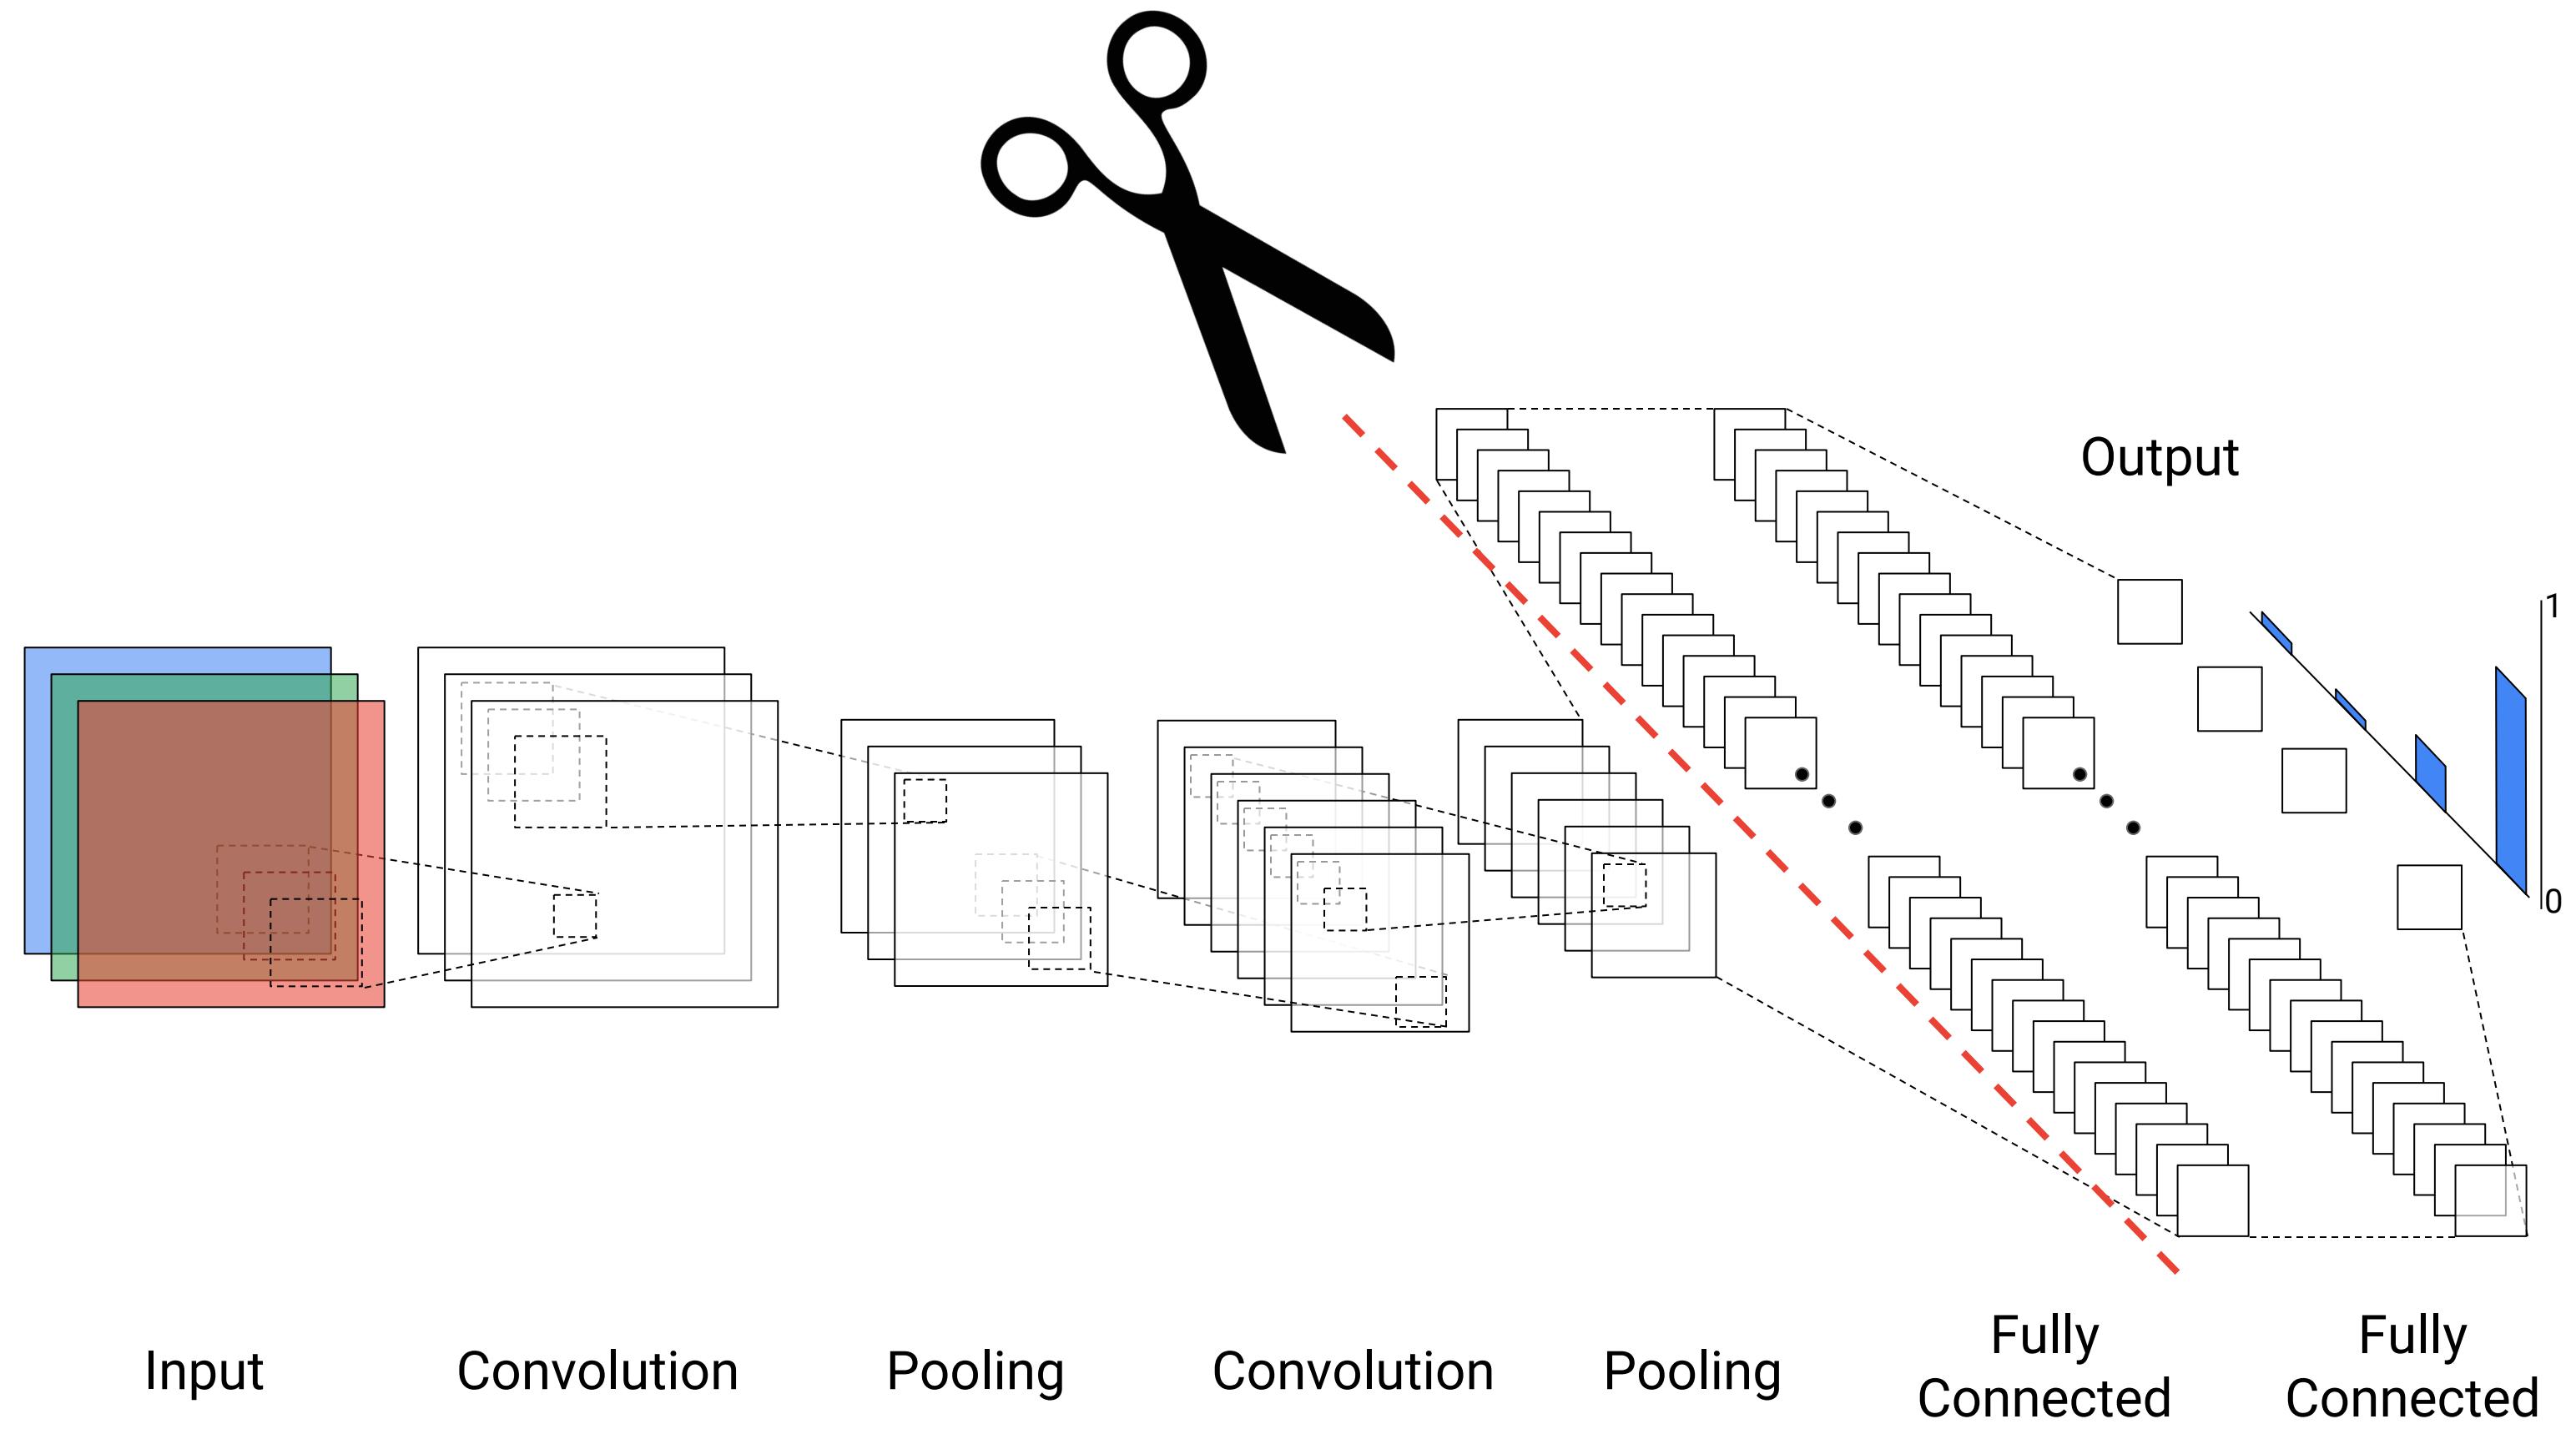


CNNs learn a hierarchy of features

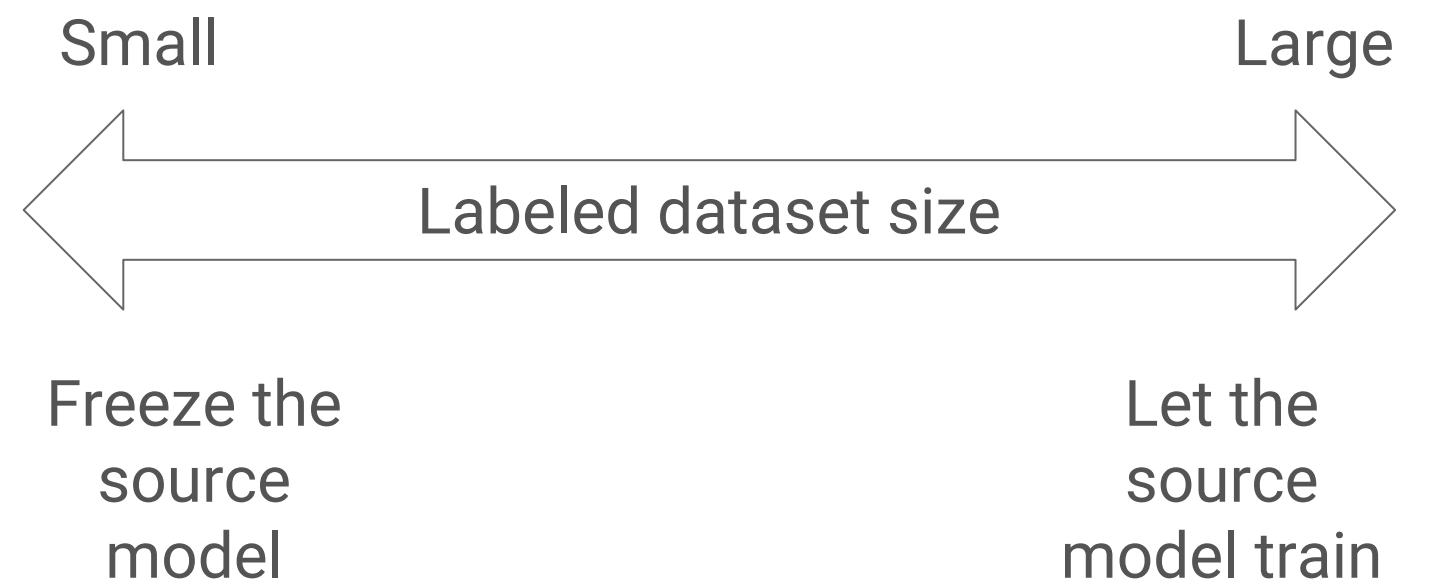


Flow of data in the network

<https://distill.pub/2017/feature-visualization/>



To freeze or not to freeze?



File Name:

T-ICML-O_3_I6_lab_intro:_implementing_transfer_learning

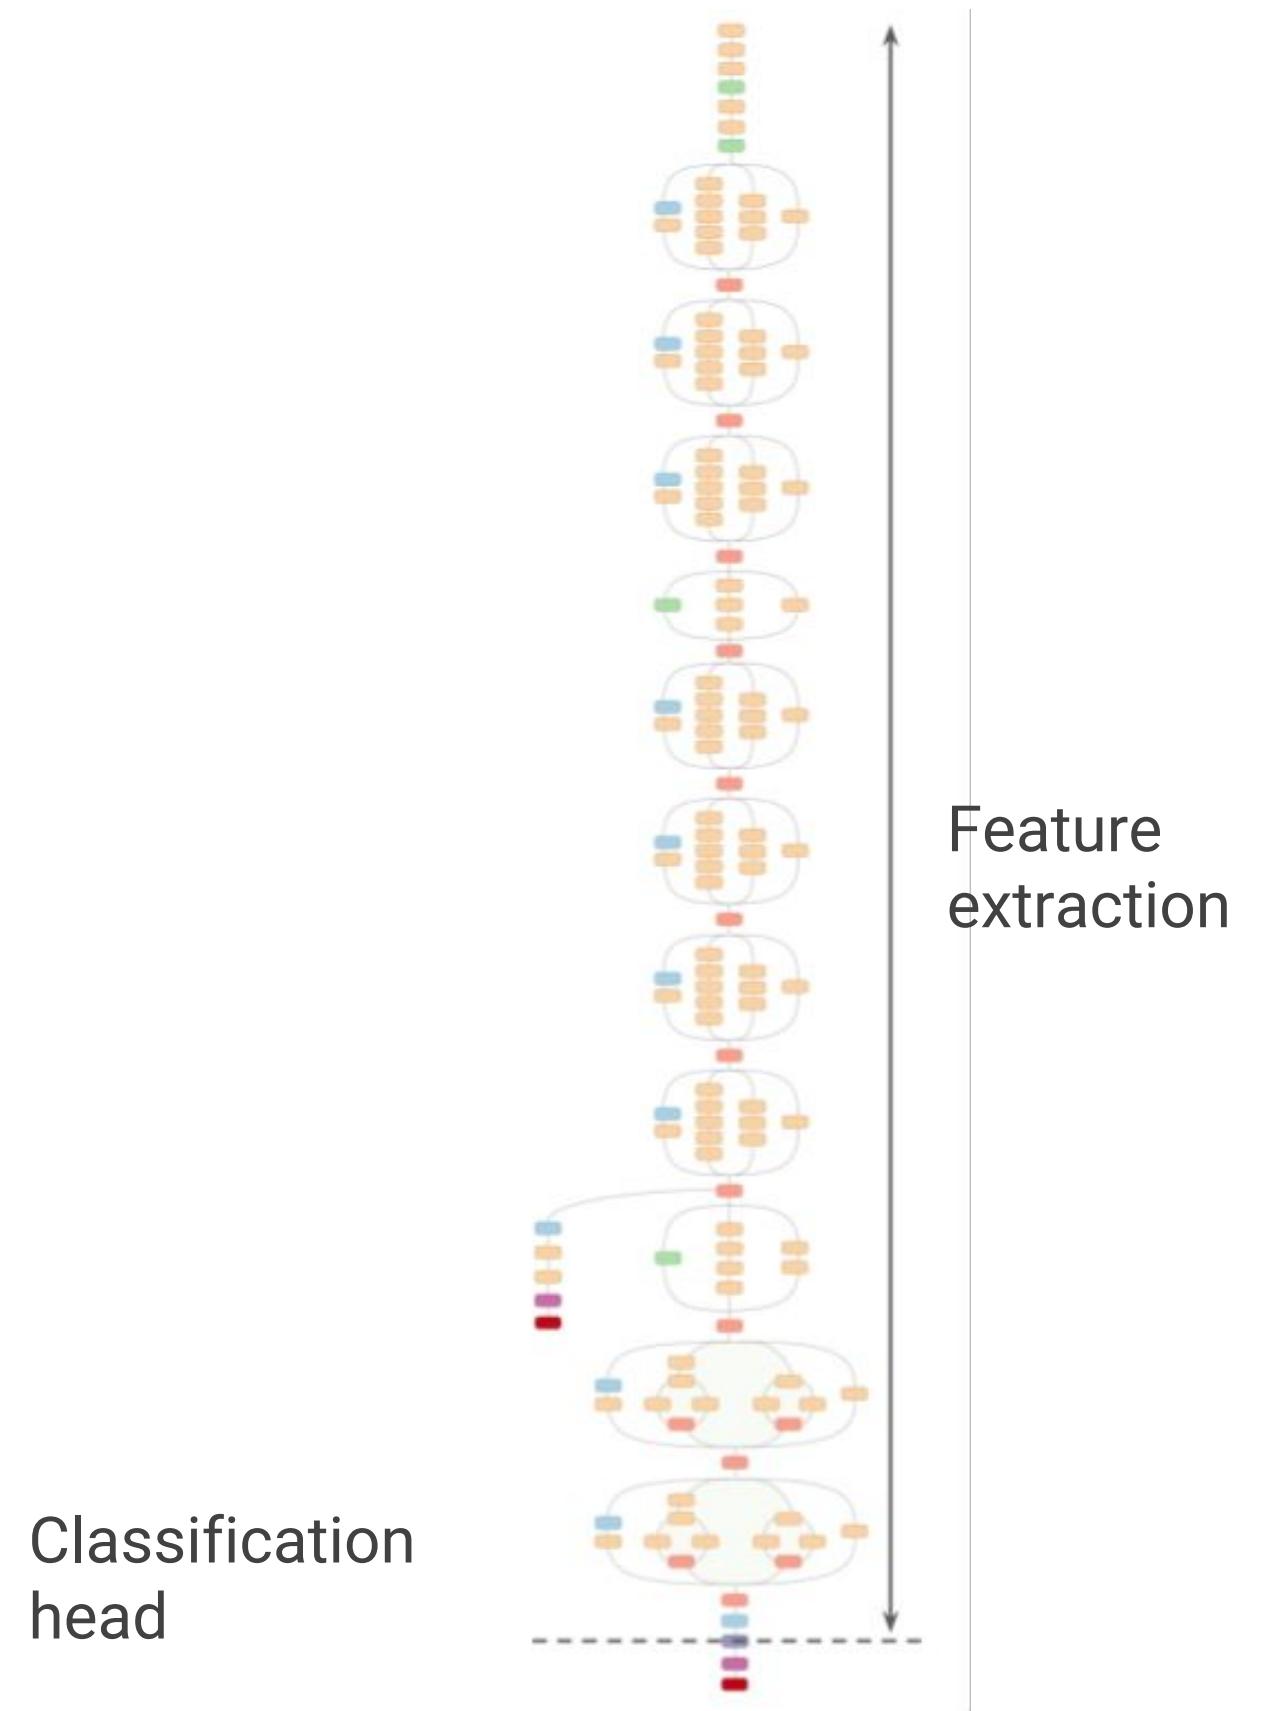
Format: Presenter in Studio

Presenter: Max Lotstein

Lab

Implementing transfer
learning

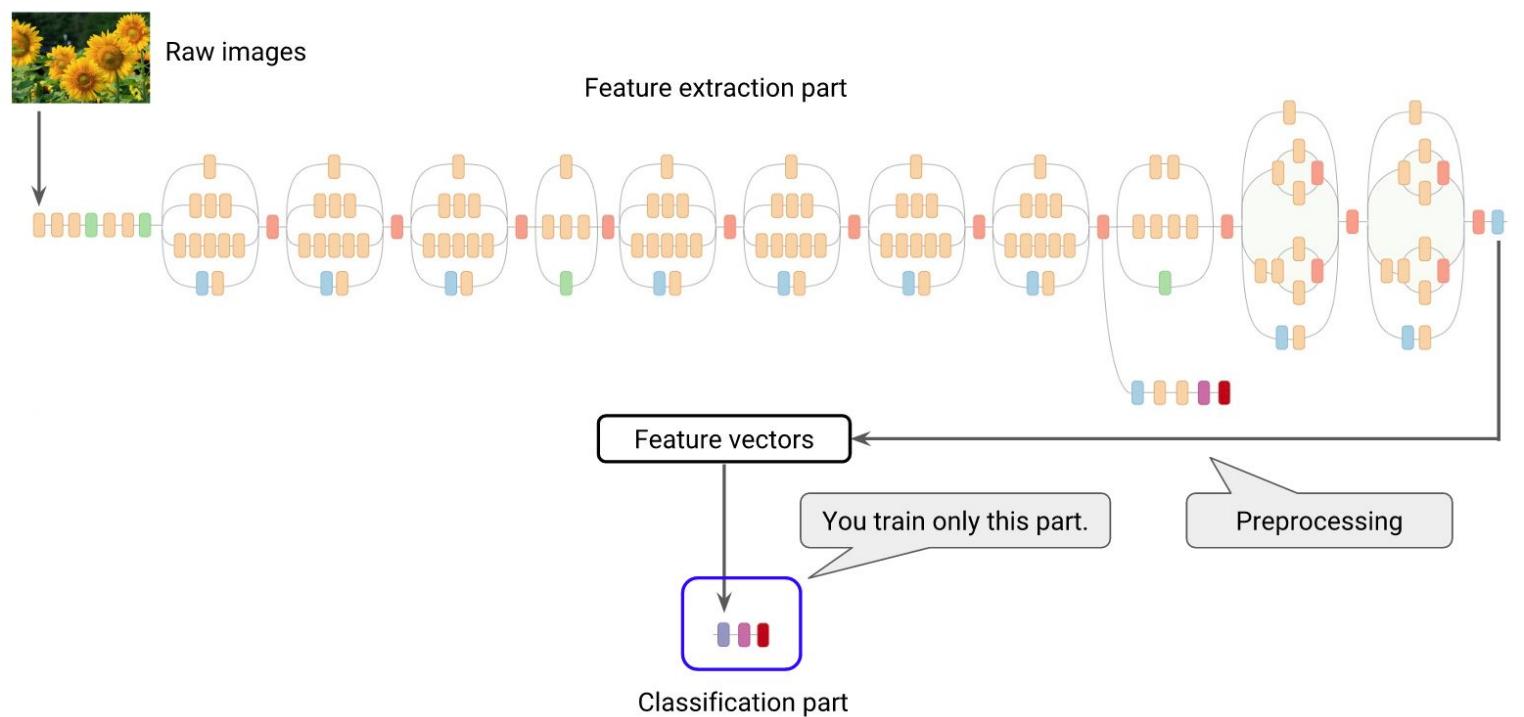
Max Lotstein



First, we'll preprocess using Dataflow

```
_ = (p
  | 'Read input' >> read_input_source
  | 'Parse input' >> beam.Map(lambda line:
    csv.reader([line]).next())
  | 'Extract label ids' >>
    beam.ParDo(ExtractLabelIdsDoFn(),
    beam.pvalue.AsIter(labels))
  | 'Read and convert to JPEG'
    >> beam.ParDo(ReadImageAndConvertToJpegDoFn())
  | 'Embed and make TFExample' >>
    beam.ParDo(TFExampleFromImageDoFn())
  | 'SerializeToString' >> beam.Map(lambda x:
    x.SerializeToString())
  | 'Save to disk'
    >> beam.io.WriteToTFRecord(opt.output_path,
    file_name_suffix='.tfrecord.gz'))
```

Then, we'll train the model

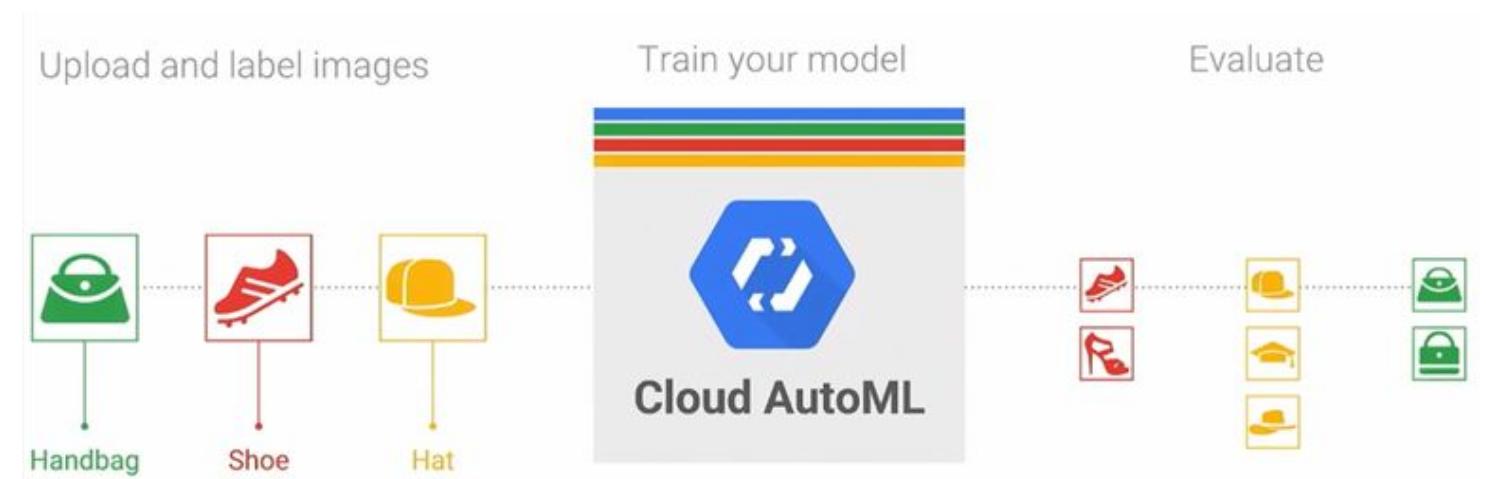


File Name:
T-ICML-O_3_I7_lab_solution:_implementing_transfer_learning

Format: Presenter in Studio

Presenter: Max Lotstein

AutoML is like codeless transfer learning



File Name: T-ICML-O_3_l8_no_data,_no_problem

Format: Presenter in Studio

Presenter: Max Lotstein

Agenda

The Data Scarcity Problem

Data Augmentation

Transfer Learning

No Data, No Problem

No Data, No Problem

- 1 Use a pre-trained model

Google has cutting-edge machine learning APIs



Cloud
Vision API



Cloud
Speech API



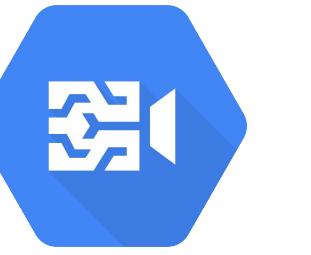
Cloud
Jobs API



Cloud
Translation API



Cloud
Natural
Language API



Cloud Video
Intelligence

Image Recognition with the Cloud Vision API

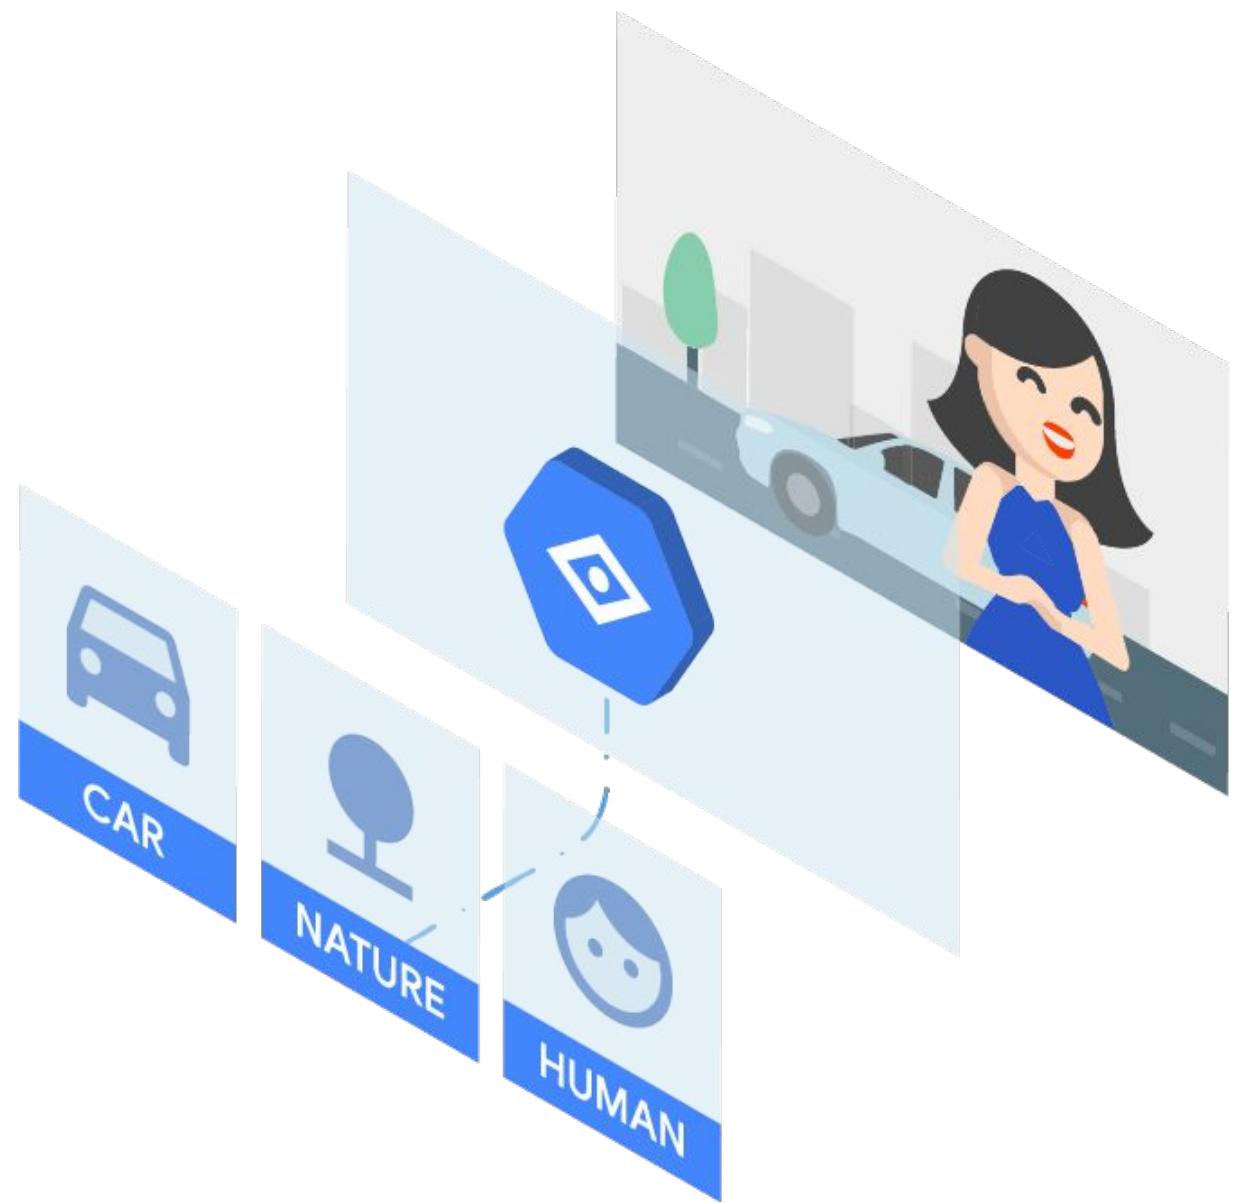
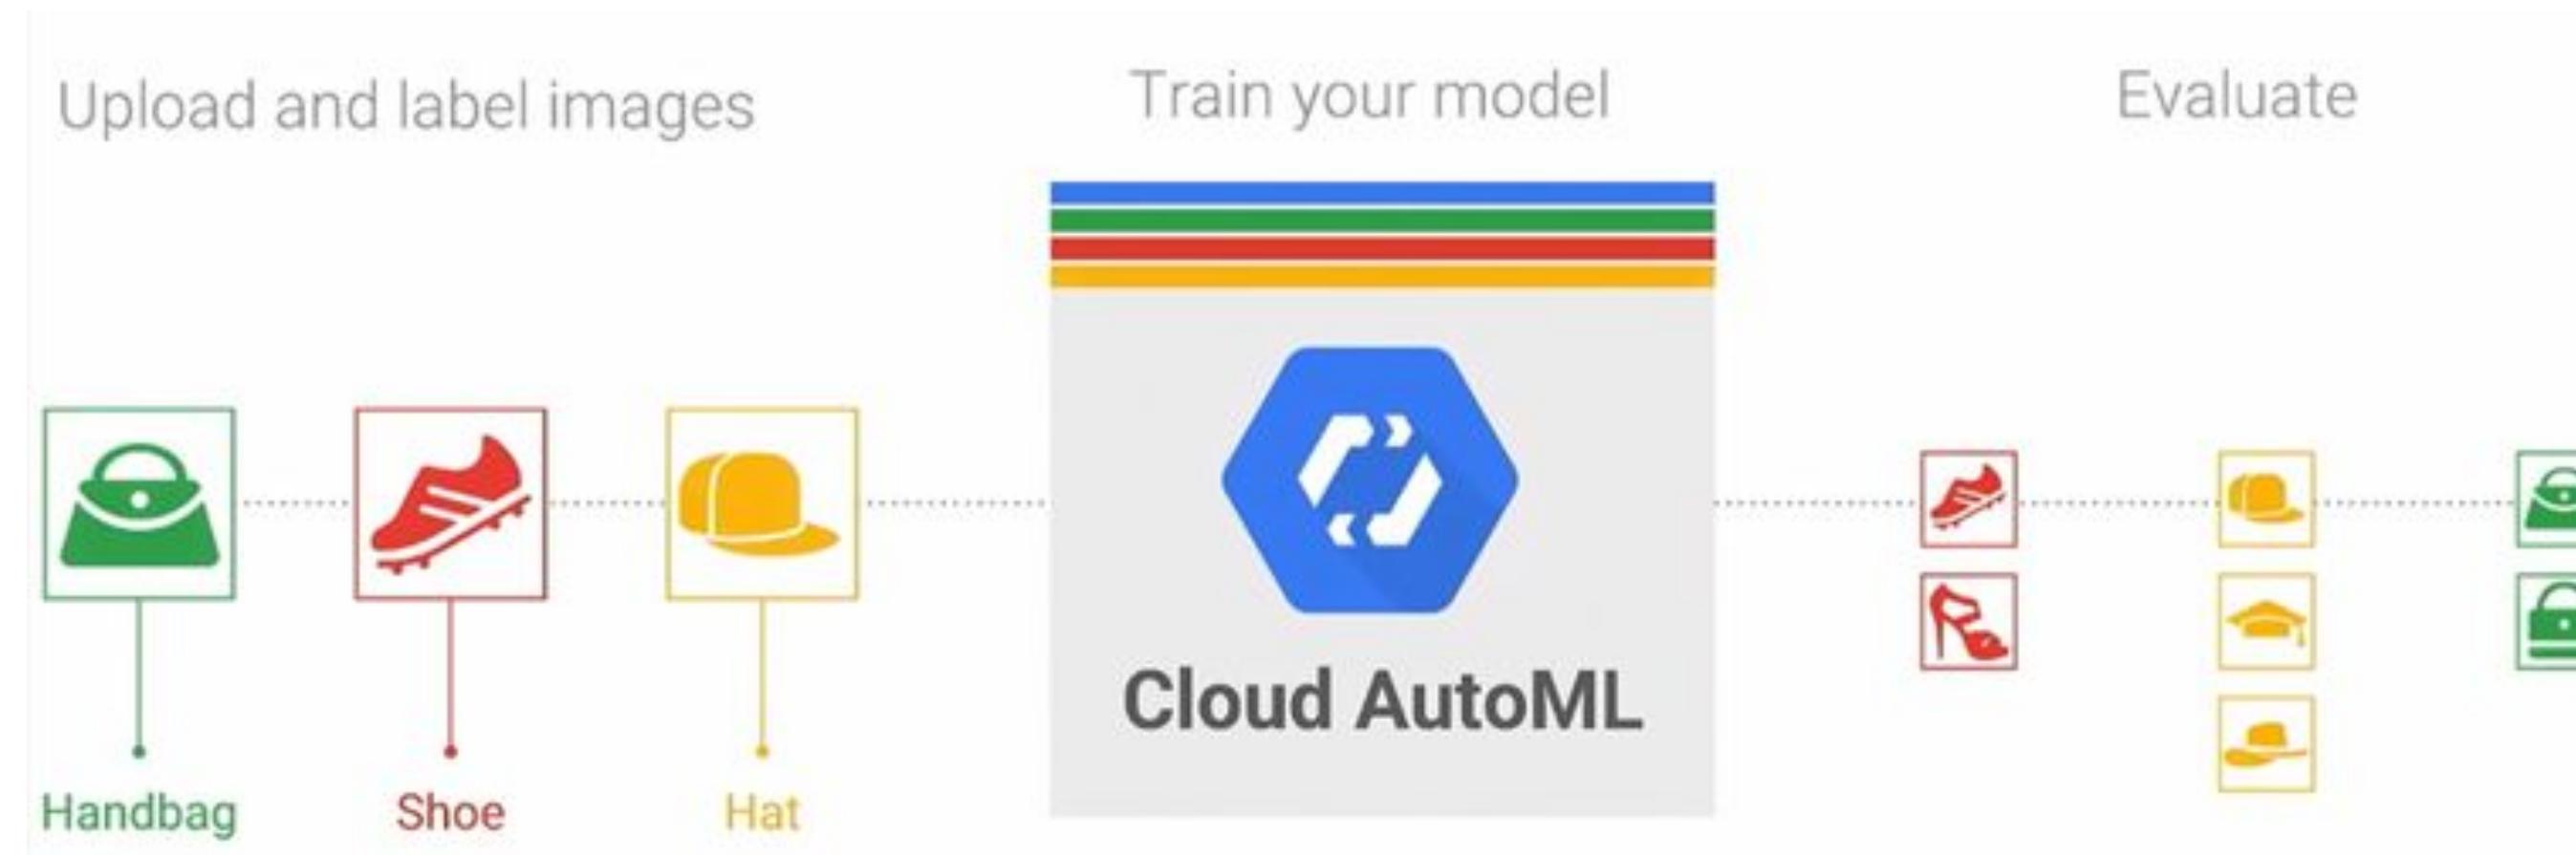


Image Recognition with the Cloud Vision API



No Data, No Problem

- 1 Use a pre-trained model
- 2 Find a partner to label your data

Data scarcity is not as bad as you might think

If you have some labeled data:

- Data augmentation helps you get more
- Transfer learning helps you need less
 - AutoML makes transfer learning codeless

Data scarcity is not as bad as you might think

If you have some labeled data:

- Data augmentation helps you get more
- Transfer learning helps you need less
 - AutoML makes transfer learning codeless

If you don't have any labeled data:

- Machine learning APIs are state-of-the-art
- Partners can help you economically label your data