



---

## Dealing with Longer Sequences



# Advanced ML with TensorFlow on GCP

---

End-to-End Lab on Structured Data ML

Production ML Systems

Image Classification Models

**Sequence Models**

Recommendation Systems



# Agenda

---

**LSTMs and GRUs**

Deep RNNs

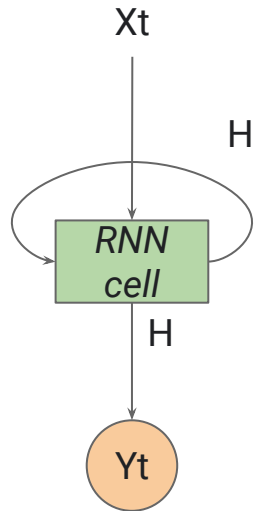
Improving our loss function

Working with real world data

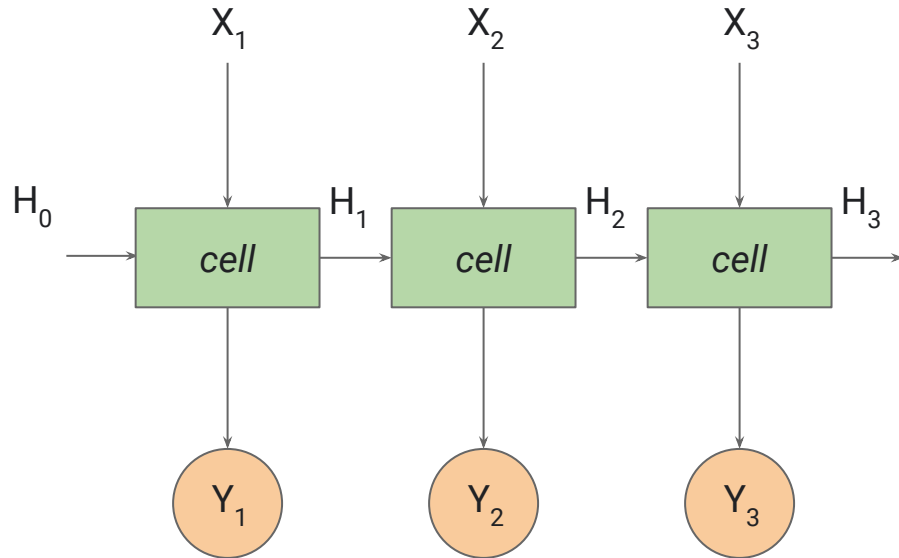


# Problem with RNNs: Long term dependencies

Recurrent View

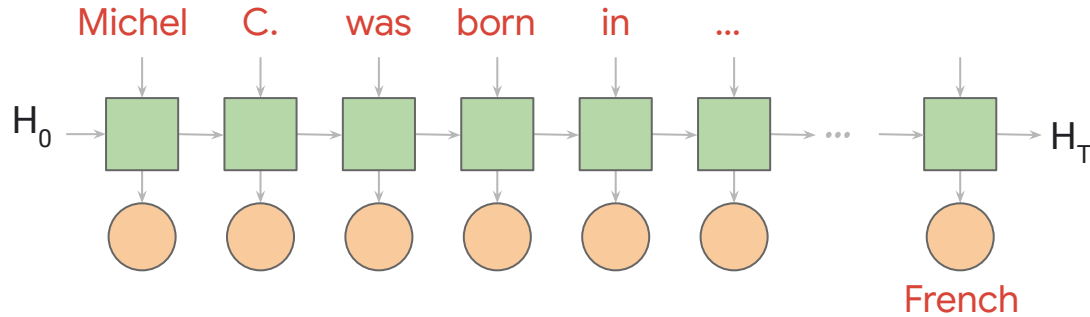


Unrolled View



# Problem with RNNs: Long term dependencies

*"Michel C. was born in Paris, France. His mother tongue is ? ? ? ?"*

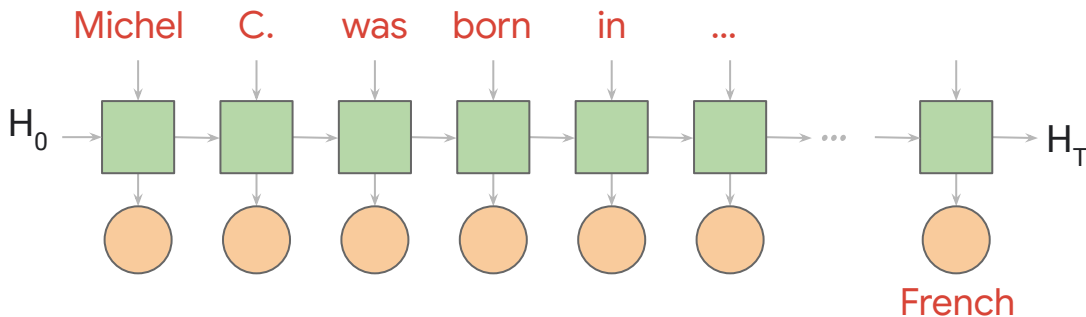


$T$  = max sequence length

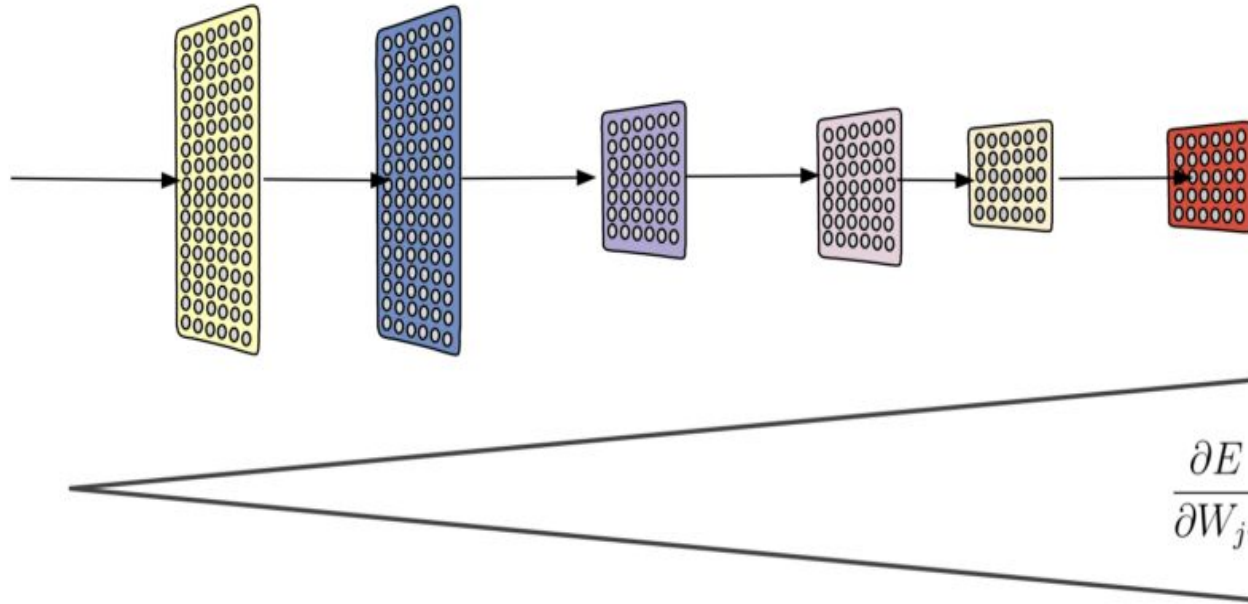


# Problem with RNNs: Long term dependencies

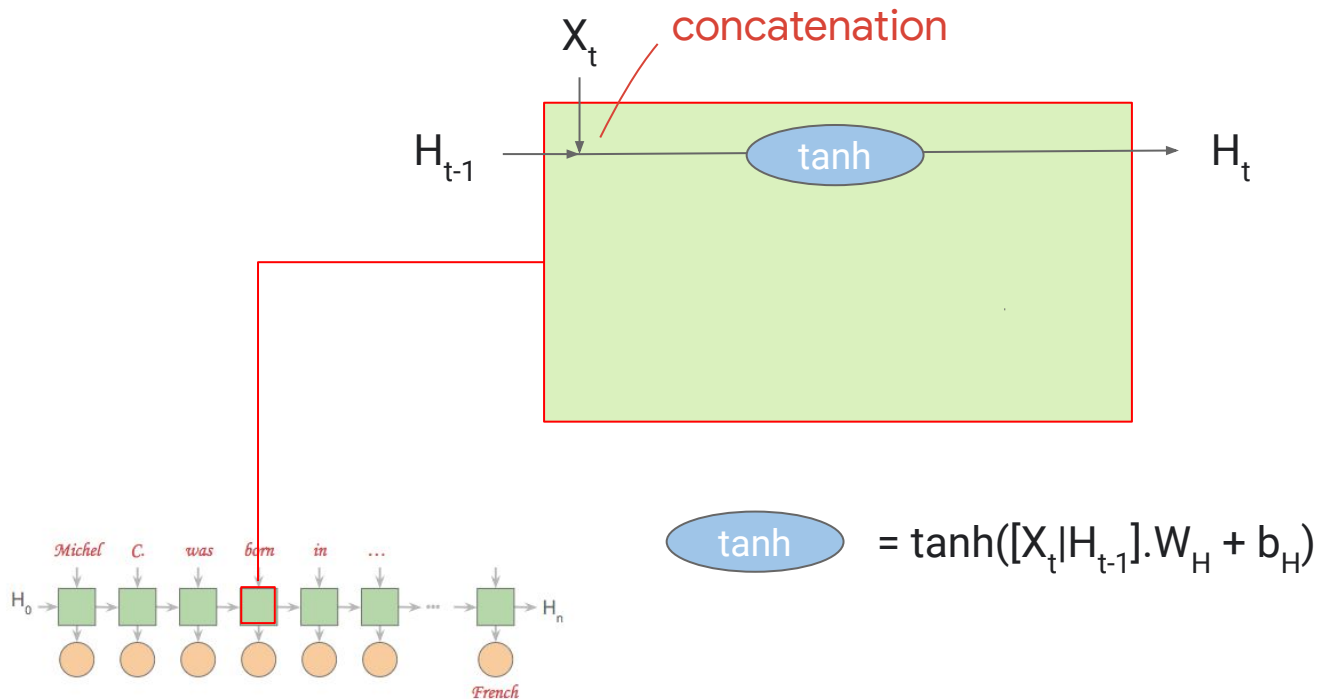
*"Michel C. was born in Paris, France. He is married and has three children. He received a M.S. in neurosciences from the University Pierre & Marie Curie and the Ecole Normale Supérieure in 1987, and then spent most of his career in Switzerland, at the Ecole Polytechnique de Lausanne. He specialized in child and adolescent psychiatry and his first field of research was severe mood disorders in adolescent, topic of his PhD in neurosciences (2002). His mother tongue is ???"*



# Problem with RNNs: Vanishing gradients

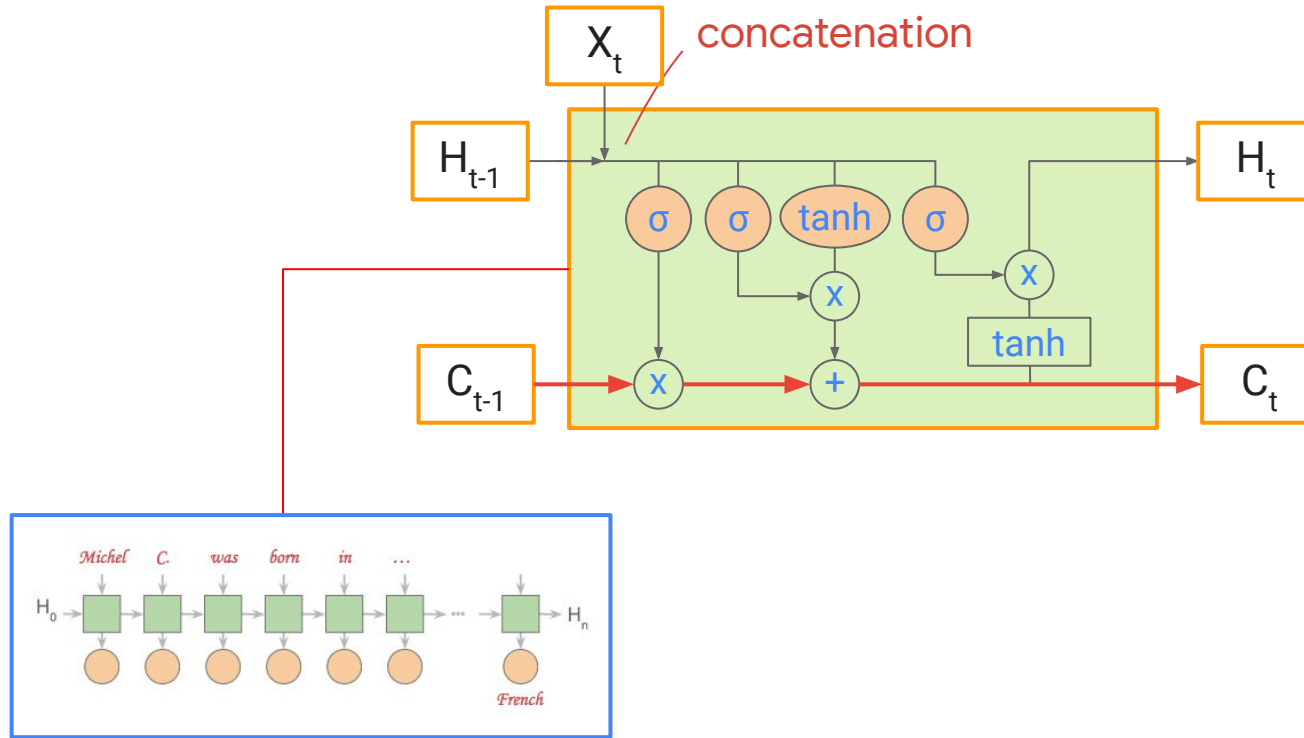


# Simple RNN cell





# LSTM cell



# LSTM cell

LSTM = Long Short Term Memory

concatenate :

$$X = X_t \parallel H_{t-1}$$

vector sizes

$p+n$

forget gate :

$$f = \sigma(X.W_f + b_f) \quad n$$

update gate :

$$u = \sigma(X.W_u + b_u) \quad n$$

output gate :

$$o = \sigma(X.W_o + b_o) \quad n$$

input :

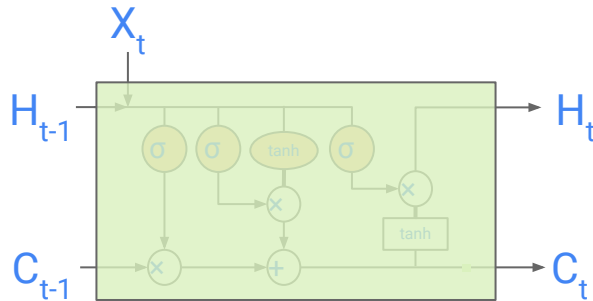
$$X' = \tanh(X.W_c + b_c) \quad n$$

new C :

$$C_t = f * C_{t-1} + u * X' \quad n$$

new H :

$$H_t = o * \tanh(C_t) \quad n$$

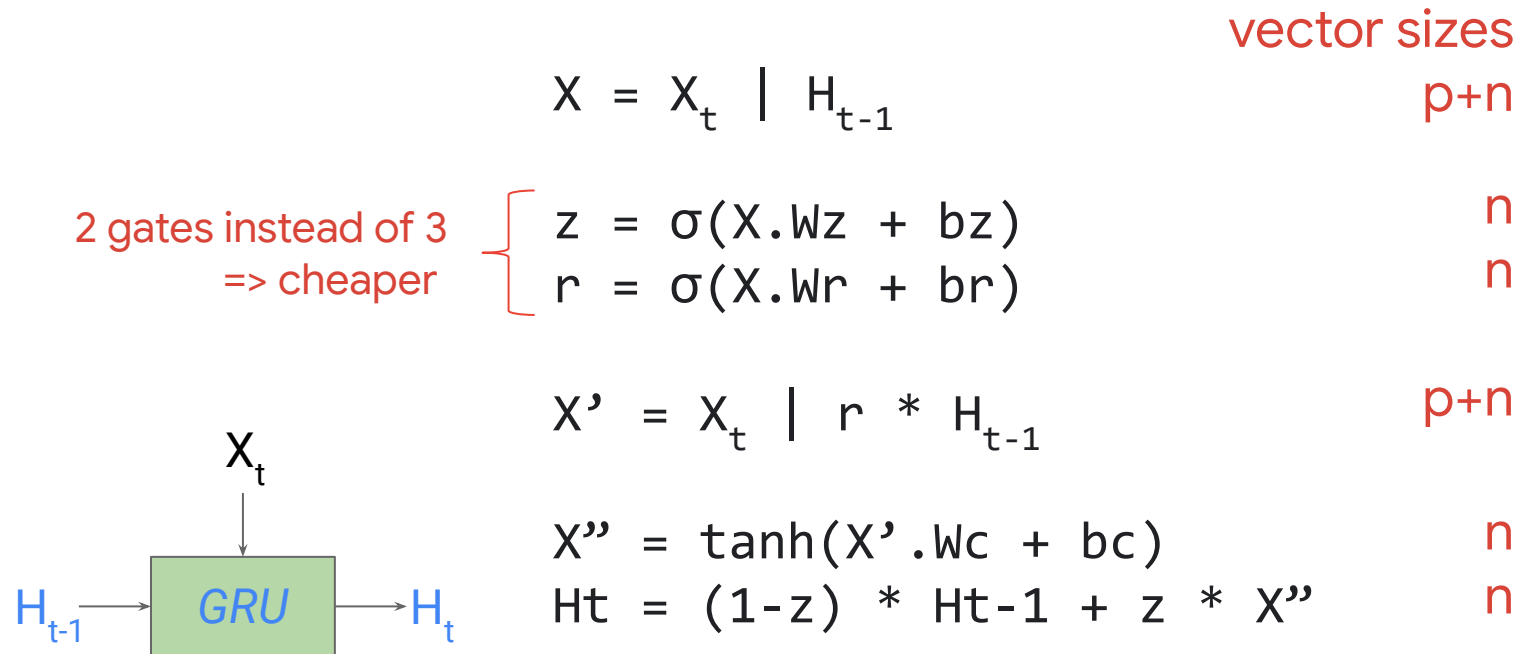


# Our cell has been given control abilities

- 1 What to forget from the cell state.
- 2 What new data to store into the cell state.
- 3 What data from the cell state to expose to the hidden state.



# Gated Recurrent Unit (GRU)



Quiz: How many weight matrices do a simple RNN, LSTM, and GRU cell have respectively?

1. 1, 4, 3
2. 1, 4, 4
3. 1,1,1
4. Depends on the number of time steps



Quiz: How many weight matrices do a simple RNN, LSTM, and GRU cell have respectively?

1. 1, 4, 3

2. 1, 4, 4

3. 1,1,1

4. Depends on the number of time steps



# Use an RNN in TensorFlow

```
CELL_SIZE = 32 # size of the cell's internal state

# 1. Choose RNN Cell type
cell = tf.nn.rnn_cell.GRUCell(CELL_SIZE) # or BasicLSTMCell or BasicRNNCell
# 2. Create RNN by passing cell and tensor of features (x)
outputs, state = tf.nn.dynamic_rnn(cell, x, dtype=tf.float32)

# x needs shape: [BATCH_SIZE, MAX_SEQUENCE_LENGTH, INPUT_DIM]
# outputs has shape: [BATCH_SIZE, MAX_SEQUENCE_LENGTH, CELL_SIZE]

# state has shape: [BATCH_SIZE, CELL_SIZE]
# 3. Pass rnn state through a DNN to get prediction
h1 = tf.layers.dense(state, DNN, activation=tf.nn.relu)
predictions = tf.layers.dense(h1, 1, activation=None) # (BATCH_SIZE, 1)
return predictions
```



# Lab

---

Time series prediction:  
end-to-end (rnn)

In this lab we will continue with the sprinkler problem, this time implementing an RNN to predict the sprinkler height.





# Lab Steps

---

1. Complete `rnn_model` function in `model.py`.
2. Run RNN model locally to make sure code works.



# Agenda

---

LSTMs and GRUs

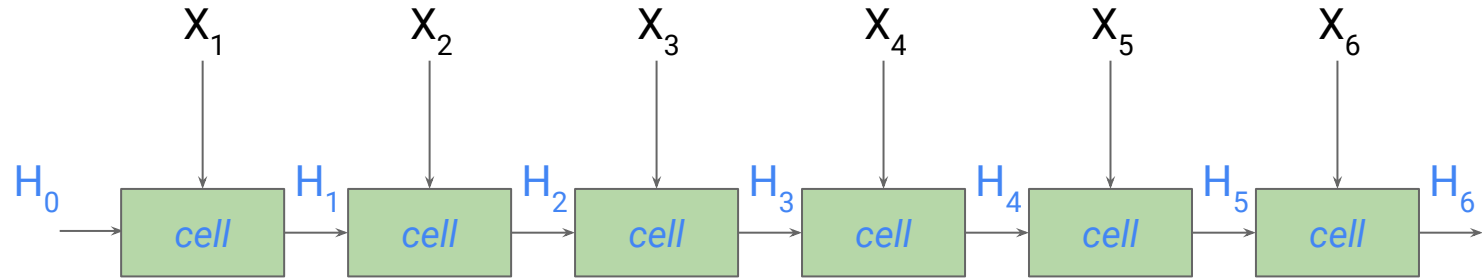
## **Deep RNNs**

Improving our loss function

Working with real world data



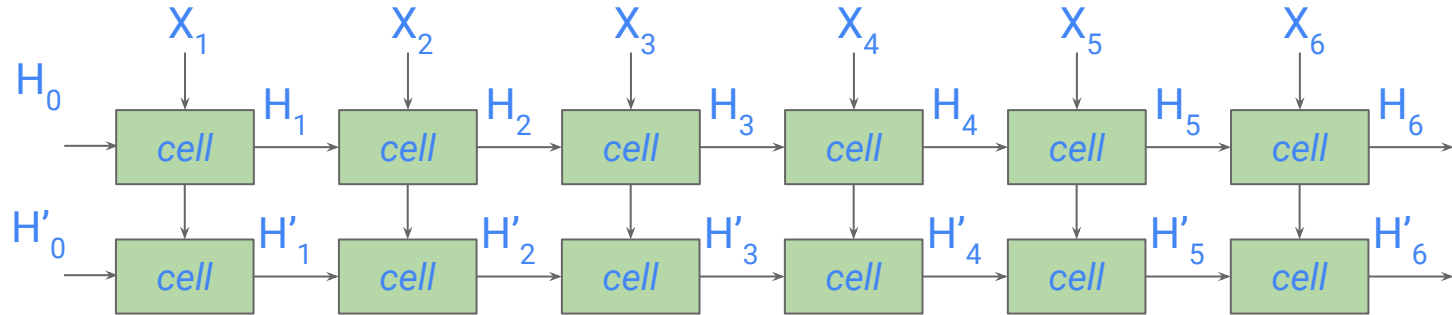
# Shallow RNN



The same weights and biases shared and updated in training loop (across iterations).



# Two layer RNN



$$H_t = \tanh([X_t | H_{t-1}] \cdot W + b)$$

$$H'_t = \tanh([H_t | H'_{t-1}] \cdot W' + b')$$

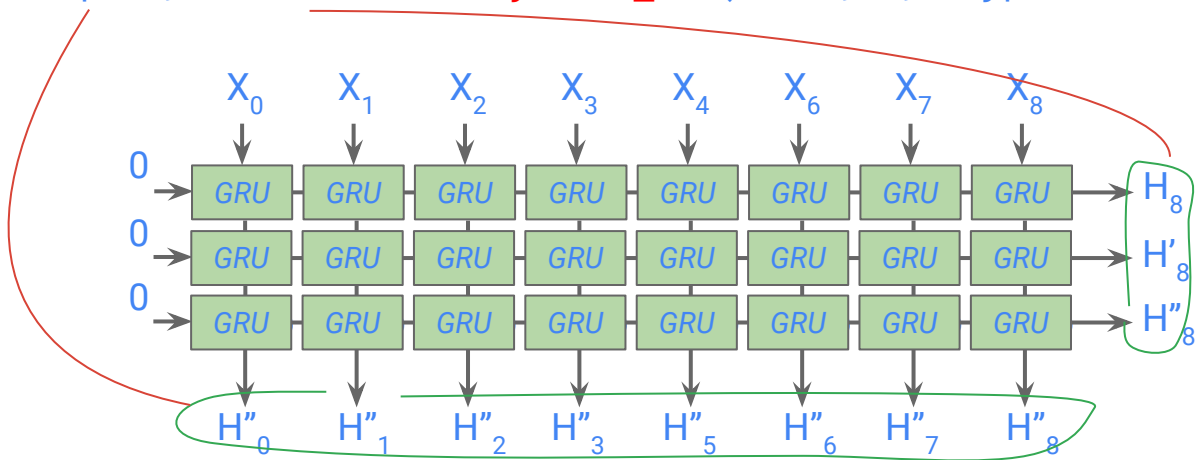


# Deep RNNs in Tensorflow

```
cells = [tf.nn.rnn_cell.GRUCell(CELLSIZE) for _ in range(NLAYERS)]
```

```
mcell = tf.nn.rnn_cell.MultiRNNCell(cells)
```

```
outputs, state = tf.nn.dynamic_rnn(mcell, X, dtype=tf.float32)
```



# Lab

---

Time series prediction:  
end-to-end (rnn2)

In this lab, you will define a  
rnn2 model to find next value  
of a time series.



# Lab Steps

---

1. Complete `rnn_model` function in `model.py`.
2. Run RNN model locally to make sure code works.
3. Run all models in in the cloud.
4. Compare results using Tensorboard.



# Agenda

---

LSTMs and GRUs

Deep RNNs

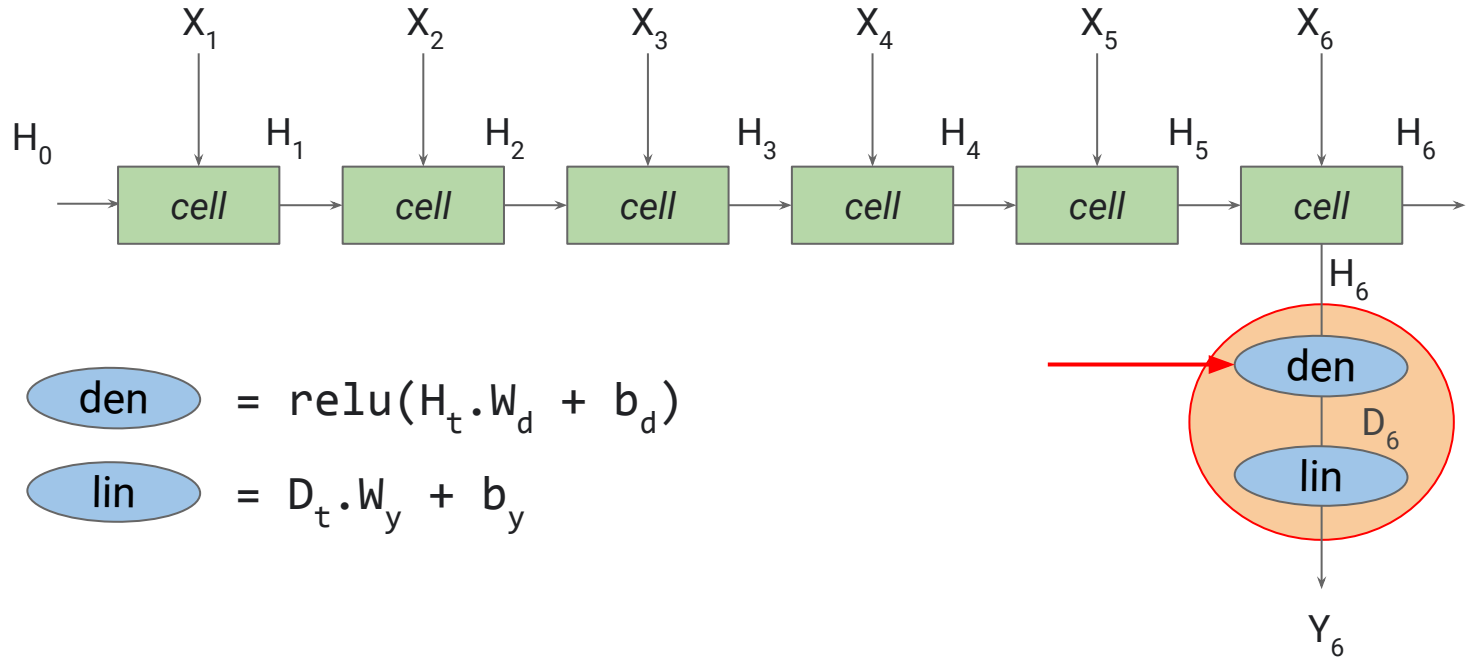
**Improving our loss function**

Working with real world data

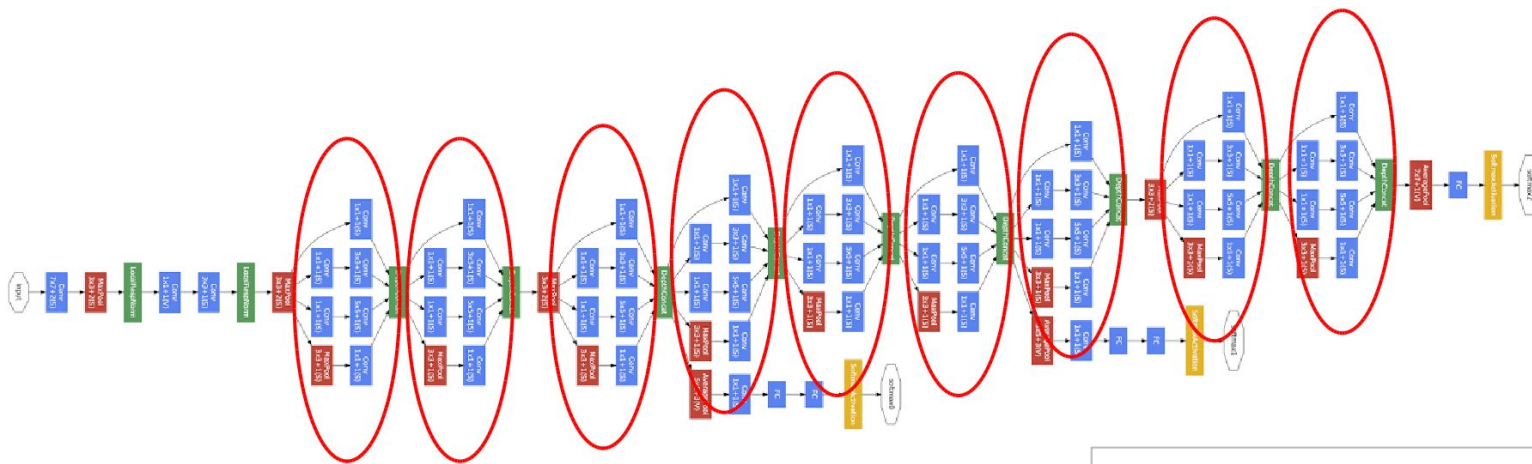




# Improving our loss function



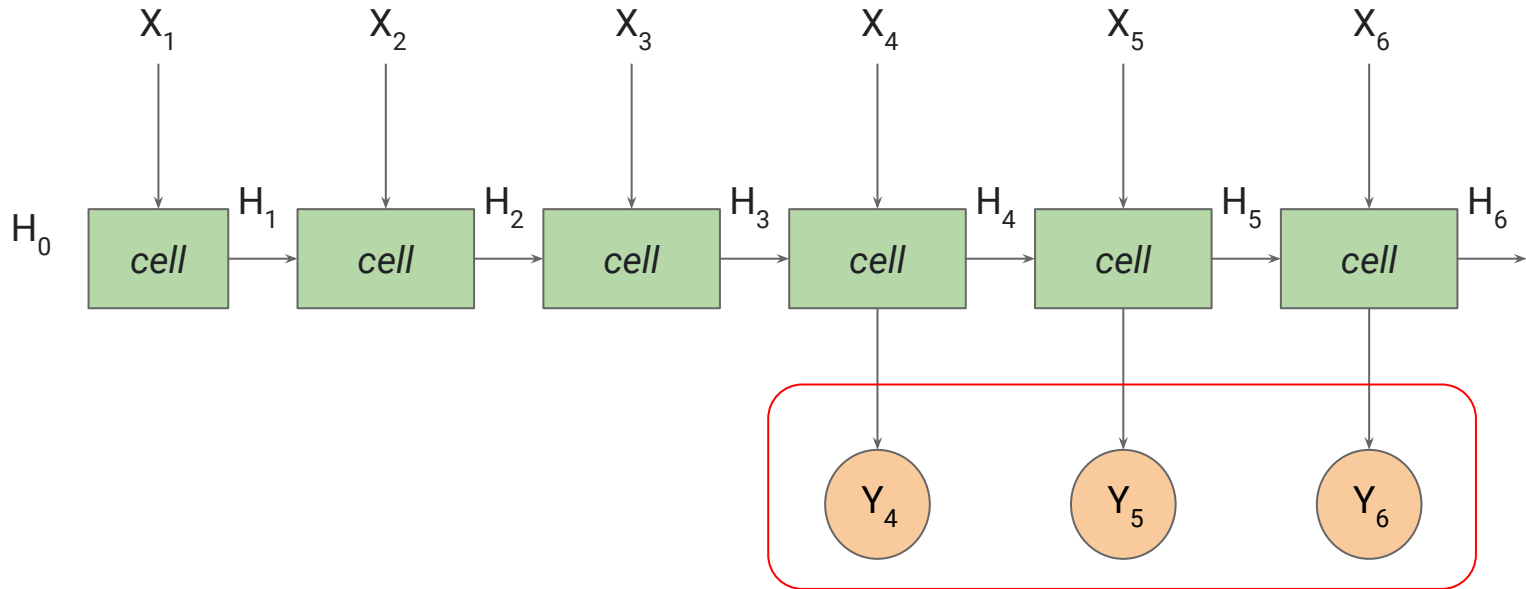
# Inception CNN taking intermediate outputs



Convolution  
Pooling  
Softmax  
Concat/Normalize



# Improving our loss function



Average loss over multiple predictions



# Agenda

---

LSTMs and GRUs

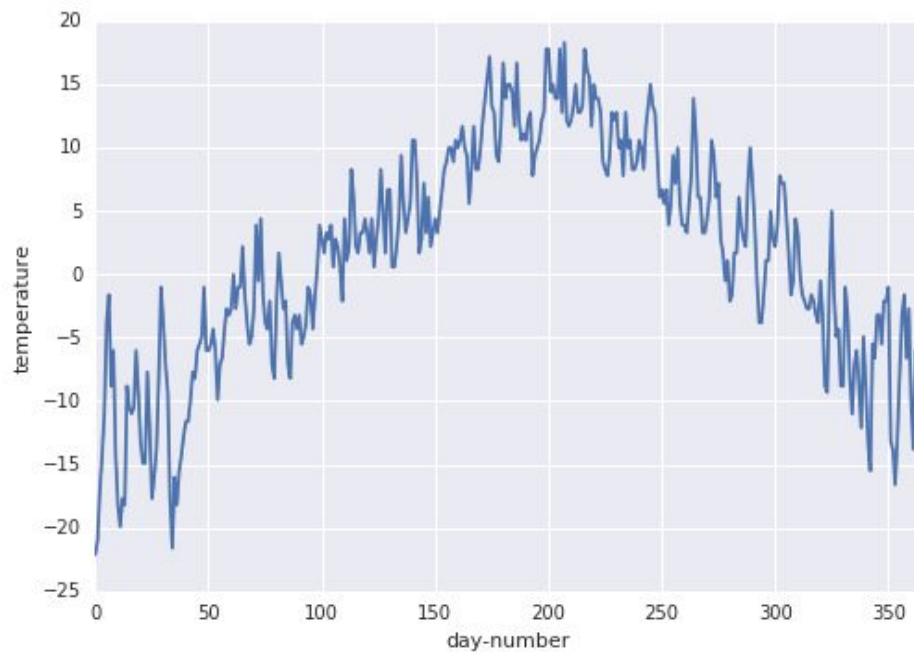
Deep RNNs

Improving our loss function

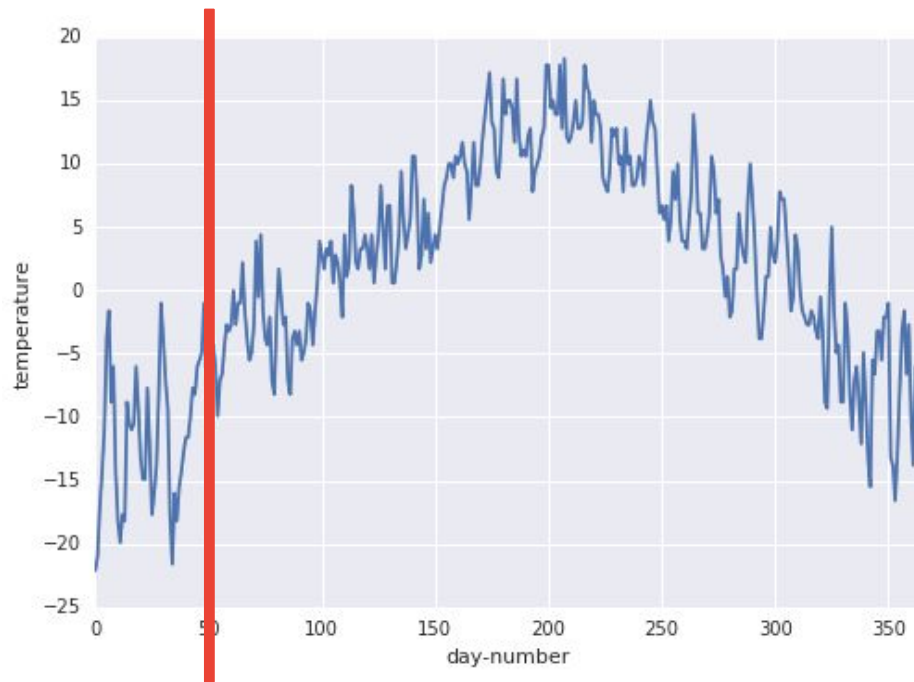
**Working with real world data**



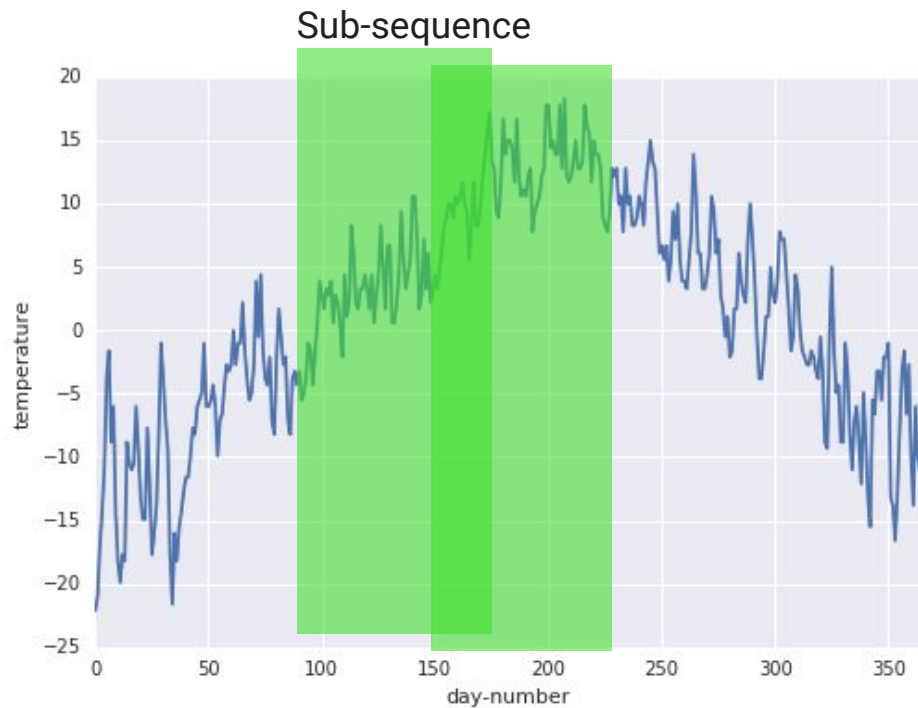
# Working with real world sequential data



# Splitting up sequences



# Splitting up sequences



# Splitting up sequences

Original Sequence

1, 2, 3, 4, 5, 6

Non-overlapping

Sequence 1

1,2,3

Sequence 2

4, 5, 6

---

Overlapping

Sequence 3

3, 4, 5

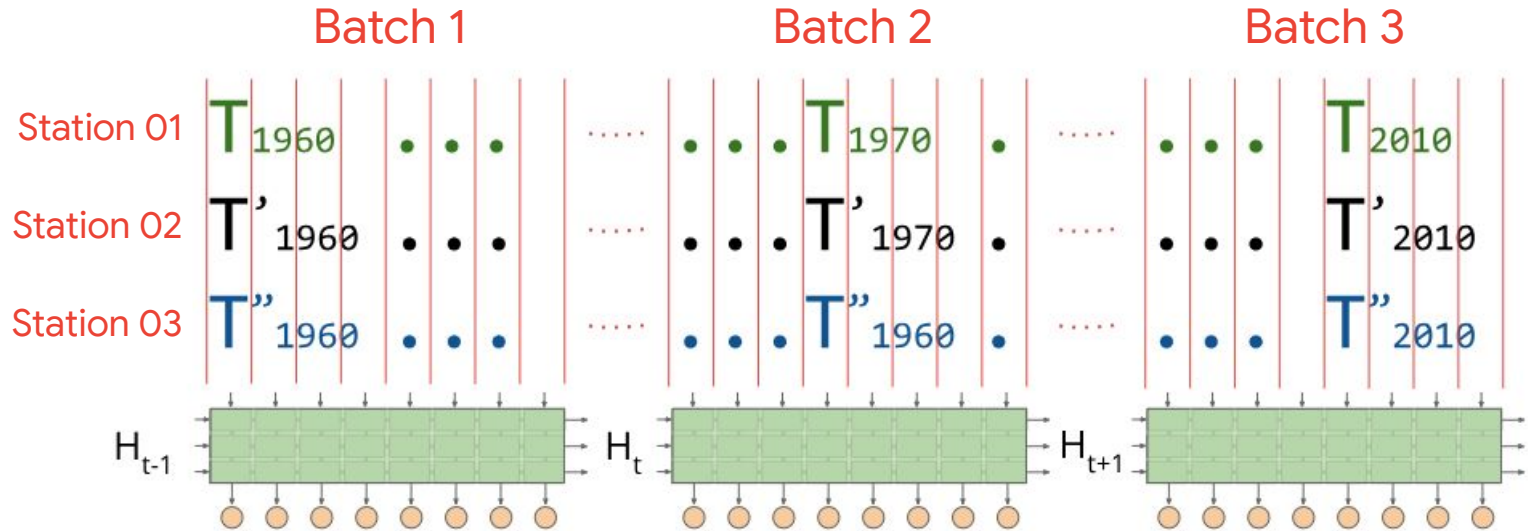
Sequence 4

4, 5, 6

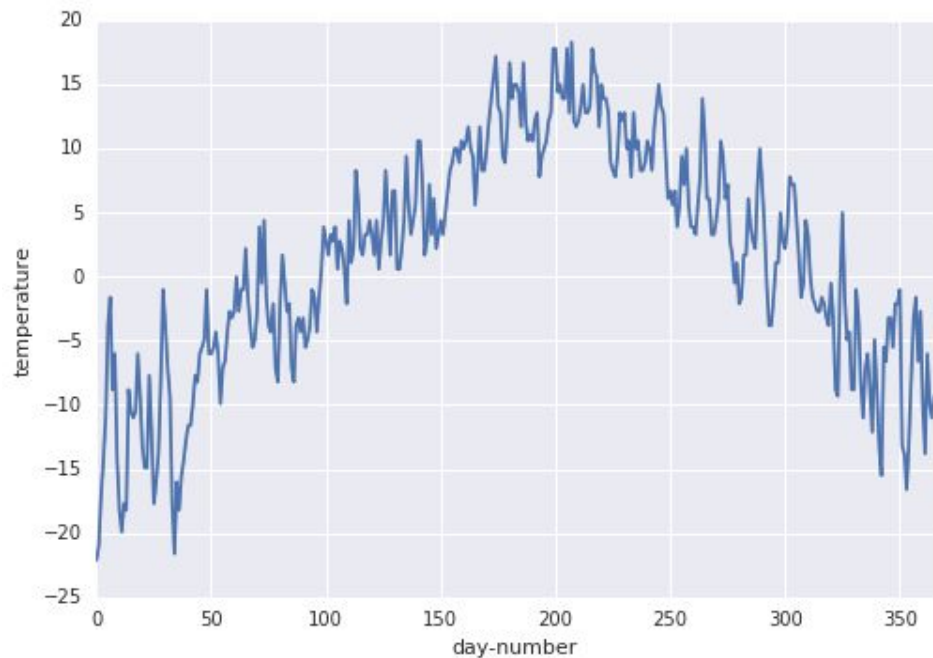




# Propagating state between batches

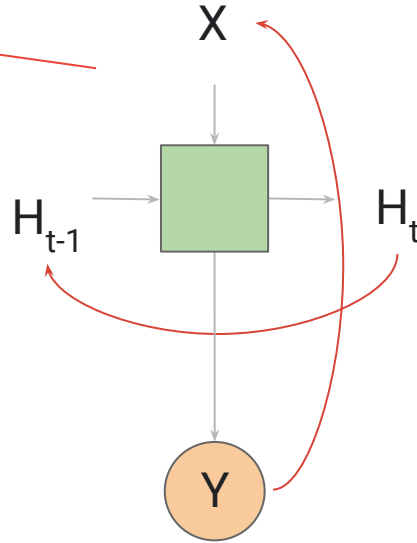


# Splitting up sequences



# Predicting multiple time steps ahead

One value at  
a time



# Other considerations

- 1 Resampling data.
- 2 One model versus multiple models.
- 3 Incorporating non-sequential data.



# Resampling data



# One versus multiple models



Model 1

Station  
#1

Model 2

Station  
#2

Model 3

Station  
#3



One Model

Station  
#1

Station  
#2

Station  
#3



One Model

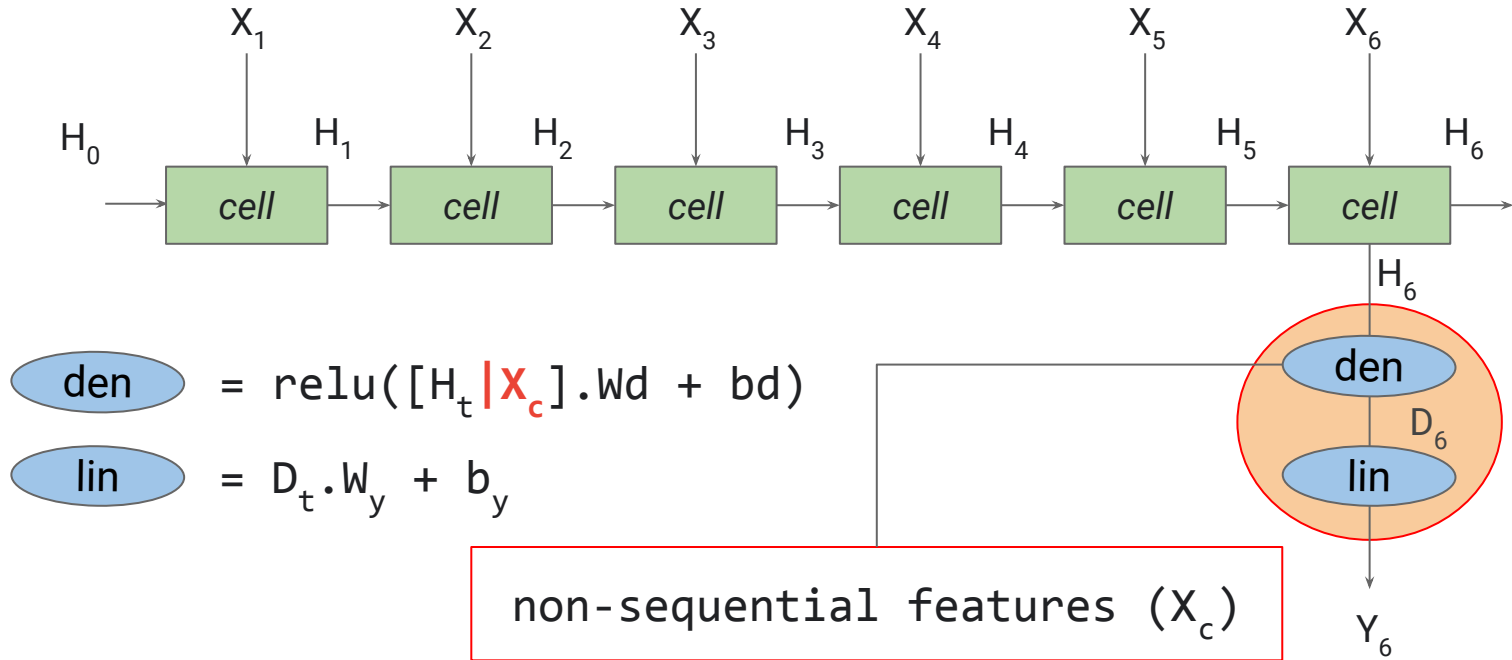
Station  
#1  
LAT  
LONG

Station  
#2  
LAT  
LONG

Station  
#3  
LAT  
LONG



# Incorporating non-sequential data



# Lab

---

Time series prediction:  
Temperature from weather  
stations

In this lab you'll see how to  
resample data, split up  
sequences, generate a  
sequence of predictions, and  
more.





# Lab Steps

---

1. Run the notebook as it is. Look at the data visualisations. Then look at the predictions at the end.
2. Play with the data to find good values for RESAMPLE\_BY and SEQLEN in hyperparameters.
3. Predict N data points ahead instead of only 1 ahead.
4. Adjust hyperparameters and add regularization.





cloud.google.com

